

Data Compression Explained

[Matt Mahoney](#)

Copyright (C) 2010-2012, [Dell, Inc.](#) You are permitted to copy and distribute material from this book provided (1) any material you distribute includes this license, (2) the material is not modified, and (3) you do not charge a fee or require any other considerations for copies or for any works that incorporate material from this book. These restrictions do not apply to normal "fair use", defined as cited quotations totaling less than one page. This book may be downloaded without charge from <http://mattmahoney.net/dc/dce.html>.

Last update: Apr. 15, 2013.

e-reader translations by Alejo Sanchez, Oct. 29, 2011: [mobi](#) (Kindle) and [epub](#) (other readers).

About this Book

This book is for the reader who wants to understand how data compression works, or who wants to write data compression software. Prior programming ability and some math skills will be needed. Specific topics include:

- [1. Information theory](#)
 - [No universal compression](#)
 - [Coding is bounded](#)
 - [Modeling is not computable](#)
 - [Compression is an artificial intelligence problem](#)
- [2. Benchmarks](#)
 - [Calgary corpus results](#)
- [3. Coding](#)
 - [Huffman](#)
 - [arithmetic](#)
 - [asymmetric binary](#)
 - [numeric codes](#) ([unary](#), [Rice](#), [Golomb](#), [extra bit](#))
 - [archive formats](#) ([error detection](#), [encryption](#))
- [4. Modeling](#)
 - [Fixed order](#): [bytewise](#) [bitwise](#), [indirect](#)
 - [Variable order](#): [DMC](#), [PPM](#), [CTW](#)
 - [Context mixing](#): [linear mixing](#), [logistic mixing](#), [SSE](#), [indirect SSE](#), [match](#), [PAQ](#), [ZPAQ](#), [Crinkler](#)
- [5. Transforms](#)
 - [RLE](#)
 - [LZ77](#) ([LZSS](#), [deflate](#), [LZMA](#), [LZX](#), [ROLZ](#), [LZP](#), [snappy](#), [deduplication](#))
 - [LZW](#) and [dictionary encoding](#)
 - [Symbol ranking](#)
 - [BWT](#) ([context sorting](#), [inverse](#), [bzip2](#), [BBB](#), [MSufSort v2](#) and [v3](#), [Itoh-Tanaka](#), [DivSufSort](#), [Bijective](#))
 - [Predictive filtering](#) ([delta coding](#), [color transform](#), [linear filtering](#))
 - [Specialized transforms](#) ([E8E9](#), [precomp](#))
 - [Huffman pre-coding](#)
- [6. Lossy compression](#)
 - [Images](#) ([BMP](#), [GIF](#), [PNG](#), [TIFF](#), [JPEG](#))
 - [JPEG recompression](#) ([Stuffit](#), [PAQ](#), [WinZIP PackJPG](#))
 - [Video](#) ([NTSC](#), [MPEG](#))
 - [Audio](#) ([CD](#), [MP3](#), [AAC](#), [Dolby](#), [Vorbis](#))
- [Conclusion](#)
- [Acknowledgements](#)
- [References](#)

This book is intended to be self contained. Sources are linked when appropriate, but you don't need to click on them to understand the material.

1. Information Theory

Data compression is the art of reducing the number of bits needed to store or transmit data. Compression can be either lossless or lossy. Losslessly compressed data can be decompressed to exactly its original value. An example is 1848 [Morse Code](#). Each letter of the alphabet is coded as a sequence of dots and dashes. The most common letters in English like E and T receive the shortest codes. The least common like J, Q, X, and Z are assigned the longest codes.

All data compression algorithms consist of at least a model and a coder (with optional preprocessing transforms). A model estimates the probability distribution (E is more common than Z). The coder assigns shorter codes to the more likely symbols. There are efficient and optimal solutions to the coding problem. However, optimal modeling has been proven not computable. Modeling (or equivalently, prediction) is both an artificial intelligence (AI) problem and an art.

Lossy compression discards "unimportant" data, for example, details of an image or audio clip that are not perceptible to the eye or ear. An example is the 1953 [NTSC](#) standard for broadcast color TV, used until 2009. The human eye is less sensitive to fine detail between colors of equal brightness (like red and green) than it is to brightness (black and white). Thus, the color signal is transmitted with less resolution over a narrower frequency band than the monochrome signal.

Lossy compression consists of a transform to separate important from unimportant data, followed by lossless compression of the important part and discarding the rest. The transform is an AI problem because it requires understanding what the human brain can and cannot perceive.

Information theory places hard limits on what can and cannot be compressed losslessly, and by how much:

1. There is no such thing as a "universal" compression algorithm that is guaranteed to compress any input, or even any input above a certain size. In particular, it is not possible to compress random data or compress recursively.
2. Given a model (probability distribution) of your input data, the best you can do is code symbols with probability p using $\log_2 1/p$ bits. Efficient and optimal codes are known.
3. Data has a universal but uncomputable probability distribution. Specifically, any string x has probability (about) $2^{-|M|}$ where M is the shortest possible description of x , and $|M|$ is the length of M in bits, almost independent of the language in which M is written. However there is no general procedure for finding M or even estimating $|M|$ in any language. There is no algorithm that tests for randomness or tells you whether a string can be compressed any further.

1.1. No Universal Compression

This is proved by the [counting argument](#). Suppose there were a compression algorithm that could compress all strings of at least a certain size, say, n bits. There are exactly 2^n different binary strings of length n . A universal compressor would have to encode each input differently. Otherwise, if two inputs compressed to the same output, then the decompressor would not be able to decompress that output correctly. However there are only $2^n - 1$ binary strings shorter than n bits.

In fact, the vast majority of strings cannot be compressed by very much. The fraction of strings that can be compressed from n bits to m bits is at most 2^{m-n} . For example, less than 0.4% of strings can be compressed by one byte.

Every compressor that can compress any input must also expand some of its input. However, the expansion never needs to be more than one symbol. Any compression algorithm can be modified by adding one bit to indicate that the rest of the data is stored uncompressed.

The counting argument applies to systems that would recursively compress their own output. In general, compressed data appears random to the algorithm that compressed it so that it cannot be compressed again.

1.2. Coding is Bounded

Suppose we wish to compress the digits of π , e.g. "314159265358979323846264...". Assume our model is that each digit occurs with probability 0.1, independent of any other digits. Consider 3 possible binary codes:

Digit	BCD	Huffman	Binary
0	0000	000	0
1	0001	001	1
2	0010	010	10
3	0011	011	11
4	0100	100	100
5	0101	101	101
6	0110	1100	110
7	0111	1101	111
8	1000	1110	1000
9	1001	1111	1001
---	----	----	----
bpc	4.0	3.4	not valid

Using a BCD (binary coded decimal) code, π would be encoded as 0011 0001 0100 0001 0101... (Spaces are shown for readability only). The compression ratio is 4 bits per character (4 bpc). If the input was ASCII text, the output would be compressed 50%. The decompressor would decode the data by dividing it into 4 bit strings.

The [Huffman code](#) would code π as 011 001 100 001 101 1111... The decoder would read bits one at a time and decode a digit as soon as it found a match in the table (after either 3 or 4 bits). The code is uniquely decodable because no code is a prefix of any other code. The compression ratio is 3.4 bpc.

The binary code is not uniquely decodable. For example, 111 could be decoded as 7 or 31 or 13 or 111.

There are better codes than the Huffman code given above. For example, we could assign Huffman codes to pairs of digits. There are 100 pairs each with probability 0.01. We could assign 6 bit codes (000000 through 011011) to 00 through 27, and 7 bits (0111000 through 1111111) to 28 through 99. The average code length is 6.72 bits per pair of digits, or 3.36 bpc. Similarly, coding groups of 3 digits using 9 or 10 bits would yield 3.3253 bpc.

Shannon and Weaver (1949) proved that the best you can do for a symbol with probability p is assign a code of length $\log_2 1/p$. In this example, $\log_2 1/0.1 = 3.3219$ bpc.

Shannon defined the expected information content or equivocation (now called [entropy](#)) of a random variable X as its expected code length. Suppose X may have values $X_1, X_2, \dots$ and that each X_i has probability $p(i)$. Then the entropy of X is $H(X) = E[\log_2 1/p(X)] = \sum_i p(i) \log_2 1/p(i)$. For example, the entropy of the digits of π , according to our model, is 10 $(0.1 \log_2 1/0.1) = 3.3219$ bpc. There is no smaller code for this model that could be decoded unambiguously.

The information content of a set of strings is at most the sum of the information content of the individual strings. If X and Y are strings, then $H(X, Y) \leq H(X) + H(Y)$. If they are equal, then X and Y are *independent*. Knowing one string would tell you nothing about the other.

The *conditional entropy* $H(X|Y) = H(X, Y) - H(Y)$ is the information content of X given Y . If X and Y are independent, then $H(X|Y) = H(X)$.

If X is a string of symbols $x_1 x_2 \dots x_n$, then by the [chain rule](#), $p(X)$ may be expressed as a product of the sequence of symbol predictions conditioned on previous symbols: $p(X) = \prod_i p(x_i | x_{1..i-1})$. Likewise, the information content $H(X)$ of random string X is the sum of the conditional entropies of each symbol given the previous symbols: $H(X) = \sum_i H(x_i | x_{1..i-1})$.

Entropy is both a measure of uncertainty and a lower bound on expected compression. The entropy of an information source is the expected limit to which you can compress it. There are efficient coding methods, such as arithmetic codes, which are for all practical purposes optimal in this sense. It should be emphasized, however, that entropy can only be calculated for a known probability distribution. But in general, the model is not known.

1.3. Modeling is Not Computable

We modeled the digits of π as uniformly distributed and independent. Given that model, Shannon's coding theorem places a hard limit on the best compression that could be achieved. However, it is possible to use a better model. The digits of π are not really random. The digits are only unknown until you compute them. An intelligent compressor might recognize the digits of π and encode it as a description meaning "the first million digits of pi", or as a program that reconstructs the data when run. With our previous model, the best we could do is $(10^6 \log_2 10)/8 \approx 415,241$ bytes. Yet, there are very small programs that can output π , some as small as [52 bytes](#).

The counting argument says that most strings are not compressible. So it is a rather remarkable fact that most strings that we care about, for example English text, images, software, sensor readings, and DNA, are in fact compressible. These strings generally have short descriptions, whether they are described in English or as a program in C or x86 machine code.

Solomonoff (1960, 1964), Kolmogorov (1965), and Chaitin (1966) independently proposed a universal a-priori probability distribution over strings based on their minimum description length. The [algorithmic probability](#) of a string x is defined as the fraction of random programs in some language L that output x , where each program M is weighted by $2^{-|M|}$ and $|M|$ is the length of M in bits. This probability is dominated by the shortest such program. We call this length the [Kolmogorov complexity](#) $K_L(x)$ of x .

Algorithmic probability and complexity of a string x depend on the choice of language L , but only by a constant that is independent of x . Suppose that M_1 and M_2 are encodings of x in languages L_1 and L_2 respectively. For example, if L_1 is C++, then M_1 would be a program in C++ that outputs x . If L_2 is English, the M_2 would be a description of x with just enough detail that it would allow you to write x exactly. Now it is possible for any pair of languages to write in one language a compiler or interpreter or rules for understanding the other language. For example, you could write a description of the C++ language in English so that you could (in theory) read any C++ program and predict its output by "running" it in your head. Conversely, you could (in theory) write a program in C++ that input an English language description and translated it into C++. The size of the language description or compiler does not depend on x in any way. Then for any description M_1 in any language L_1 of any x , it is possible to find a description M_2 in any other language L_2 of x by appending to M_1 a fixed length description of L_1 written in L_2 .

It is not proven that algorithmic probability is a true universal prior probability. Nevertheless it is widely accepted on empirical grounds because of its success in sequence prediction and machine learning over a wide range of data types. In [machine learning](#), the [minimum description length](#) principle of choosing the simplest hypothesis consistent with the training data applies to a wide range of algorithms. It formalizes [Occam's Razor](#). Occam noted in the 14th century, that (paraphrasing) "the simplest answer is usually the correct answer". Occam's Razor is universally applied in all of the sciences because we know from experience that the simplest or most elegant (i.e. shortest) theory that explains the data tends to be the best predictor of future experiments.

To summarize, the best compression we can achieve for any string x is to encode it as the shortest program M in some language L that outputs x . Furthermore, the choice of L becomes less important as the strings get longer. All that remains is to find a procedure that finds M for any x in some language L . However, Kolmogorov proved that there is no such procedure in any language. In fact, there is no procedure that just gives you $|M|$, the length of the shortest possible description of x , for all x . Suppose there were. Then it would be possible to describe "the first string that cannot be described in less than a million bits" leading to the paradox that we had just done so. (By "first", assume an ordering over strings from shortest to longest, breaking ties lexicographically).

Nor is there a general test for randomness. Li and Vitanyi (2007) give a simple proof of a stronger variation of [Gödel's first incompleteness theorem](#), namely that in any consistent, formal system (a set of axioms and rules of inference) powerful enough to describe statements in arithmetic, that there is at least one true statement that cannot be proven. Li and Vitanyi prove that if the system is sound (you can't prove any false statements) then there are an infinite number of true but unprovable statements. In particular, there are only a finite number of statements of the form " x is random" ($K(x) \geq |x|$) that can be proven, out of an infinite number of possible finite strings x . Suppose otherwise. Then it would be possible to enumerate all proofs (lists of axioms and applications of the rules of inference) and describe "the string x such that it is the first to be proven to be a million bits long and random" in spite of the fact that we just gave a short description of it. If F describes any

formal system, then the longest string that can be proven to be random is never much longer than F , even though there are an infinite number of longer strings and most of them are random.

Nor is there a general test to determine whether the compression of a string can be improved any further, because that would be equivalent to testing whether the compressed string is random.

Because determining the length of the shortest descriptions of strings is not computable, neither is optimal compression. It is not hard to find difficult cases. For example, consider the short description "a string of a million zero bytes compressed with AES in CBC mode with key 'foo'". To any program that does not know the key, the data looks completely random and incompressible.

1.4. Compression is an Artificial Intelligence Problem

Prediction is intuitively related to understanding. If you understand a sequence, then you can predict it. If you understand a language, then you could predict what word might appear next in a paragraph in that language. If you understand an image, then you could predict what might lie under portions of the image that are covered up. Alternatively, random data has no meaning and is not predictable. This intuition suggests that prediction or compression could be used to test or measure understanding. We now formalize this idea.

Optimal compression, if it were computable, would optimally solve the artificial intelligence (AI) problem under two vastly different definitions of "intelligence": the [Turing test](#) (Turing, 1950), and [universal intelligence](#) (Legg and Hutter, 2006).

Turing first proposed a test for AI to sidestep the philosophically difficult question (which he considered irrelevant) of whether machines could think. This test, now known as the Turing test, is now widely accepted. The test is a game played by two humans who have not previously met and the machine under test. One human (the judge) communicates with the other human (the confederate) and with the machine through a terminal. Both the confederate and the machine try to convince the judge that each is human. If the judge cannot guess correctly which is the machine 70% of the time after 10 minutes of interaction, then the machine is said to have AI. Turing gave the following example of a possible dialogue:

Q: Please write me a sonnet on the subject of the Forth Bridge.

A: Count me out on this one. I never could write poetry.

Q: Add 34957 to 70764.

A: (Pause about 30 seconds and then give as answer) 105621.

Q: Do you play chess?

A: Yes.

Q: I have K at my K1, and no other pieces. You have only K at K6 and R at R1. It is your move. What do you play?

A: (After a pause of 15 seconds) R-R8 mate.

It should be evident that compressing transcripts like this requires the ability to compute a model of the form $p(A|Q) = p(QA)/P(Q)$ where Q is the context up to the current question, and A is the response. But if a model could make such predictions accurately, then it could also generate responses indistinguishable from that of a human.

Predicting transcripts is similar to the problem to predicting ordinary written language. It requires in either case vast, real-world knowledge. Shannon (1950) estimated that the information content of written case-insensitive English without punctuation is 0.6 to 1.3 bits per character, based on experiments in which human subjects guessed successive characters in text with the help of letter n-gram frequency tables and dictionaries. The uncertainty is due not so much to variation in subject matter and human skill as it is due to the fact that different probability assignments lead to the same observed guessing sequences. Nevertheless, the best text compressors are only now compressing near the upper end of this range.

Legg and Hutter proposed the second definition, universal intelligence, to be far more general than Turing's human intelligence. They consider the problem of reward-seeking agents in completely arbitrary environments described by random programs. In this model, an agent communicates with an environment by sending and receiving symbols. The environment also sends a reinforcement or reward signal to the agent. The goal of the agent is to maximize accumulated reward. Universal intelligence is defined as the expected reward over all possible environments, where the probability of each environment described by a program M is algorithmic,

proportional to $2^{-|M|}$. Hutter (2004, 2007) proved that the optimal (but not computable) strategy for the agent is to guess after each input that the distribution over M is dominated by the shortest program consistent with past observation.

Hutter calls this strategy [AIXI](#). It is, of course, is just our uncomputable compression problem applied to a transcript of past interaction. AIXI may also be considered a formal statement and proof of Occam's Razor. The best predictor of the future is the simplest or shortest theory that explains the past.

1.5. Summary

There is no such thing as universal compression, recursive compression, or compression of random data.

Most strings are random. Most meaningful strings are not.

Compression = modeling + coding. Coding is a solved problem. Modeling is provably not solvable.

Compression is both an art and an artificial intelligence problem. The key to compression is to understand the data you want to compress.

2. Benchmarks

A data compression benchmark measures compression ratio over a data set, and sometimes memory usage and speed on a particular computer. Some benchmarks evaluate size only, in order to avoid hardware dependencies. Compression ratio is often measured by the size of the compressed output file, or in bits per character (bpc) meaning compressed bits per uncompressed byte. In either case, smaller numbers are better. 8 bpc means no compression. 6 bpc means 25% compression or 75% of original size.


Generally there is a 3 way trade off between size, speed, and memory usage. The top ranked compressors by size require a lot of computing resources.

2.1. Calgary Corpus

The [Calgary corpus](#) is the oldest compression benchmark still in use. It was created in 1987 and described in a survey of text compression models in 1989 (Bell, Witten and Cleary, 1989). It consists of 14 files with a total size of 3,141,622 bytes as follows:

111,261	BIB	- ASCII text in UNIX "refer" format - 725 bibliographic references.
768,771	BOOK1	- unformatted ASCII text - Thomas Hardy: Far from the Madding Crowd.
610,856	BOOK2	- ASCII text in UNIX "troff" format - Witten: Principles of Computer Speech.
102,400	GEO	- 32 bit numbers in IBM floating point format - seismic data.
377,109	NEWS	- ASCII text - USENET batch file on a variety of topics.
21,504	OBJ1	- VAX executable program - compilation of PROGP.
246,814	OBJ2	- Macintosh executable program - "Knowledge Support System".
53,161	PAPER1	- UNIX "troff" format - Witten, Neal, Cleary: Arithmetic Coding for Data Compression.
82,199	PAPER2	- UNIX "troff" format - Witten: Computer (in)security.
513,216	PIC	- 1728 x 2376 bitmap image (MSB first): text in French and line diagrams.
39,611	PROGC	- Source code in C - UNIX compress v4.0.
71,646	PROGL	- Source code in Lisp - system software.
49,379	PROGP	- Source code in Pascal - program to evaluate PPM compression.
93,695	TRANS	- ASCII and control characters - transcript of a terminal session.

The struture of the corpus is shown in the diagram below. Each pixel represents a match between consecutive occurrences of a string. The color of the pixel represents the length of the match: black for 1 byte, red for 2, green for 4 and blue for 8. The horizontal axis represents the position of the second occurrence of the string. The vertical axis represents the distance back to the match on a logarithmic scale. (The image was generated by the [fv](#) program with labels added by hand).



String match structure of the Calgary corpus

The horizontal dark line at around 60-70 in BOOK1 is the result of linefeed characters repeating at this regular interval. Dark lines at 1280 and 5576 in GEO are due to repeated data in block headers. The dark bands at multiples of 4 are due to the 4 byte structure of the data. PIC has a dark band at 1 due to long runs of zero bytes, and lines at multiples of 216, which is the width of the image in bytes. The blue curves at the top of the image show matches between different text files, namely BIB, BOOK1, BOOK2, NEWS, PAPER1, PAPER2, PROGC, PROGL, PROGP and TRANS, all of which contain English text. Thus, compressing these files together can result in better compression because they contain mutual information. No such matches are seen between GEO, OBJ2, and PIC with other files. Thus, these files should be compressed separately.

Other structure can be seen. For example, there are dark bands in OBJ2 at 4, 6, 8, and higher even numbers due to the lengths of 68000 machine instructions. The dark band in BIB at around 150-200 represents the average length of a bibliographic entry. This knowledge is useful in designing compression algorithms.

2.1.1. BOOK1

A sample of text from BOOK1:

```
was receiving a percentage from the farmer till such
time as the advance should be cleared off Oak found+
that the value of stock, plant, and implements which
were really his own would be about sufficient to pay his
debts, leaving himself a free man with the clothes he
stood up in, and nothing more.
<C vi>
<P 88>
THE FAIR -- THE JOURNEY -- THE FIRE
TWO months passed away. We are brought on to a
day in February, on which was held the yearly statute
or hiring fair in the county-town of Casterbridge.
At one end of the street stood from two to three
hundred blithe and hearty labourers waiting upon Chance
-- all men of the stamp to whom labour suggests nothing
worse than a wrestle with gravitation, and pleasure
nothing better than a renunciation of the same among
these, carters and waggoners were distinguished by
having a piece of whip-cord twisted round their hats;
```

The tables below show then 10 most frequent n-grams (n byte sequences) in BOOK1 for n from 1 to 5, and the 10 most frequent words, bigrams, and trigrams. (2 or 3 word sequences). Words (but not n-grams) were counted as sequences of letters with other punctuation ignored.

Count n=1	Count n=2	Count n=3	Count n=4	Count n=5	Words	Bigrams	Trigrams
-----	-----	-----	-----	-----	-----	-----	-----
125551	20452 e	11229 th	7991 the	5869 the	7079 the	890 of the	127 I don t
72431 e	17470 he	9585 the	6366 the	3238 and	4071 and	648 in the	38 I can t
50027 t	17173 t	8989 he	3771 and	1918 , and	3760 of	401 to the	36 one of the
47836 a	15995 th	4728 ing	3647 and	1590 was	3701 a	297 and the	34 don t know
44795 o	13649 a	4666 and	3373 ing	1407 that	3565 to	292 to be	30 would have been
40919 n	13255 d	4564 nd	3179 of	1356 n the	2265 in	268 in a	30 the end of
37561 h	11153 in	4539 an	2966 to	1262 that	2217 I	254 of a	30 It was a
37007 i	10937 t	4501 ed	1942 , an	1056 d the	1966 was	231 don t	28 out of the
36788 s	10749 s	3936 to	1853 in	1050 with	1536 that	225 at the	26 to be a
32889 r	9974 er	3756 ng	1843 was	1030 of th	1430 her	201 from the	24 I won t

2.1.2. PIC

PIC is an uncompressed FAX scan of a page of a French textbook shown below. The image is represented by scanning in rows from the top left using a 1 for black or 0 for white. The bits are packed 8 to a byte (MSB first) with 216 bytes per row. The file has no header. It can be viewed in some imaging software as a **.pbm** file by appending the following two lines of text to the beginning.

```
P4
1728 2376
```

Cela est d'autant plus valable que $T\Delta f$ est plus grand. A cet égard la figure 2 représente la vraie courbe donnant $|\phi(f)|$ en fonction de f pour les valeurs numériques indiquées page précédente.




FIG. 2

Dans ce cas, le filtre adapté pourra être constitué, conformément à la figure 3, par la cascade :

— d'un filtre passe-bande de transfert unité pour $f_0 \leq f \leq f_0 + \Delta f$ et de transfert quasi nul pour $f < f_0$ et $f > f_0 + \Delta f$, filtre ne modifiant pas la phase des composants le traversant ;




FIG. 3

— filtre suivi d'une ligne à retard (LAR) dispersive ayant un temps de propagation de groupe T_R décroissant linéairement avec la fréquence f suivant l'expression :

$$T_R = T_0 + (f_0 - f) \frac{T}{\Delta f} \quad (\text{avec } T_0 > T)$$

(voir fig. 4),




FIG. 4

telle ligne à retard est donnée par :

$$\varphi = -2\pi \int_0^f T_R df$$

$$\varphi = -2\pi \left[T_0 + \frac{f_0 - f}{\Delta f} T \right] f + \pi \frac{T}{\Delta f} f^2$$

Et cette phase est bien l'opposé de $\phi(f)$,

à un déphasage constant près (sans importance) et à un retard T_0 près (inévitables).

Un signal utile $S(f)$ traversant un tel filtre adapté donne à la sortie (à un retard T_0 près et à un déphasage près de la porteuse) un signal dont la transformée de Fourier est réelle, constante entre f_0 et $f_0 + \Delta f$, et nulle de part et d'autre de f_0 et de $f_0 + \Delta f$, c'est-à-dire un signal de fréquence porteuse $f_0 + \Delta f/2$ et dont l'enveloppe a la forme indiquée à la figure 5, où l'on a représenté simultanément le signal $S(f)$ et le signal $S_r(f)$ correspondant obtenu à la sortie du filtre adapté. On comprend le nom de récepteur à compression d'impulsion donné à ce genre de filtre adapté : la « largeur » (à 3 dB) du signal comprimé étant égale à $1/\Delta f$, le rapport de compression est de

$$\text{est de } \frac{T}{1/\Delta f} = T\Delta f$$





FIG. 5

On saisit physiquement le phénomène de compression en réalisant que lorsque le signal $S(f)$ entre dans la ligne à retard (LAR) la fréquence qui entre la première à l'instant 0 est la fréquence basse f_0 , qui met un temps T_0 pour traverser. La fréquence f entre à l'instant $t = (f - f_0) \frac{T}{\Delta f}$ et elle met un temps


$$T_0 - (f - f_0) \frac{T}{\Delta f} \text{ pour traverser, ce qui la fait ressortir à l'instant } T_0 \text{ également. Ainsi donc, le signal } S(f)$$



Left: PIC (reduced in size). Right: closeup of PIC.

2.1.3. GEO

GEO contains seismic data in IBM 32 bit floating point format, which is now obsolete. Each number has a sign bit (1 if negative), a 7 bit exponent, and a 24 bit mantissa representing a fraction between 0 and 1. The number is decoded as $\text{sign} \times \text{mantissa} \times 16^{\text{exponent} - 64}$. The decoded data is graphed below. Each line represents one data sequence of either 1344 or 320 samples starting at the top. The file has a repeating structure with a 200 byte header (not shown) preceding each sequence.



GEO

The numeric values range from -147504 to 159280. Values have at most 16 bits of precision. Thus, the last 8 bits of all of the mantissa values are 0. The points are further rounded as follows: between -240 and 240, all values are rounded to a multiple of 1/64. Other values with magnitudes less than 960 are rounded to 1/16. Between 960 and 4048, values are rounded to multiples of 1/4. Between 4048 and 61440, values are rounded to integers. Above 61440, values are rounded to multiples of 16.

2.1.4. Results

Early tests sometimes used an 18 file version of the corpus that included 4 additional papers (PAPER3 through PAPER6). Programs were often ranked by measuring bits per character (bpc) on each file separately and reporting them individually or taking the average. Simply adding the compressed sizes is called a "weighted average" since it is weighted toward the larger files.

The Calgary corpus is no longer widely used due to its small size. However, it has been used since 1996 in an ongoing [compression challenge](#) run by Leonid A. Broukhis with small cash prizes. The best compression ratios established as of Feb. 2010 are as follows.

Table. Calgary Compression Challenge History

Size	Date	Name
759,881	Sep 1997	Malcolm Taylor
692,154	Aug 2001	Maxim Smirnov
680,558	Sep 2001	Maxim Smirnov

653,720	Nov 2002	Serge Voskoboynikov
645,667	Jan 2004	Matt Mahoney
637,116	Apr 2004	Alexander Rhatushnyak
608,980	Dec 2004	Alexander Rhatushnyak
603,416	Apr 2005	Przemyslaw Skibinski
596,314	Oct 2005	Alexander Rhatushnyak
593,620	Dec 2005	Alexander Rhatushnyak
589,863	May 2006	Alexander Rhatushnyak
580,170	Jul 2010	Alexander Rhatushnyak

The rules of the Calgary challenge specify that the compressed size include the size of the decompression program, either as a Windows or Linux executable file or as source code. This is to avoid programs that cheat by hiding information from the corpus in the decompression program. Furthermore, the program and compressed files must either be packed in an archive (in one of several specified formats), or else 4 bytes plus the length of each file name is added. This is to prevent cheating by hiding information in the file names and sizes. Without such precautions, programs like [barf](#) could claim to compress to zero bytes.

Submissions prior to 2004 are custom variants of compressors by the authors based on PPM algorithms (rk for Taylor, slim for Voskoboynikov, ppmn for Smirnov). Subsequent submissions are variants of the open source [paq6](#), a context mixing algorithm. The 2010 submission is based on paq8.

The following are the compressed sizes and times for some compressors on the file [calgary.tar](#), a [TAR](#) file containing a concatenation of the 14 files. Compression times (CT) and decompression times (DT) are process times (not wall times) in seconds and are measured on a 2 GHz T3200 under Windows Vista. Underlined values indicate the [Pareto frontier](#), meaning that no other program is both faster and higher ranked. When options are shown, they are usually set for maximum or nearly maximum compression at the expense of speed and memory (up to 1 GB). However, memory usage is not critical because of the small size of the test data. Some programs are tested again with default options. The decompression program size is not included. The year is for the version tested. Many programs have more than one author. The listed author is for the tested version.

calgary.tar	CT	DT	Program	Options	Algorithm	Year	Author
595,533	<u>411</u>	<u>401</u>	paq8l	-6	CM	2007	Matt Mahoney
598,983	462	463	paq8px_v67	-6	CM	2009	Jan Ondrus
605,656	<u>182</u>	<u>197</u>	paq8f	-6	CM	2005	Matt Mahoney
610,871	549	528	paqar 4.1	-6	CM	2006	Alexander Rhatushnyak
643,746	<u>59</u>	<u>60</u>	fp8	-6	CM	2010	Jan Ondrus
644,190	<u>19.9</u>	<u>19.9</u>	zpaq 1.10	ocmax.cfg	CM	2009	Matt Mahoney
647,110	140	138	paq6	-8	CM	2003	Matt Mahoney
647,440	<u>7.6</u>	<u>7.3</u>	nanozip 0.07a	-cc	CM	2009	Sami Runsas
666,216	9.1	9.4	durilca 0.5		PPM	2006	Dmitry Shkarin
669,497	8.5	8.1	ppmonstr J		PPM-12	2006	Dmitry Shkarin
674,145	22.6	22.8	slim 23d		PPM-32	2004	Serge Voskoboynikov
676,904	12.4	12.3	paq9a	-6	LZP+CM	2007	Matt Mahoney
681,419	10.9	10.2	xwrt 3.2	-l14	Dict+CM	2007	Przemyslaw Skibinski
682,512	8.7	9.3	lpaq1	6	CM	2007	Matt Mahoney
682,581	51.2	51.6	pimple2		CM	2006	Ilia Muraviev
686,139	9.5	8.8	lpaq9m	8	CM	2009	Alexander Rhatushnyak
694,050	9.0	8.6	epm r2		PPM	2003	Serge Osnach
696,850	43.9	44.5	ash 04a		CM	2003	Eugene D. Shelwien
699,191	7.6	7.4	zpaq 1.10	ocmid.cfg	CM	2009	Matt Mahoney
705,962	<u>6.9</u>	<u>7.0</u>	bit 0.7	-p=5	CM	2008	Osman Turan
706,911	<u>5.0</u>	<u>5.2</u>	cmm4		CM	2008	Christopher Mattern
716,731	18.7	18.7	paq1		CM	2001	Matt Mahoney
716,734	103	18.2	enc		CM	2003	Serge Osnach
720,980	8.7	8.5	tc 5.2dev2		CM	2007	Ilia Muraviev
725,423	5.0	<u>4.3</u>	ccmx	7	CM	2008	Christian Martelock
731,926	10.2	10.0	bee 0.79	-m3 -d8	PPM	2005	Andrew Filinsky, Melchiorre Caruso
733,864	5.0	<u>4.1</u>	uharc 0.6b	-mx -md32768	PPM	2005	Uwe Herklotz
738,016	<u>3.8</u>	<u>3.6</u>	ccm	7	CM	2008	Christian Martelock
741,419	64.4	65.5	ocamyd 1.66 test1	-m0 -s8 -f	DMC	2008	Frank Schwellinger
743,070	<u>1.7</u>	<u>1.4</u>	ctxf	-me	PPM	2003	Nikita Lesnikov
743,351	<u>1.4</u>	<u>0.53</u>	nanozip 0.07a		BWT	2009	Sami Runsas
745,162	662	448	bwmonstr 0.02		BWT	2009	Sami Runsas
754,037	1.8	1.8	ppmvc	-o16 -m256	ROLZ+PPM	2006	Przemyslaw Skibinski
754,243	1.4	1.4	ppmd J	-o16 -m256	PPM	2006	Dmitry Shkarin

769,400	8.5	0.62	mcomp_x32 2.00	-mf3x	ROLZ	2008	Malcolm Taylor
773,141	<u>0.93</u>	0.98	ppms J	-o6	PPM	2006	Dmitry Shkarin
774,397	1.1	0.62	bsc 1.0.3		BWT	2010	Ilya Grebnov
775,647	3.6	3.7	ppmx 0.5		PPM	2010	Ilia Muraviev
776,262	79.4	0.87	quark 0.95r beta	-m1 -d25 -l8	LZ77	2006	Frederic Bautista
776,273	9.6	9.7	boa 0.58b		PPM	1998	Ian Sutton
778,760	18.7	18.7	acb 2.00a		LZ77	1997	George Buyanovsky
780,549	0.93	0.71	bsc 1.0.3	-m2	ST5	2010	Ilya Grebnov
783,884	1.4	1.1	bcm 0.10-x86		BWT	2009	Ilia Muraviev
784,981	3.9	4.0	px 1.0		CM	2006	Ilia Muraviev
785,427	14.1	6.9	m03exp	32MB	BWT	2005	Michael Maniscalco, mij4x
785,541	2.3	1.8	m03 0.2 alpha	4000000	BWT	2009	Michael Maniscalco
787,007	1.6	0.72	sbc 0.970r2	-m3 -b63	BWT	2002	Sami Runsas
787,653	<u>0.73</u>	0.73	dark 0.51		BWT	2007	Malyshev Dmitry Alexandrovich
791,446	2.3	2.0	uhbc 1.0	-m3 -cf	BWT	2003	Uwe Herklotz
794,255	0.84	0.70	GRZipII 0.2.4		LZP+BWT	2004	Ilya Grebnov
798,705	10.3	3.8	bbb	cf	BWT	2006	Matt Mahoney
800,402	0.96	0.87	mcomp_x32 2.00	-mw	BWT	2008	Malcolm Taylor
804,992	<u>0.57</u>	0.67	ppmd J		PPM-4	2006	Dmitry Shkarin
806,208	1.1	0.57	bssc 0.95a	-b16383	BWT	2005	Sergeo Sizikov
812,123	3.1	2.8	lzipxj	9	LZP+PPM	2007	Ilia Muraviev, Jan Ondrus
816,982	3.1	3.3	mcomp_x32 2.00	-mc	DMC	2008	Malcolm Taylor
820,307	3.6	0.42	xwrt 3.2	-l6	Dict+LZMA	2007	Przemyslaw Skibinski
824,573	1.5	<u>0.16</u>	7zip 9.08a		LZMA	2009	Igor Pavlov
831,828	<u>0.36</u>	0.50	rings 1.6	10	BWT	2009	Nania Francesco Antonio
832,035	6.4	6.6	p12		NN	2000	Matt Mahoney
842,265	5.9	5.9	p6		NN	2000	Matt Mahoney
842,696	0.95	0.70	szip	-b41	ST6	1998	Michael Schindler
851,043	1.9	1.8	hook 1.4	512	LZP+DMC	2009	Nania Francesco Antonio
858,791	8.7	0.23	lzipm 0.15	9	ROLZ	2008	Ilia Muraviev
860,097	0.60	0.39	bzip2 1.0.3	-9	BWT	2005	Julian Seward
860,984	4.6	0.53	flashzip 0.99b8	-m3 -c7 -b8	ROLZ	2010	Nania Francesco Antonio
863,350	3.0	<u>0.06</u>	cabarc 1.00.0601	-m lzx:21	LZX	1997	Microsoft Corp.
864,499	17.0	4.2	qc	-8	LZP?	2006	Denis Kyznetsov
870,948	1.9	0.37	balz 1.15		LZ77	2008	Ilia Muraviev
872,028	5.2	0.31	lzturbo 0.9	-m59	LZ77	2007	Hamid Bouzidi
873,681	1.4	0.44	quad 1.12	-x	ROLZ	2007	Ilia Muraviev
876,234	8.8	7.3	ha 0.98	a2	PPM-4	1993	Harri Hirvola
885,273	1.3	1.3	tarsalzp		LZP	2007	Piotr Tarsa
885,954	6.9	2.2	qazar	-d9 -x7 -l7	LZP	2006	Denis Kyznetsov
900,048	2.6	2.9	dmc	100000000	DMC	1987	Gordon V. Cormack
905,889	1.2	0.17	csc31	-m3 -d7	LZ77	2009	Fu Siyuan
917,136	8.3	0.16	packet 0.91	-m6 -s9	LZ77	2009	Nania Francesco Antonio
935,598	43.5	43.6	fpaq2		PPM+CM	2006	Nania Francesco Antonio
948,568	0.67	0.75	winturtle 1.60	512MB	PPM	2008	Nania Francesco Antonio
975,889	35.8	0.09	kzip		LZ77	2006	Ken Silverman
979,349	0.37	0.40	sr2		SR	2007	Matt Mahoney
996,074	1.2	0.15	rar 2.50		LZ77	1999	Eugene Roshal
997,381	33.5	33.9	ctw		CTW	2002	Erik Franken, Marcel Peeters
1,001,364	<u>0.31</u>	0.12	thor	e5	LZ77	2007	Oscar Garcia
1,004,728	0.33	0.36	bzp		LZP	2008	Nania Francesco Antonio
1,011,786	0.47	<u>0.03</u>	tornado		LZ77	2008	Bulat Ziganshin
1,016,779	3.0	0.90	ha 0.98		LZ77	1993	Harri Hirvola
1,020,896	5.6	1.0	cm -m25659172	-t3152896 -o4	PPM-4	1997	Bulat Ziganshin
1,021,485	1.0	0.32	xwrt	-l3	Dict+LZ77	2007	Przemyslaw Skibinski
1,021,821	0.64	0.17	pkzip 2.04e	-ex	LZ77	1993	Phil Katz, PKWARE Inc.
1,022,414	0.09	0.11	slug 1.27b		ROLZ	2007	Christian Martelock
1,022,810	0.68	0.11	gzip 1.3.5	-9	LZ77	2002	Jean-loup Gailly, Mark Adler
1,023,037	0.68	0.10	zip 2.32	-9	LZ77	2005	Jean-loup Gailly, Mark Adler
1,029,644	0.34	0.11	zip 2.32		LZ77	2005	Jean-loup Gailly, Mark Adler
1,045,532	0.31	0.14	thor	e4	LZP	2007	Oscar Garcia
1,053,371	4.2	4.2	fpaq3		o3	2006	Nania Francesco Antonio
1,093,944	<u>0.12</u>	0.09	etincelle		ROLZ	2010	Yann Collet
1,093,958	3.6	5.1	lzgt3a		LZ77	2008	Gerald R. Tamayo
1,112,774	1.4	0.17	lzuf		LZ77	2009	Gerald R. Tamayo
1,131,100	23.1	24.7	lcssr 0.2		SR	2007	Frank Schwellinger
1,137,453	30.9	0.03	lzss 0.01	ex	LZSS	2008	Ilia Muraviev
1,140,859	2.8	0.46	ulz	c6	LZ77	2010	Ilia Muraviev
1,155,175	8.4	9.1	arbc2z		PPM-2	2006	David A. Scott
1,156,301	0.50	0.31	lzc 0.08		LZ77	2007	Nania Francesco Antonio
1,162,904	1.2	0.04	lzop 1.01		LZ77	2003	Markus F.X.J. Oberhumer

1,237,612	0.08	0.08	zhuff		LZ77	2009	Yann Collet
1,243,434	6.7	0.09	lzw		LZW	2008	Ilia Muraviev
1,267,256	0.20	0.20	srank	-C8 -03	SR	1997	Peter Fenwick
1,319,521	1.6	0.22	compress		LZW	1985	Spencer Thomas et. al.
1,331,783	1.0	1.0	fastari		o2	2006	Nania Francesco Antonio
1,337,403	0.14	0.03	quicklz 1.30		LZ77	2007	Lasse Mikkel Reinhold
1,340,008	0.16	0.09	brieflz		LZ77	2004	Joergen Ibsen
1,401,857	0.12	0.06	lzw3-a		LZ77	1991	Ross Williams
1,430,672	1.0	0.05	lzss	e	LZSS	2008	Ilia Muraviev
1,537,191	1.8	<u>0.02</u>	bpe	5000 4096 200 3	BPE	1994	Philip Gage
1,545,939	<u>0.08</u>	0.06	lzw3		LZ77	1991	Ross Williams
1,601,207	0.65	0.64	fpaq0f2		o0 (ind)	2008	Matt Mahoney
1,615,386	<u>0.06</u>	0.04	lzw2		LZ77	1991	Ross Williams
1,669,277	0.08	0.04	snappy 1.0.1		LZ77	2011	Google
1,670,642	0.50	0.06	lzw5		LZW	1991	Ross Williams
1,685,882	0.36	0.40	fpaq0p		o0 (adapt)	2007	Ilia Muraviev
1,691,924	0.21	0.14	flzp		LZP	2008	Matt Mahoney
1,725,464	0.06	0.06	lzw1-a		LZ77	1991	Ross Williams
1,737,702	<u>0.04</u>	0.03	lzw1		LZ77	1991	Ross Williams
1,884,160	<u>0.03</u>	0.03	lzt1/NTFS		LZSS	1995	Microsoft Corp.
1,903,024	0.70	0.76	fpaq0		o0 (stat)	2004	Matt Mahoney
1,938,635	0.06	0.03	lzp2		LZP	2009	Yann Collet
3,152,896			uncompressed				

Algorithms (when known) are as follows. They are described in more detail throughout the rest of this book. "CM" means "context mixing". "Dict" means dictionary encoding. "PPM-N" means order-N PPM. "oN" means order-N context modeling (stationary, adaptive, or indirect). "STn" means an order-n sort (Schindler) transform. "NN" means neural network. "SR" means symbol ranking.

2.2. Large Text Compression Benchmark


The [Large Text Compression Benchmark](#) consists of a single Unicode encoded XML file containing a dump of Wikipedia text from Mar. 3, 2006, truncated to 10^9 bytes (the file *enwik9*). Its stated goal is to encourage research into artificial intelligence, specifically, natural language processing. As of Feb. 2010, 128 different programs (889 including different versions and options) were evaluated for compressed size (including the decompression program source or executable and any other needed files as a zip archive), speed, and memory usage. The benchmark is open, meaning that anyone can submit results.

Programs are ranked by compressed size with options selecting maximum compression where applicable. The best result obtained is 127,784,888 bytes by Dmitry Shkarin for a customized version of durilca using 13 GB memory. It took 1398 seconds to compress and 1797 seconds to decompress using a size-optimized decompression program on a 3.8 GHz quad core Q9650 with 16 GB memory under 64 bit Windows XP Pro on July 21, 2009. The data was preprocessed with a custom dictionary built from the benchmark and encoded with order 40 PPM. durilca is a modified version of ppmonstr by the same author. ppmonstr is in turn a slower but better compressing ppmd program which is used for maximum compression in several archivers such as rar, WinZip, 7zip, and freearc.

By comparison, zip -9 compresses to 322,649,703 bytes in 104 seconds and decompresses in 35 seconds using 0.1 MB memory. It is ranked 92'nd.


The benchmark shows a 3 way trade off between compressed size, speed, and memory usage. The two graphs below show the Pareto frontier, those compressors for which no other compressors both compress smaller and faster (or smaller and use less memory). The graphs are from Aug. 2008, but the current data shows a similar trend. In particular, no single algorithm (shown in parenthesis) is the "best".

LTCB Size vs. Speed



Pareto frontier: compressed size vs. compression time as of Aug. 18, 2008 (options for maximum compression).

LTCB Size vs Memory




Pareto frontier: compressed size vs. memory as of Aug. 18, 2008 (options for maximum compression).

Note that speed tests may be run on different machines, and that only the options for maximum compression for each program are used. Nevertheless, the general trend remains valid. Individual compressors often have options that allow the user to make the same 3 way trade off.

The reason that the benchmark is so dependent on memory is that there are string matches that span the entire length of the 1 GB file. This can be seen in the analysis below. The blue area at the top right of the image represents matches between the beginning and end of the file. The dark band in the middle of the image is due

to several thousand Wikipedia articles for U.S. cities that were apparently generated from census data by a program. This region is highly compressible.



String repeat analysis of enwik9 with fv.

2.3. Hutter Prize

The [Hutter prize](#) is based on the first 10^8 bytes (the file *enwik8*) of the Large Text Compression benchmark with similar rules and goals. It is a contest in which prize money (500 euros per 1% gain) is awarded for improvements of 3% or more over the previous submission, subject to time and memory limits on the test computer. The best result is 15,949,688 bytes for an archive and a decompressor submitted by Alexander Rhatushnyak on May 23, 2009. It requires 7608 seconds and 936 MB memory to decompress on a 2 GHz dual core T3200 under 32 bit Windows Vista. The submission is based on two open source, context mixing programs paq8hp12 and lpaq9m with a custom dictionary for preprocessing.

By comparison, zip -9 compresses the same data to 36,445,373 bytes and uncompresses in 3.5 seconds using 0.1 MB memory.

2.4. Maximum Compression

The [maximum compression benchmark](#) has two parts: a set of 10 public files totaling 53 MB, and a private collection of 510 files totaling 301 MB. In the public data set (SFC or single file compression), each file is compressed separately and the sizes added. Programs are ranked by size only, with options set for best compression individually for each file. The set consists of the following 10 files:

842,468 a10.jpg	- a high quality 1152 x 864 baseline JPEG image of a fighter jet.
3,870,784 acro32.exe	- x86 executable code - Acrobat Reader 5.0.
4,067,439 english.dic	- an alphabetically sorted list of 354,941 English words.
4,526,946 FlashMX.pdf	- PDF file with embedded JPEG and zipped BMP images.
20,617,071 fp.log	- web server log, ASCII text.
3,782,416 mso97.dll	- x86 executable code from Microsoft Office.
4,168,192 ohs.doc	- Word document with embedded JPEG images.
4,149,414 rafale.bmp	- 1356 x 1020 16 bit color image in 24 bit RGB format.
4,121,418 vcfiu.hlp	- OCX help file - binary data with embedded text.
2,988,578 world95.txt	- ASCII text - 1995 CIA World Factbook.

The top ranked program as of Dec. 31, 2009 with a total size of 8,813,124 bytes is paq8px, a context mixing algorithm with specialized models for JPEG images, BMP images, x86 code, text, and structured binary data. WinRK 3.1.2, another context mixing algorithm, is top ranked on 4 of the files (txt, exe, dll, pdf). WinRK uses a dictionary which is not included in the total size. 208 programs are ranked. zip 2.2 is ranked 163 with a size of 14,948,761.

In the second benchmark or MFC (multiple file compression), programs are ranked by size, compression speed, decompression speed, and by a formula that combines size and speed with time scaled logarithmically. The

data is not available for download. Files are compressed together to a single archive. If a compressor cannot create archives, then the files are collected into an uncompressed archive (TAR or QFC), which is compressed.

In the MFC test, paq8px is top ranked by size. freearc is top ranked by combined score, followed by nanozip, winrar, and 7zip. All are archivers that detect file type and apply different algorithms depending on type.

2.5. Generic Compression Benchmark

The [Generic Compression Benchmark](#) has the goal of evaluating compression algorithms in the context of universal prediction or intelligence, as defined by Legg and Hutter (2006). By this definition, data sources are assumed to have a universal Solomonoff distribution, i.e. generated by random programs with a preference for smaller or simpler programs. The evidence for such a distribution is the success of applying Occam's Razor to machine learning and to science in general: the simplest theories that fit the observed data tend to be the best predictors. The purpose of the test is to find good compression algorithms that are not tuned to specific file types.

The benchmark does not publish any test data. Rather, it publishes a program to generate the data from a secret seed or an internally hardware generated random number. The data consists of the bit string outputs of one million random Turing machines, truncated to 256 bits and packed into null terminated byte strings. The average output size is about 6.5 MB. The test allows public verification while eliminating the need to measure the decompressor size because it is not possible to hide the test data in the decompressor without knowing the cryptographic random number seed. The test produces repeatable results with about 0.05% accuracy. Programs are ranked by the ratio of compressed output to the compressed output of a reference compressor (ppmonstr) to improve repeatability.

Unfortunately the benchmark fails to completely eliminate the problem of tuning compressors to public benchmarks. The top ranked program is a stationary context mixing model configuration implemented in [zpaq](#) using a preprocessor by J. Ondrus that splits each string into an incompressible prefix and a bit repeat instruction. Its score is 0.8750, compared to 1.3124 for zip -9. Generally, however, the rank order of compressors is similar to that of other benchmarks.

2.6. Compression Ratings

[Compression Ratings](#) by Sami Runsas ranks programs on 5.4 GB of various data types from public sources using a score that combines size and speed. The data includes English text, Windows executable code, RGB and grayscale images, CD quality audio, and a mix of data types from two video games. Some of the 10 data sets contains multiple files. Single file compressors are tested on an equivalent [tar](#) file. Compressed sizes include the size of the decompression program source code or executable program compressed with 7zip. Run times are measured from RAM-disk to avoid I/O delays.

Programs must pass a qualifying round with minimum compression ratio and time requirements on a small data set. The benchmark includes a calculator that allows the user to rank compressors using different weightings for the importance of size, compression speed, and decompression speed. The default scale is "1/10 lower ratio or twice the total time is worth half the rating". This is effectively the same formula used by maximumcompression.com. As of June 2012, 130 programs were tested with 682 combinations of options. The top ranked programs for the default settings were nanozip followed by freearc, CCM, flashzip, and 7-zip. (Runsas is the author of nanozip.) There are additional benchmarks that compare BWT based compressors and bzip2-compatible compressors, and additional benchmarks for some special file types.

2.7. Other Benchmarks

Some other benchmarks are mentioned briefly.

[Squeeze Chart](#) by Stephen Busch, ranks programs on 6.4 GB of mostly private data of various types by size only. The top ranked is paq8px_v67 as of Dec. 28, 2009.

[Monster of Compression](#) by Nania Francesco Antonio, ranks programs by size on 1,061,420,156 bytes of mostly public data of various types with a 40 minute time limit. There are separate tests for single file

compressors and archivers. As of Dec. 20, 2009 the top ranked archiver is nanzip 0.7a and the top ranked file compressor is ccmx 1.30c. Both use context mixing.

[UCLC](#) by Johan de Bock contains several benchmarks of public data for compressors with a command line interface (which is most of them). As of Feb. 2009, paq8i or paq8p was top ranked by size on most of them.

[Metacompressor](#) is an automated benchmark that allows developers to submit programs and test files and compare size (excluding the decompressor), compression time, and decompression time.

[Xtreme Compression](#) compares 60 compression/option combinations with their own product for size and speed on a 80 GB synthetic database file. After their product, which works only on database tables, nanzip -nm -cc ranks next, followed by zpaq (mid.cfg), as of Nov. 2011.

2.8. File System Benchmarks

Meyer and Bolosky (2011), in a study of practical deduplication, examined the contents of 857 Windows file systems among Microsoft employees from a wide range of departments in 2009. The table below gives the approximate distribution by type, as compared to similar studies by them in 2000 and 2004.

2000	2004	2009	Data type
2%	3%	13%	No filename extension
13%	10%	10%	.dll (x86 dynamic library)
3%	4%	8%	.lib (x86 static library)
	5%	7%	.vhd (virtual hard drive)
7%	5%	7%	.pdb (database)
6%	4%	4%	.exe (x86 executable)
4%		3%	.pch (precompiled C/C++ header file)
3%	4%	2%	.cab (compressed archive)
	4%	1%	.wma (compressed audio)
		1%	.iso (CD/DVD image)
5%	4%		.pst (email application database)
4%	3%		.mp3 (compressed audio)
3%			.chm (compressed HTML help file)
49%	55%	43%	Other

The average file system in 2009 had a capacity of 150 GB and was 40% full. The fraction in use remained nearly constant since 2000 in spite of a doubling of disk capacity almost every year. Average file size also remained nearly constant at about 4 KB since 1981. However the number of files increased, to an average of 225,000 files in 36,000 directories in 2009. File size varies widely. About half of all data is in files larger than 25 MB. One fourth is in files smaller than 3 MB and one fourth in files larger than 1 GB. These percentiles are also increasing as the number of larger files grow. Half of all data was in files larger than 1 MB in 2000 and 5 MB in 2004.

An average file system in 2009 could be compressed by 22% by whole file deduplication, i.e. reduced to 78% of original size by replacing identical copies of files with links. Compression could be increased to 28% by dividing files into 64 KB blocks and removing duplicate blocks, or to 31% using 8 KB blocks. If alignment restrictions are removed, then compression increases to 39% by replacing duplicate segments larger than 64 KB, or 42% using 8 KB segments. If links are allowed to point to duplicates on other file systems, then compression improves by about 3 percentage points for each doubling of the number of computers. For the entire 857 file systems, whole file deduplication compressed by 50%, and 8K segments by 69%.

The obvious application of deduplication is incremental backups. It is only necessary to back up data that has changed. In the 2009 study, the median time since last modification was 100 days. 7% of files of files were modified within one week, and 75% within one year.

3. Coding

A code is an assignment of bit strings to symbols such that the strings can be decoded unambiguously to recover the original data. The optimal code for a symbol with probability p will have a length of $\log_2 1/p$ bits. Several efficient coding algorithms are known.

3.1. Huffman Coding

Huffman (1952) developed an [algorithm](#) that calculates an optimal assignment over an alphabet of n symbols in $O(n \log n)$ time. deflate (zip) and bzip2 use Huffman codes. However, Huffman codes are inefficient in practice because code lengths must be rounded to a whole number of bits. If a symbol probability is not a power of $1/2$, then the code assignment is less than optimal. This coding inefficiency can be reduced by assigning probabilities to longer groups of symbols but only at the cost of an exponential increase in alphabet size, and thus in run time.

The algorithm is as follows. We are given an alphabet and a probability for each symbol. We construct a binary tree by starting with each symbol in its own tree and joining the two trees that have the two smallest probabilities until we have one tree. Then the number of bits in each Huffman code is the depth of that symbol in the tree, and its code is a description of its path from the root (0 = left, 1 = right). For example, suppose that we are given the alphabet $\{0,1,2,3,4,5,6,7,8,9\}$ with each symbol having probability 0.1. We start with each symbol in a one-node tree:

```

.1 .1 .1 .1 .1 .1 .1 .1 .1 .1
0 1 2 3 4 5 6 7 8 9

```

Because each small tree has the same probability, we pick any two and combine them:

```

      .2
     / \
    .1 .1
   /  \
  0    1
.1 .1 .1 .1 .1 .1 .1 .1 .1 .1
2 3 4 5 6 7 8 9

```

Continuing,

```

      .2      .2      .2      .2
     / \     / \     / \     / \
    .1 .1   .1 .1   .1 .1   .1 .1
   /  \   /  \   /  \   /  \
  0    1  2    3  4    5  6    7
.1 .1 .1 .1 .1 .1 .1 .1 .1 .1
8 9

```

At this point, 8 and 9 have the two lowest probabilities so we have to choose those:

```

      .2      .2      .2      .2      .2
     / \     / \     / \     / \     / \
    .1 .1   .1 .1   .1 .1   .1 .1   .1 .1
   /  \   /  \   /  \   /  \   /  \
  0    1  2    3  4    5  6    7    8 9

```

Now all of the trees have probability .2 so we choose any pair of them:

```

      .4
     / \
    .2 .2
   /  \ /  \
  .1 .1 .1 .1
 /  \ /  \
0    1 2    3
.1 .1 .1 .1 .1 .1 .1 .1 .1 .1
4 5 6 7 8 9

```

We choose any two of the three remaining trees with probability .2:

```

      .4      .4      .2
     / \     / \     / \
    .2 .2   .2 .2   .1 .1
   /  \ /  \ /  \ /  \
  .1 .1 .1 .1 .1 .1 .1 .1
 /  \ /  \ /  \
0    1 2    3 4    5 6    7
.1 .1 .1 .1 .1 .1 .1 .1 .1 .1
8 9


```

Now the two smallest probabilities are .2 and one of the .4:


```

      .6
     / \

```



Now the two smallest are .4 and .6. After this step, the tree is finished. We can label the branches 0 for left and 1 for right, although the choice is arbitrary.



From this tree we construct the code:

Symbol	Code
0	000
1	001
2	010
3	011
4	1000
5	1001
6	1010
7	1011
8	110
9	111

A code may be static or dynamic. A static code is computed by the compressor and transmitted to the decompressor as part of the compressed data. A dynamic code is computed by the compressor and periodically updated, but not transmitted. Instead, the decompressor reconstructs the code using exactly the same algorithm using the previously decoded data to estimate the probabilities. Neither method compresses better because any space saved by not transmitting the model is paid back by having less data with which to estimate probabilities.

Huffman codes are typically static, mainly for speed. The compressor only needs to compute the code once, using the entire input to compute probabilities. To transmit a Huffman table, it is only necessary to send the size of each symbol, for example: (3,3,3,3,4,4,4,4,3,3). Both the compressor and decompressor would then assign codes by starting with the shortest symbols, counting up from 0, and appending a 0 bit whenever the code gets longer. This would result in the following different but equally effective code:

Symbol	Size	Code
0	3	000
1	3	001
2	3	010

3	3	011
8	3	100
9	3	101
4	4	1100
5	4	1101
6	4	1110
7	4	1111

For file compression, Huffman coded data still needs to be packed into bytes. JPEG packs bits in MSB (most significant bit) to LSB (least significant bit) order. For example, the codes 00001 00111 would be packed as 00001001 11..... . The [deflate](#) format used in zip, gzip, and png files packs bits in LSB to MSB order, as if each byte is written backward, i.e. 1001000011 .

One other complication is that the last byte has to be padded in such a way that it is not interpreted as a Huffman code. JPEG does this by not assigning any symbol to a code of all 1 bits, and then padding the last byte with 1 bits. Deflate handles this by reserving a code to indicate the end of the data. This tells the decoder not to decode the remaining bits of the last byte.

3.2. Arithmetic Coding

Huffman coding has the drawback that code lengths must be a whole number of bits. This effectively constrains the model to probabilities that are powers of 1/2. The size penalty for modeling errors is roughly proportional to the square of the error. For example, a 10% error results in a 1% size penalty. The penalty can be large for small codes. For example, the only possible ways to Huffman code a binary alphabet is to code each bit as itself (or its opposite), resulting in no compression.

[Arithmetic coding](#) (Rissanen, 1976), also called range coding, does not suffer from this difficulty. Let P be a model, meaning that for any string x , $P(x)$ is the probability of that string. Let $P(< x)$ be the sum of the probabilities of all strings lexicographically less than x . Let $P(\leq x) = P(< x) + P(x)$. Then the arithmetic code for a string x is the shortest binary number y such that $P(< x) \leq y < P(\leq x)$. Such a number can always be found that is no more than 1 bit longer than the Shannon limit $\log_2 1/P(x)$.

An arithmetic code can be computed efficiently by expressing $P(x)$ as a product of successive symbol predictions by the chain rule, $P(x) = \prod_i P(x_i | x_1 x_2 \dots x_{i-1})$ where x_i means the i 'th symbol (bit or character) in x . Then the arithmetic code can be computed by updating a range [low, high) (initially [0, 1)) for each symbol by dividing the range in proportion to the probability distribution for that symbol. Then the portion of the range corresponding to the symbol to be coded is used to update the range. As the range shrinks, the leading bits of low and high match and can be output immediately because the code y must be between them. The decompressor is able to decode y by making an identical sequence of predictions and range reductions.

Most modern data compressors use arithmetic coding. Early compressors used Huffman coding because arithmetic coding was [patented](#) and because its implementation required multiplication operations, which was slow on older processors. Neither of these issues are relevant today because the important patents have expired and newer processors have fast multiply instructions (faster than memory access).

The most common arithmetic coders code one byte at a time (PPM) or one bit at a time (CM, DMC). Free source code with no licensing restrictions for a bitwise encoder can be found in the source code for [ppmd](#) (D. Shkarin). Bitwise coders licensed under GPL can be found in the source code for [PAQ](#) based compressors including FPAQ, LPAQ, and ZPAQ, the BWT compressor BBB, and the symbol ranking compressor SR2. The simplest of these is the order 0 coder [fpq0](#).

The following is the arithmetic coder from [zpaq 1.10](#). It encodes bit y (0 or 1) with probability $p = P(1) * 65536$ (a scaled 16 bit number) and codes to FILE* out. The encoder range is represented by two 32 bit integers (unsigned int) low and high, which are initially 1 and 0xffffffff respectively. After the range is split, a 1 is coded in the lower part of the range and a 0 in the upper part. After the range is split and reduced, it is normalized by shifting out any leading bytes that are identical between low and high. The low bits shifted in are all 0 bits for low and all 1 bits for high.

```
// Initial state
unsigned int low = 1, high = 0xffffffff;
```

```
// Encode bit y with probability p/65536
inline void Encoder::encode(int y, int p) {
    assert(out); // output file
    assert(p>=0 && p<65536);
    assert(y==0 || y==1);
    assert(high>low && low>0);
    unsigned int mid=low+((high-low)>>16)*p+((((high-low)&0xffff)*p)>>16); // split range
    assert(high>mid && mid>=low);
    if (y) high=mid; else low=mid+1; // pick half
    while ((high^low)<0x1000000) { // write identical leading bytes
        putc(high>>24, out); // same as low>>24
        high=high<<8|255;
        low=low<<8;
        low+=(low==0); // so we don't code 4 0 bytes in a row
    }
}
```

The range split is written to avoid 32 bit arithmetic overflow. It is equivalent to:

```
unsigned int mid=low+((unsigned long long)(high-low)*p>>16);
```

where (unsigned long long) is a 64 bit unsigned type, which not all compilers support. The initialization of low to 1 instead of 0 and the statement

```
low+=(low==0);
```

discard a tiny bit of the range to avoid writing 4 consecutive 0 bytes, which the ZPAQ compressor uses to mark the end of the encoded data so it can be found quickly without decoding. This is not a requirement in general. It is not used in the rest of the PAQ series. The decoder looks like this:

```
// Initial state
unsigned int low=1, high=0xffffffff, curr;
for (int i=0; i<4; ++i)
    curr=curr<<8|getc(in); // first 4 bytes of input

// Return decoded bit from file 'in' with probability p/65536
inline int Decoder::decode(int p) {
    assert(p>=0 && p<65536);
    assert(high>low && low>0);
    if (curr<low || curr>high) error("archive corrupted");
    assert(curr>=low && curr<=high);
    unsigned int mid=low+((high-low)>>16)*p+((((high-low)&0xffff)*p)>>16); // split range
    assert(high>mid && mid>=low);
    int y=curr<=mid;
    if (y) high=mid; else low=mid+1; // pick half
    while ((high^low)<0x1000000) { // shift out identical leading bytes
        high=high<<8|255;
        low=low<<8;
        low+=(low==0);
        int c=getc(in);
        if (c==EOF) error("unexpected end of file");
        curr=curr<<8|c;
    }
    return y;
}
```

The decoder receives as input p, the 16 bit probability that the next bit is a 1, and returns the decoded bit. The decoder has one additional variable in its state, the 32 bit integer curr, which is initialized to the first 4 bytes of compressed data. Each time the range is rescaled, another byte is read from FILE* in. high and low are initialized as in the encoder.

One additional detail is how to handle the end of file. Most of the PAQ series compressors encode the file size separately and perform 8 encoding operations per byte. After the last encoding operation, the 4 bytes of either high or low or some value in between must be flushed to the archive because the decoder will read these 4 bytes in.

ZPAQ encodes 9 bits per byte, using a leading 1 bit modeled with probability $p = 0$ to mark the end of file. The effect on the encoder is to set `mid = low` and cause 4 bytes to be flushed. After that, 4 zero bytes are written to mark the end of the compressed data. When the end of file bit is decoded, the decoder reads these 4 zero bytes into `curr`, resulting in `low = 1`, `curr = 0`, `high = 0xffffffff`. Any further decoding would result in an error because the condition $low \leq curr \leq high$ fails.

The assertions are not necessary to make the code work. Rather, they are useful because they make the programmer's assumptions explicit in the form of self documenting run time tests. The code is compiled with `-DNDEBUG` to remove them after debugging.

3.3. Asymmetric Binary Coding

Most high end compressors use arithmetic coding. However, another possibility with the same theoretical coding and time efficiency for bit strings is asymmetric binary coding or ABC (Duda, 2007). An asymmetric coder has a single n -bit integer state variable x , as opposed to two variables (`low` and `high`) in an arithmetic coder. This allows a lookup table implementation. A bit y (0 or 1) with probability $p = P(y = 1)$ ($0 < p < 1$, a multiple of 2^{-n}) is coded in state x , initially 2^n :

```
if y = 0 then x := ceil((x+1)/(1-p)) - 1
if y = 1 then x := floor(x/p)
```

To decode, given x and p

```
y = ceil((x+1)*p) - ceil(x*p) (0 if fract(x*p) < 1-p, else 1)
if y = 0 then x := x - ceil(x*p)
if y = 1 then x := ceil(x*p)
```

x is maintained in the range 2^n to $2^{n+1} - 1$ by writing the low bits of x prior to encoding y and reading into the low bits of x after decoding. Because compression and decompression are reverse operations of each other, they must be performed in reverse order by storing the predictions and coded bits in a stack in either the compressor or the decompressor.

The coder is implemented in the order-0 compressors `fpaqa`, `fpaqb`, and `fpaqc` and the context mixing compressor `lpaq1a` from the [PAQ series](#). `fpaqa` uses lookup tables for both compression and decompression. It uses a 10 bit state and the probability is quantized to 7 bits on a nonlinear scale (finer near 0 and 1 and coarse near 1/2). The stack size is 500,000. Increasing these numbers would result in better compression at the expense of a larger table. `fpaqb` uses direct calculations except for division, which uses a table of inverses because division is slow. `fpaqc` is `fpaqb` with some speed optimizations. The coder for `fpaqc` is used in `lpaq1a` (a context mixing program), replacing the arithmetic coder in `lpaq1`.

Although a lookup table implementation might be faster on some small processors, it is slower on modern computers because multiplication is faster than random memory access. Some benchmarks are shown for `enwik8` (100 MB text) on a 2.0 GHz T3200 processor running on one of two cores. Ratio is fraction of original size. Compression and decompression times are nanoseconds per byte.

	Ratio	Comp	Decomp	Coder
	-----	----	----	----
<code>fpaqa</code>	.61310	247	238	ABC lookup table
<code>fpaqb</code>	.61270	244	197	ABC direct calculation
<code>fpaqc</code>	.61270	246	173	ABC direct calculation
<code>fpaqa -DARITH</code>	.61280	130	112	arithmetic (fpaqa compiled with -DARITH)

For high end compressors, CPU time and memory are dominated by the model, so the choice of coder makes little difference. `lpaq1` is a context mixing compressor, a predecessor of `lpaq9m`, ranked third of 127 on the large text benchmark as of Feb. 2010. `lpaq1a` is the same except that the arithmetic coder was replaced by the asymmetric binary coder from `fpaqb`. (Timed on a 2.188 GHz Athlon-64 3500+).

	Ratio	Comp	Decomp	Coder
	-----	----	----	----
<code>lpaq1a</code>	.19755	3462	3423	ABC direct calculation (fpaqb)
<code>lpaq1</code>	.19759	3646	3594	arithmetic

The arithmetic coder in lpaq1 and fpaqa -DARITH compresses slightly worse than the ABC coder because it uses 12 bits of precision for the probability, rather than 16 bits as in ZPAQ.

3.4. Numeric Codes

Numeric codes are Huffman codes for numbers with some common probability distributions, such as might appear in coding prediction errors in images or audio, or run lengths in highly redundant data. Generally such codes have simpler descriptions to transmit to the decoder than a full Huffman code length table.

3.4.1. Unary Codes

A [unary code](#) for a non negative integer n is n ones followed by a zero (or n zeros followed by a one). For example, 5 would be coded as 111110. A unary code implies a probability $P(n) = 1/2^{n+1} = 1/2, 1/4, 1/8, 1/16$, etc.

Signed integers can be encoded by mapping positive n to $2n$ and negative n to $-2n-1$ to form the sequence 0, -1, 1, -2, 2, -3, 3, etc.

3.4.2. Rice and Golomb Codes

A [Golomb code](#) for non negative integer n with parameter M divides n into a quotient $q = \text{floor}(n/M)$ and a remainder $r = n - Mq$. Then q is encoded in unary and r in binary. If M is a power of 2, then the code is called a "Rice code", and all of the remainders are coded in $\log_2 M$ bits. Otherwise, the smaller remainders are coded using $\text{floor}(\log_2 M)$ bits and the larger ones with one more bit. A Golomb code with $M = 1$ is just a unary code. For example:

n	M=1 (unary) q r code	M=2 (Rice) q r code	M=3 (Golomb) q r code	M=4 (Rice) q r code
0	0 0 0	0 0 00	0 0 00	0 0 000
1	1 0 10	0 1 01	0 1 010	0 1 001
2	2 0 110	1 0 100	0 2 011	0 2 010
3	3 0 1110	1 1 101	1 0 100	0 3 011
4	4 0 11110	2 0 1100	1 1 1010	1 0 1000
5	5 0 111110	2 1 1101	1 2 1011	1 1 1001
6	6 0 1111110	3 0 11100	2 0 1100	1 2 1010

A Golomb code with parameter M increases in length once every M values. Thus, it implies an approximate geometric distribution where n has probability proportional to $2^{-n/M}$.

3.4.3. Extra Bit Codes

An "extra bit" code uses a Huffman code to encode a range of numbers, then encodes a value within that range (the extra bits) in binary. This method can be used to approximate any distribution that is approximately smoothly varying.

[JPEG](#) uses extra bit codes to encode discrete cosine transform coefficients. The ranges are 0, 1, 2..3, 4..7, 8..15, increasing by powers of 2. It also uses a sign bit for all of the ranges except 0. The [deflate](#) format used in zip and gzip uses extra bit codes to encode match lengths and offsets using a more complex set of ranges.

A floating point number is a special case of an extra-bit code where the reals are approximated by mapping to integers on a logarithmic scale, and where each range is the same size and equally probable. A IEEE 64 bit floating point number (type double in C/C++) has one sign bit s representing -1 or 1, an 11 bit exponent e representing a number in the range -1022 to +1023, and a 52 bit mantissa m representing a number in the range $1/2$ to 1 (or 0 to 1 for the smallest exponent). It encodes the real number $sm2^e$. A [μ-law](#) or mu-law code, used to encode sampled telephone speech in North America and Japan, is an 8 bit floating point number with one sign bit, 3 exponent bits and 4 mantissa bits.

3.5. Archive Formats

Compressed files are stored in archives. An archive may contain a single file or multiple files. For example, [gzip](#) compresses a file, adds a **.gz** extension to the name and deletes the original. Decompression restores the original and deletes the compressed file. [zip](#) stores multiple files in an archive with a **.zip** file name extension. It has commands to list, add, update, delete, and extract files from the archive (without deleting the input files). However, both programs use the same compression algorithm.

Compression can sometimes be improved by compressing similar files together to take advantage of mutual information. Some archivers support this using a "solid" format, which requires that all of the files be compressed or extracted at the same time. [zip](#) is optimized for speed, so its archives are not solid. [PAQ](#) is optimized for maximum compression, so archives are always solid. It is optional for some archivers such as [WinRAR](#).

A file compressor such as [gzip](#) can be used to create solid archives by first collecting the input files together using a non-compressing archiver such as [tar](#) and then compressing. The combination of [tar](#) + [gzip](#) is commonly given a file name extension of **.tar.gz** or **.tgz**.

The [tar](#) format was originally designed for UNIX or Linux backups on tape. Thus, it is possible to make a [tar](#) file of an entire file system and restore its structure unchanged. The [tar](#) program appends a 512 byte header to each file containing the file name, path, size, timestamp, owner, group, permissions, and padding with zero bytes. Then the files are all appended together. The files are also padded with zero bytes to a multiple of 512. These zero bytes are easily compressed.

Archivers [zip](#), [WinRAR](#), and [7-zip](#) have features similar to [tar](#) in that they preserve timestamps and can compress and extract directory trees. However they do not store owner, group, and permission information in order to be compatible with Windows, which has a different security model.

3.5.1. Error Detection

Archivers sometimes add a checksum to detect transmission errors. The checksum is computed and stored before compression and verified after decompression. The need for a checksum is questionable because networks and storage media typically use error correction so that errors are very unlikely. Decompression errors are more likely to result from software bugs or deliberate manipulation. A single bit change in an archive can produce a completely different file. A checksum increases the archive size slightly (usually by 4 bytes) and also increases compression and decompression time by the time it takes to compute it. Some common checksum functions:

3.5.1.1. Parity

The checksum is computed by counting the number of 1 bits in the input stream. The checksum is 1 if the total is odd or 0 if even. All single bit errors are detected, but other errors are detected only with probability 1/2.

3.5.1.2. CRC-32

CRC-32 is a commonly used [cyclic redundancy check](#). It detects all burst errors up to 31 bits and all other errors with probability $1 - 2^{-32}$.

The input is a stream of bits which are shifted through a 32 bit window. Each time a 1 bit is shifted out, the new window contents is XORed with 0x04C11DB7, otherwise it is unchanged. The final contents of the window is the 32 bit checksum. A fast implementation will shift and XOR one byte at a time using a 256 by 32 bit lookup table.

3.5.1.3. Adler-32

The [Adler-32](#) checksum is used in the [zlib](#) implementation of deflate, which is used in [zip](#) and [gzip](#). It is a 32 bit checksum which detects errors with probability of $1 - 65521^{-2}$, which is almost $1 - 2^{-32}$.

The checksum has two 16-bit parts. The first part is 1 plus the sum of all input bytes, modulo 65521. The second part is the sum of all the intermediate sums of the first part, modulo 65521. The computation is very fast

because the (slow) mod operation only needs to be calculated infrequently using 32 bit arithmetic. 65521 is the largest prime number less than 2^{16} .

3.5.1.4. Cryptographic Hash Functions

CRC-32 and Adler-32 detect most accidental errors but it is easy to deliberately alter the input without changing the checksum. [Cryptographic hash functions](#) are designed to resist even deliberate attempts to find collisions (two different inputs that produce the same checksum). They are designed so that the best possible algorithms are no better than randomly guessing the inputs and testing the outputs. The disadvantage is that they are larger (usually 20 to 32 bytes) and take longer to compute (although still reasonably fast).

[ZPAQ](#) optionally adds a 160 bit (20 byte) [SHA-1](#) checksum. It detects errors with probability $1 - 2^{-160}$, even if deliberately introduced. The checksum calculation adds 5.6 ns per byte on a 2 GHz T3200, compared to 5.0 ns for CRC-32 and 5.3 ns for Adler-32 using the [SlavaSoft fsุม](#) v2.51 implementation and testing on a 1 GB text file.

3.5.2. Encryption

Some archives will encrypt data as well as compress it. I do not recommend adding this feature unless you know what you are doing. It is tempting to design your own algorithm, perhaps combining it with compression by altering the initial state of the model in a way that depends on a password. Such an attempt will almost certainly fail. Just because the output looks random doesn't mean it is secure. There is no test for randomness.

If you add encryption, use a well tested algorithm such as [AES](#). Compress your data before encryption, because encrypted data cannot be compressed. Use published encryption code. Don't write it yourself. Then publish *all* your source code and have security experts look at it. Others will try to break your code, and you should make their job as easy as possible by clearly documenting what your code does and how it works. Just because you use a secure, well tested algorithm doesn't mean the rest of your program is secure. For example, if your archiver deletes the plaintext, it might still be recovered from unallocated disk sectors or the swap file. Likewise for the password or password derived data in memory that gets swapped to disk. There are many subtle things that can go wrong.

Do not believe that a secret algorithm or an executable only release is more secure. An encryption algorithm is not considered secure unless it can withstand chosen plaintext attacks by adversaries with full knowledge of the algorithm. Most security experts will not even waste their time trying to attack a closed system. They will just use other, well tested solutions. If your system ever does become popular, then others *will* reverse engineer it and they *will* break it.

4. Modeling

A model is an estimate of the probability distribution of inputs to a compressor. Usually this is expressed as a sequence of predictions of successive symbols (bits, bytes, or words) in the input sequence given the previous input as context. Once we have a model, coding is a solved problem. But (as proved by Kolmogorov) there is no algorithm for determining the best model. This is the hard problem in data compression.

A model can be static or adaptive. In the static case, the compressor analyzes the input, computes the probability distribution for its symbols, and transmits this data to the decompressor followed by the coded symbols. Both the compressor and decompressor select codes of the appropriate lengths using identical algorithms. This method is often used with Huffman coding.

Typically the best compressors use dynamic models and arithmetic coding. The compressor uses past input to estimate a probability distribution (prediction) for the next symbol without looking at it. Then it passes the prediction and symbol to the arithmetic coder, and finally updates the model with the symbol it just coded. The decompressor makes an identical prediction using the data it has already decoded, decodes the symbol, then updates its model with the decoded output symbol. The model is unaware of whether it is compressing or decompressing. This is the technique we will use in the rest of this chapter.

4.1. Fixed Order Models

The simplest model is a fixed order model. An *order n* model inputs the last *n* bytes or symbols (the context) into a table and outputs a probability distribution for the next symbol. In the update phase, the predicted symbol is revealed and the table is updated to increase its probability when the same context next appears. An order 0 model uses no context.

4.1.1. Bytewise Encoding

A probability distribution is typically computed by using a counter for each symbol in the alphabet. If the symbols are bytes, then the size of the alphabet is 256. The prediction for each symbol is the count for that symbol divided by the total count. The update procedure is to increment the count for that symbol. If the arithmetic coder codes one byte at a time, then you pass the array of counts and the total to the arithmetic coder. For compression, you also pass the byte to be coded. For decompression, it returns the decoded byte. The procedure looks like this:

```
const int CONTEXT_SIZE = 1 << (n*8); // for order n
int count[CONTEXT_SIZE][256] = {{0}}; // symbol counts
int context = 0; // last n bytes packed together

// Update the model with byte c
void update(int c) {
    ++count[context][c];
    context = (context << 8 | c) % CONTEXT_SIZE;
}

// compress byte c
void compress(int c) {
    encode(count[context], c); // predict and encode
    update(c);
}

// decompress one byte and return it
int decompress() {
    int c = decode(count[context]);
    update(c);
    return c;
}
```

The functions `encode()` and `decode()` are assumed to be encoding and decoding procedures for a bitwise arithmetic coder. They each take an array of 256 counts and divide the current range in proportion to those counts. `encode()` then updates the range to the *c*'th subrange. `decode()` reads the compressed data and determines that it is bounded by the *c*'th range and returns *c*. The `update()` procedure stores the last *n* bytes in the low bits of the context.

There are a few problems with this method. First, what happens if a byte has a probability of zero? An ideal encoder would give it a code of infinite size. In practice, the encoder would fail. One fix is to initialize all elements of count to 1. Sometimes it improves compression if the initial count were smaller, say 0.5 or 0.1. This could be done effectively by increasing the increment to, say, 2 or 10.

A second problem is that a count can eventually overflow. One solution is that when a count becomes too large, to rescale all of the counts by dividing by 2. Setting a small upper limit (typically 30 to several hundred) can improve compression of data with blocks of mixed types (like text with embedded images) because the statistics reflect recent input. This is an *adaptive* or *nonstationary* model. Data with uniform statistics such as pure text are compressed better with *stationary* models, where the counts are allowed to grow large. In this case, probabilities depend equally on all of the past input.

A third problem is that the table of counts grows exponentially with the context order. Some memory can be saved by changing `count[]` to unsigned char and limiting counts to 255. Another is to replace the context with a hash, for example:

```
const int k = 5; // bits of hash per byte
const int CONTEXT_SIZE = 1 << (n*k); // order n
```

```
...
context = (context * (3 << k) + c) % CONTEXT_SIZE; // update context hash
```

The multiplier can be any odd number left shifted by another number k in the range 1 through 8. Then the last nk bits of context depend on the last n bytes of input. A larger k will result in better compression at the expense of more memory.

A fourth problem is that straightforward bitwise arithmetic coding is inefficient. The decoder must compute 256 range intervals to find the one containing the compressed data. This could be solved by using cumulative counts, i.e. $\text{count}[\text{context}][c]$ is the sum of counts for all byte values $\leq c$, but that only moves the inefficiency to the $\text{update}()$ function, which must increment up to 256 values. One solution is to encode one bit at a time using the bitwise encoder like the one described in section 3.2.

4.1.2. Bitwise encoding

The idea is to encode one bit at a time by using the previous bits of the current byte as additional context. Only two values are stored: a count of ones, count_1 , and a total count. The prediction is $\text{count}_1/\text{count}$. The update procedure is to increment count and to increment count_1 if the bit is 1. We handle zero probabilities, overflow, and large contexts as before.

Alternatively, we can avoid a (slow) division operation by storing the prediction directly. Each bitwise context is associated with a prediction that the next bit will be a 1 (initially $1/2$) and an update count (initially 0). The update rate is initially fast and decreases as the count is incremented, resulting in a stationary model. Alternatively, the count can be bounded, resulting in an adaptive model.

```
// Prediction and count for one bitwise context
struct Model {
    double prediction; // between 0 and 1 that next bit will be a 1
    int count;         // number of updates
    Model(): prediction(0.5), count(0) {}
};

Model model[CONTEXT_SIZE][256]; // context, bit_context -> prediction and count
int context = 0;                 // bitwise order n context

// Compress byte c in MSB to LSB order
void compress(int c) {
    for (int i=7; i>=0 --i) {
        int bit_context = c+256 >> i+1;
        int bit = (c >> i) % 2;
        encode(bit, model[context][bit_context].prediction);
        update(bit, model[context][bit_context]);
    }
    context = (context << 8 | c) % CONTEXT_SIZE;
}

// Decompress and return a byte
int decompress() {
    int c; // bit_context
    int bit;
    for (c = 1; c < 256; c = c * 2 + bit) {
        bit = decode(model[context][c].prediction);
        update(bit, model[context][c]);
    }
    c -= 256; // decoded byte
    context = (context << 8 | c) % CONTEXT_SIZE;
    return c;
}

// Update the model
void update(int bit, Model& m) {
    const double DELTA = 0.5;
    const int LIMIT = 255;
    if (m.count < LIMIT) ++m.count;
    m.prediction += (bit - m.prediction) / (m.count + DELTA);
}
```

The `compress()` function takes a byte `c` and compresses it one bit at a time starting with the most significant bit. At each of the 8 steps, the previously coded bits are packed into a number in the range (1..255) as a binary number 1 followed by up to 7 earlier bits. For example, if `c = 00011100`, then `bit_context` takes the 8 successive values 1, 10, 100, 1000, 10001, 100011, 1000111, 10001110. In `decompress()`, `c` plays the same role. After 8 decoding operations it has the value 100011100 and the leading 1 is stripped off before being returned.

As before, the context may also be a hash.

The update function computes the prediction error (`bit - m.prediction`) and adjusts the prediction in inverse proportion to the count. The count is incremented up to a maximum value. At this point, the model switches from stationary to adaptive.

DELTA and LIMIT are tunable parameters. The best values depend on the data. A large LIMIT works best for stationary data. A smaller LIMIT works better for mixed data types. On stationary sources, the compressed size is typically larger by $1/\text{LIMIT}$. The choice of DELTA is less critical because it only has a large effect when the data size is small (relative to the model size). With $\text{DELTA} = 1$, a series of zero bits would result in the prediction sequence $1/2, 1/4, 1/6, 1/8, 1/10$. With $\text{DELTA} = 0.5$, the sequence would be $1/2, 1/6, 1/10, 1/14, 1/18$. Cleary and Teahan (1995) measured the actual probabilities in English text and found a sequence near $1/2, 1/30, 1/60, 1/90$... for zeros and $1/2, 19/20, 39/40, 59/60$... for consecutive ones. This would fit DELTA around 0.07 to 0.1.

A real implementation would use integer arithmetic to represent fixed point numbers, and use a lookup table to compute $1/(\text{m.count} + \text{DELTA})$ in `update()` to avoid a slow division operation. ZPAQ packs a 22 bit prediction and 10 bit count into a 32 bit model element. As a further optimization, the model is stored as a one dimensional array aligned on a 64 byte cache line boundary. The bitwise context is updated once per byte as usual, but the extra bits are expanded in groups of 4 in a way that causes only two cache misses per byte. The leading bits are expanded to 9 bits as shown below, then exclusive-ORed with the bitwise context address.

```
0 0000 0001
0 0000 001x
0 0000 01xx
0 0000 1xxx
1 xxxx 0001
1 xxxx 001x
1 xxxx 01xx
1 xxxx 1xxx
```

ZPAQ fixes DELTA at $1/2$ but LIMIT is configurable to 4, 8, 12,..., 1020. The following table shows the effect of varying LIMIT for an order 0 model on 10^6 digits of π (stationary) and orders 0 through 2 on the 14 file Calgary corpus concatenated into a single data stream (nonstationary). Using a higher order model can improve compression at the cost of memory. However, direct lookup tables are not practical for orders higher than about 2. The order 2 model in ZPAQ uses 134 MB memory. The higher orders have no effect on π because the digits are independent (short of actually computing π).

	pi	Calgary corpus		
LIMIT	order-0	order-0	order-1	order-2
4	455,976	1,859,853	1,408,402	1,153,855
8	435,664	1,756,081	1,334,979	1,105,621
16	425,490	1,704,809	1,306,838	<u>1,089,660</u>
32	420,425	1,683,890	<u>1,304,204</u>	1,091,029
64	417,882	<u>1,680,784</u>	1,315,988	1,101,612
128	416,619	1,686,478	1,335,080	1,115,717
256	415,990	1,696,658	1,357,396	1,129,790
512	415,693	1,710,035	1,379,823	1,141,800
1020	<u>415,566</u>	1,726,280	1,399,988	1,150,737


4.1.3. Indirect Models

An indirect context model answers the question of how to map a sequence of bits to a prediction for the next bit. Suppose you are given a sequence like 0000000001 and asked to predict what bit is next. If we assume that the source is stationary, then the answer is 0.1 because 1 out of 10 bits is a 1. If we assume a nonstationary source then the answer is higher because we give preference to newer history. How do we decide?

An indirect model learns the answer by observing what happened after similar sequences appeared. The model uses two tables. The first table maps a context to a bit history, a state representing a past sequence of bits. The second table maps the history to a prediction, just like a direct context model.

Indirect models were introduced in [paq6](#) in 2004. A bit history may be written in the form (n_0, n_1, LB) which means that there have been n_0 zeros, n_1 ones, and that the last bit was LB (0 or 1). For example, the sequence 00101 would result in the state $(3, 2, 1)$. The initial state is $(0, 0, -)$ meaning there is no last bit.

In paq6 and its derivatives (including ZPAQ), a bit history is stored as 1 byte, which limits the number of states to 256. The state diagram below shows the allowed states in ZPAQ with n_0 on the horizontal axis and n_1 on the vertical axis. Two dots (:) represents two states for $LB=0$ and $LB=1$. A single dot represents a single state where LB can take only one value because the state is reachable with either a 0 or 1 but not both. (LB is the larger of the two counts). In general, an update with a 0 moves to the right and an update with a 1 moves up. The initial state is marked with a 0 in the lower left corner. The diagram is symmetric about the diagonal. There are a total of 219 states.



There are some exceptions to the update rule. Since it is not possible to go off the end of the diagram, the general rule is to move back to the nearest allowed state in the direction of the lower left corner (preserving the ratio of n_0 to n_1). There is another rule intended to make the model somewhat nonstationary, and that is when one of the counts is large and the other is incremented, then the larger count is reduced. The specific rule from the [ZPAQ standard](#) is that if the larger count is 6 or 7 it is decremented, and if it is larger than 7 then it is reduced to 7. This rule is applied first, prior to moving backward from an invalid state. For example, a sequence of 10 zeros, 43 ones and a zero results in:

Input	State	Rule
0000000000	(10,0,0)	Normal case, move right
1	(7,1,1)	Discount larger count
1	(6,2,1)	Discount
1	(5,3,1)	Discount
1	(5,4,1)	Normal case, move up
1	(5,5,1)	Normal case, move up
1	(4,5,1)	Move up off diagram, then back

```

1      (4,6,1)  Normal case, move up
1      (3,5,1)  Move up off diagram, then back
111    (3,8,1)  Normal, move up
1      (2,6,1)  Move off diagram and back
11111111 (2,15,1) Normal
1      (1,8,1)  Move off diagram and back
111..(20x) (1,48,1) Normal, move up
1      (1,48,1) Can't go any further
0      (2,7,0)  Discount

```

A bit history is mapped to a prediction like a direct context model, except that there is no count and the learning rate is fixed at $1/4096$ of the error. The initial prediction for each bit history is $(n_1 + 0.5)/(n_0 + n_1 + 1)$.

The details of the design were determined experimentally to improve compression slightly over the PAQ8 series, which uses a similar design.

An indirect model is more memory efficient because it uses only one byte per context instead of four. In all of the PAQ series, it is implemented as a hash table. In the PAQ8 series, the hash table is designed to allow lookups with at most 3 cache misses per byte. In ZPAQ, there are 2 cache misses per byte, similar to the direct model. The ZPAQ hash table maps a context on a 4 bit boundary to an array of 15 bit histories and an 8-bit checksum. The histories represent all 15 possible contexts that occur after up to 3 more bits. The steps are as follows:

- Once per byte, a user specified context hash is computed.
- Once every 4 bits, the context hash is combined with the first 4 bits (if any) and a hash table lookup is done.
- A hash index h and an 8 bit checksum are extracted from the 32 bit hash.
- We look for a matching checksum at addresses h , $h \text{ XOR } 1$ and $h \text{ XOR } 2$ (to stay within the cache line) and return the first match found.
- If no match is found, then the array with the smallest $n_0 + n_1$ in the current context is replaced.

In a direct context model, we don't check for hash collisions because they have only a very small effect on compression. The effect is larger for indirect models so we use an 8 bit confirmation. There is still about a 1.2% chance that a collision won't be detected but the effect on compression is very small.

The following sizes were obtained for π and the Calgary corpus with order 0 through 5 models and a hash table size of 2^{28} (268 MB). For comparison, the best results for direct context models are shown. Direct models 3, 4, and 5 use context hashes with LIMIT set to 32 and the same memory usage.

Model	Direct	Indirect
-----	-----	-----
pi order 0	<u>415,566</u>	426,343
Calgary order 0	1,680,784	1,716,974
Calgary order 1	1,304,204	1,289,769
Calgary order 2	1,089,660	1,048,050
Calgary order 3	1,017,354	<u>964,942</u>
Calgary order 4	1,069,981	1,010,329
Calgary order 5	1,232,997	1,148,677

There are many other possibilities. For example, M1, a context mixing compressor by Christopher Mattern, uses 2 dimensional context models taking 2 quantized bit histories as input.

4.2. Variable Order Models (DMC, PPM, CTW)

Fixed order models compress better using longer contexts up to a point (order 3 for the Calgary corpus). Beyond that, compression gets worse because many higher order contexts are being seen for the first time and no prediction can be made. One solution is to collect statistics for different orders at the same time and then use the longest matching context for which we know something. DMC does this for bit level predictions, and PPM for byte level predictions.


4.2.1. DMC

DMC (dynamic Markov coding) was described in a paper by Gordon Cormack and Nigel Horspool in 1987 and [implemented in C](#). The compressor [hook](#) by Nania Francesco Antonio was written in 2007 and is based on this implementation. DMC was implemented separately in ocamyd by Frank Schwelling in 2006.

DMC uses a table of variable length bit level contexts that map to a pair of counts n_0 and n_1 . It predicts the next bit with probability $n_1/(n_0+n_1)$ and updates by incrementing the corresponding count. This implements a stationary, direct context model. There are other possibilities, of course.

Contexts are not stored or looked up in the table. Rather, each table entry has a pair of pointers to the next entries corresponding to appending a 0 or 1 bit to the current context. The next context might also drop bits off the back. Normally it is arranged so that each context starts on a byte boundary. In DMC and HOOK, the initial table contains 64K entries consisting of all contexts of length 8 to 15 bits that start on a byte boundary, i.e. bitwise order 1.

In addition to updating counts, we may also add new table entries corresponding to the input just observed. We do this by "cloning" the next state with one that does not drop any bits off the back. Consider the case below where the context is 1111 and we update the model with a 0 bit. Without cloning, we would increment n_y , transition to state 110, and the next prediction would be $n_1/(n_0+n_1)$.



The cloning procedure is to allocate a new state (labeled 11100) and copy its two output pointers from 110. The new state will also "inherit" the same prediction by copying the two counts. However we proportionally reduce the two original and two new counts in proportion to the contribution from the original input (1111) and from other states. The weight $w = n_y/(n_0+n_1)$ is the fraction of counts in state 110 that came from 1111 after a 0 bit. The newly cloned state gets that fraction of the counts and the original gets the rest. This scaling maintains the condition that the input and output counts are equal. The counts are implemented as fixed or floating point numbers to allow fractional values. The transition from 1111 to 110 is changed to point to the new state 11110.

Before cloning (which uses memory), there should be sufficient justification to do so. The two conditions are that the cloned state appears often enough (n_y exceeds a threshold) and that there are sufficient other transitions to the original state that would remain after cloning ($n_0+n_1-n_y$ exceeds a threshold). In the original DMC, both thresholds are 2. Normally when the state table is full, the model is discarded and re-initialized. Setting higher thresholds can delay this from happening.

hook v1.4 also has an LZP preprocessor to encode long, repeated strings with the match length from the last matching context prior to DMC encoding. It compresses calgary.tar to 851,043 bytes in 1.9 seconds with 70 MB memory on a 2.0 GHz T3200. It compresses the files individually to a total of 840,970 bytes.

4.2.2 PPM

PPM (prediction by partial match) uses a byte-wise context model. Each context up to some maximum order is mapped to an array of counts for the next byte that occurred in this context. To predict the next byte, we find the longest context seen at least once before and allocate probabilities in proportion to those counts. The main complication is that some of the counts might be zero but we must never assign a zero probability to any byte value because if it did occur, it would have an infinite code length. This is handled by assigning an "escape" probability that the next byte will not be any of the ones that have nonzero counts, and divide those according to the frequency distribution of the next lower context. This is repeated all the way down to order 0. If any of the

256 byte values have still not been assigned a probability, then the remaining space is divided equally (an order -1 model).

Example: Suppose the input is BANANABOAT and we are at the point of coding the second B.

- The order 6 context BANANA has not been seen previously.
- The order 5 context ANANA has not been seen previously.
- The order 4 context NANA has not been seen previously.
- The order 3 context ANA has been seen once before, followed by N. The symbol we want to code is different so we code an "escape" symbol.
- The order 2 context NA was seen once, followed by N. Because N was already considered, we exclude it. Since there are no other predictions, there is nothing to code.
- The order 1 context is A. This was seen twice, in both cases followed by N which was excluded, so again there is nothing to code.
- The order 0 context was seen 6 times with values B (once), A (3 times) and N (twice). After excluding N we have 4 symbols with B having 1/4 of the probability and A having 3/4. We still need an escape probability P_{esc} that the next symbol will be other than B, A, or N. We then code B with probability $(1 - P_{\text{esc}})/4$. We would have coded A with probability $3(1 - P_{\text{esc}})/4$ and any other symbol with probability $P_{\text{esc}}/253$.

Estimating the escape probability can be complex. Suppose you draw 10 marbles from an urn. There are 8 blue, 1 green, and 1 white. What is the probability that the next marble will be a different color not seen before?

- By method "C" (Bell, Witten, and Cleary, 1989), 3 of the 10 marbles you drew had a novel color, so the probability would be 0.3.
- But novel colors would be expected to show up early. By method "X" (Witten and Bell, 1991), 2 of the colors appeared exactly once, so the probability would be $2/10 = 0.2$.
- Method "X" can fail if no colors appeared exactly once (because the escape probability would be 0). Method "XC" is to use "X" when possible, falling back to "C" when needed. Method XC was shown experimentally to compress better than C.
- A secondary escape estimation (SEE) model would look at other cases with the same or similar distribution as {8, 1, 1}, and take the fraction of those where a novel color appeared next. This improves on XC.

Method X was shown to be optimal under certain assumptions, including that the source is stationary. Of course, that is not always the case. Suppose you receive the sequence "BBBBBBBBBWG" and are asked to predict whether the next character will be novel. The answer might be different for the sequence "WBBBBBBBBB".

These methods are not used in modern PPM compressors. Rather, better compression can be obtained by learning the escape probability directly. The method is known as secondary escape estimation (SEE). Charles Bloom (1998) used SEE in [PPMZ](#). Dmitry Shkarin (2002) refined this method in [ppmd](#). ppmd variant I, released as source code in 2002, is used in several archivers such as the open source [7zip](#) and [freearc](#), and the commercial [WinZIP](#), [Stuffit](#), and [WinRAR](#). (WinRAR uses an older variant H). A newer variant ppmd J in 2006 gets slightly better compression than I. The code is the basis of slower but better compressing programs such as [ppmonstr](#) and [durilca](#). A variant of durilca using 13 GB memory is top ranked on the large text benchmark.

ppmd uses a complex SEE model. It considers 3 cases:

1. binary context - in the highest order context only one byte value has appeared (one or more times).
2. nm-context - two or more values have appeared in the highest order context.
3. m-context - two or more values have appeared, and one or more have appeared (are masked) in a higher order context.

In the binary context, a 13 bit context to a direct context model is constructed:

- 7 bits for the quantized count of the one symbol that appeared.
- 2 bits for the quantized alphabet size of the next lower order context.
- 1 bit for the quantized probability of the previously coded byte.

- 1 bit to indicate whether the two high order bits of the previous byte are 00 (to distinguish letters from other characters in text).
- 1 bit to indicate whether the two high order bits of the predicted byte are 00.
- 1 bit for the quantized number of successive bytes that were not escaped.

In the nm-context, the program fits the frequency distribution to a geometric approximation such that the n'th most frequent value is proportional to r^n . Then r is the context.

In the m-context, the SEE context is constructed from:

- The number of unmasked values with counts > 0, quantized to 25 levels.
- 2 bits based on comparison of the alphabet sizes of the current and next higher and lower order contexts.
- 1 bit to indicate the two high bits of the previous byte are 00.
- The average frequency per value, quantized to 4 levels.

ppmonstr uses an even more complex SEE context, and additionally uses interpolation to smooth some of the quantized contexts. It also adjusts the prediction for the most probable byte using secondary symbol estimation (SSE). This is a direct context model taking as input the quantized prediction and a (very complex) context and outputting a new prediction.

Both programs use other techniques to improve compression. They use partial update exclusion. When a character is counted in some context, it is counted with a weight of 1/2 in the next lower order context. Also, when computing symbol probabilities, it performs a weighted averaging with the predictions of the lower order context, with the weight of the lower order context inversely proportional to the number of different higher order contexts of which it is a suffix.

Statistics are stored in a tree which grows during modeling. When the memory limit is reached, the tree is discarded and rebuilt from scratch. Optionally, the statistics associated with each context are scaled down and those with zero counts are pruned until the tree size becomes smaller than some threshold (25%). This improves compression but takes longer.

Although bitwise arithmetic encoding can be inefficient, ppmd is in practice faster than many equivalent bitwise context mixing models. First, a byte is encoded as a sequence of escape symbols followed by a non-escape. Each of these encodings is from a smaller alphabet. Second, the alphabet within each context can be ordered so that the most likely symbols are first. This reduces the number of operations in the majority of cases.

Shown below are compressed sizes of the Calgary corpus as a tar file and separate files. Compression and decompression times are the same. Option -o16 means use maximum order 16. -m256 says use 256 MB memory. -r1 says to prune the context tree rather than discard it.

Compressor	Options	calgary.tar	14 files	Time
ppmd J	-o16 -m256 -r1	754,243	740,737	2 sec
ppmonstr J	-o16 -m256 -r1	674,704	668,459	8 sec
durilca 0.5	-o128 -m256	672,752	666,216	10 sec

Some additional programs using PPM are as follows:

- slim (Serge Voskoboynikov, 2004) is based on PPMII (ppmd/ppmonstr) with optimizations for text and other file types. The last version is 23d. Different versions are incompatible with each other.
- [epm](#) (Serge Osnach, 2003) is based on PPMII with 20 undocumented parameters specified by a single digit option. A separate program (epmopt) tries to find the best compression by compressing over and over, adjusting each option up or down one at a time until no further improvements can be found. Because this process is very slow, it is possible to specify optimization on a prefix of the input file.
- [enc](#) (Serge Osnach, 2003) tries 6 different compression methods (ppmd and 3 variants, PAQ3n, and store) and chooses the best one.
- [CTXf](#) (Nikita Lesnikov, 2003) is derived from ha and ppmd.
- [bee](#) (Andrew Filinsky, Melchiorre Caruso, 2005) is an open source archiver based on unbounded order PPM. It also includes an optimizer (beeopty) like epm.
- [uharc](#) (Uwe Herklotz, 2005) is an archiver with options for LZ77 (fastest decompression), LZW (fastest compression), and PPM (best compression).

- [ppmvc](#) (Przemysław Skibiński, 2006) uses LZP or ROLZ to encode long string matches prior to compression with ppmd (Skibiński and Grabowski, 2004).
- ppms (Dmitry Shkarin, 2006), is a small memory (1 MB) version of ppmd.
- ha (Harri Hirvola, 1993) is an open source order 4 PPM archiver.
- [cm](#) (Bulat Ziganshin, 1998) is a static PPM encoder. Rather than the decoder rebuilding the context model as usual, the encoder builds it and transmits it to the decoder prior to transmitting the data.

4.2.3. CTW

[CTW](#) (Context Tree Weighting) is an algorithm developed by Frans Willems, Yuri Shtarkov and Tjalling Tjalkens in 1995 and implemented by Erik Franken and Marcel Peeters in 2002. In the U.S. and Europe it is protected by patents EP0913035B1, EP0855803, US6150966, US5986591, US5864308 owned by Koninklijke KPN NV, Netherlands.

CTW, like DMC, predicts one bit at a time using bit count statistics for contexts that begin on byte boundaries. Like PPM and unlike DMC, it combines the predictions of different order contexts from order 0 up to some maximum order D . Each context is associated with two bit counters, n_0 and n_1 , and a mixing weight, w . Each context order i from 0 through D computes an estimator e_i that the next bit will be a 1 depending on the two counts. The highest order context predicts $p_D = e_D$. All of the lower order contexts predict p_i which is a weighted average of e_i and the next higher order p_{i+1} . The final prediction is either p_0 or p_1 .

On update with bit y (0 or 1), the count n_y is incremented and w is updated to favor the more accurate prediction in each context, either toward e_i or p_{i+1} .

Specifically, CTW uses the Krichevski-Trofimov (KT) estimator:

$$e_i(0) = (n_0 + 1/2)/(n_0 + n_1 + 1)$$

$$e_i(1) = (n_1 + 1/2)/(n_0 + n_1 + 1) = 1 - e_i(0).$$

except for the special case where one of the counts is 0, which we consider later. Then the prediction, expressed as a ratio, is:

$$p_D(0) = e_D(0)$$

$$p_D(1) = e_D(1)$$

$$p_i(1)/p_i(0) = (w_i e_i(1) + p_{i+1}(1))/(w_i e_i(0) + p_{i+1}(0)), i < D.$$

Given bit y , the weight and one count for each context is updated:

$$w_i := w_i e_i(y)/p_{i+1}(y), i < D$$

$$n_y := n_y + 1.$$

CTW uses a zero redundancy estimator for the special case where one of the two counts is zero. The idea is to predict more strongly the other bit. Consider the case of $n_0 = 0$, i.e. only ones have been observed. Then define $k(n)$ as the cumulative probability of n consecutive bits by the KT estimator, and define the ZR estimator $e()$ as:

$$k(0) = 1/2$$

$$k(i) = k(i-1)(i+1/2)/(i+1), i > 0$$

$$e(1) = (n_1 + 1/2)/(n_0 + n_1 + 1), n_1 > 0, n_0 > 0 \text{ (KT estimator)}$$

$$e(1) = k(n_1)/(2(n_1 + 1)k(n_1) + n_1 + 1), n_1 > 0, n_0 = 0 \text{ (ZR estimator)}$$

All estimators are symmetric with respect to 0 and 1. The probability of a 1 given the sequence 00000 is the same as the probability of 0 given 11111. The ZR estimator gives a smaller probability for these cases than the KT estimator. The following table gives the probability that a 1 will follow n_0 zeros and no ones with each estimator.

n_0	KT	ZR
-------	----	----

0	0.500000	0.500000
1	0.250000	0.107143
2	0.166667	0.064103
3	0.125000	0.044192
4	0.100000	0.032984
5	0.083333	0.025908
6	0.071429	0.021089
7	0.062500	0.017625
8	0.055556	0.015032
9	0.050000	0.013029
10	0.045455	0.011441
100	0.004950	0.000499
1000	0.000500	0.000017

For large runs of identical bits, the total code length for KT grows logarithmically but the ZR estimator is bounded to 2 bits. However, ZR adds 1 bit of redundancy when the other bit value eventually occurs. For large n_0 , the ZR estimated probability of a 1 approaches 0 at the rate $n_0^{-3/2}$. This is faster than n_0^{-1} for the KT estimator.

The CTW implementation uses a context tree to store the two counts and w . Each count is 8 bits. If one count exceeds 255, then both counts are divided by 2. It stores $\log(w)$ rather than w and does all arithmetic in the log domain so that multiplication and division are implemented using addition and subtraction. The KT and ZR estimators are looked up in a table indexed by the two counts.

A context tree is a 256-ary tree where each node at depth $N+1$ from the root represents an order- N context. The parent of each node is the suffix obtained by removing one byte. The children of the root node represent the 255 possible partial byte order 0 contexts of length 0 to 7 bits. To save space, children of contexts that have only occurred once are not stored. Each tree node uses 8 bytes, including a 24 bit pointer to the context in a separate file buffer.

CTW is best suited for stationary sources. The published CTW implementation compresses the Calgary corpus to 767,594 bytes as 14 separate files in 62 seconds with options `-n16M -b16M -d12` set for maximum compression. With the same options, it compresses `calgary.tar` to 965,855 bytes. `-d12` sets the maximum context order to 12. `-b16M` sets the file buffer to 16 MB. `-n16M` limits the tree to 16 million nodes. When the tree is full, no new contexts are added but the counts and weights continue to be updated. These settings require 144 MB memory, the maximum that the published implementation can use.

4.3. Context Mixing

Context mixing algorithms based on the PAQ series are top ranked on many benchmarks by size, but are very slow. These algorithms predict one bit at a time (like CTW) except that weights are associated with models rather than contexts, and the contexts need not be mixed from longest to shortest context order. Contexts can be arbitrary functions of the history, not just suffixes of different lengths. Often the result is that the combined prediction of independent models compresses better than any of the individuals that contributed to it.

DMC, PPM, and CTW are based on the premise that the longest context for which statistics is available is the best predictor. This is usually true for text but not always the case. For example, in an audio file, a predictor would be better off ignoring the low order bits of the samples in its context because they are mostly noise. For image compression, the best predictors are the neighboring pixels in two dimensions, which do not form a contiguous context. For text, we can improve compression using some contexts that begin on word boundaries and merge upper and lower case letters. In data with fixed length records such as spreadsheets, databases or tables, the column number is a useful context, sometimes in combination with adjacent data in two dimensions. PAQ based compressors may have tens or hundreds of these different models to predict the next input bit.

A fundamental question is how do we combine predictions? Suppose you are given two predictions $p_a = P(y=1|A)$ and $p_b = P(y=1|B)$, probabilities that the next bit y will be a 1 given contexts A and B . Assume that A and B have occurred often enough for the two models to make reliable guesses, but that both contexts have never occurred together before. What is $p = P(y=1|A,B)$?

Probability theory does not answer the question. It is possible to create sequences where p can be anything at all for any p_a and p_b . For example, we could have $p_a=1$, $p_b=1$, $p=0$. But intuitively, we should do some kind of averaging or weighted averaging. For example, if we wish to estimate $P(\text{car accident} \mid \text{dark and raining})$ given $P(\text{car accident} \mid \text{dark})$ and $P(\text{car accident} \mid \text{raining})$, we would expect the effects to be additive.

Furthermore, we want to mix predictions weighed by confidence. If $p_a = 0.5$ and $p_b = 0.99$, then intuitively it seems that p_b expresses greater confidence in its prediction, so it should be given greater weight. All PAQ versions do this. Early versions expressed predictions as pairs of counts (observed zeros and ones) and added them together. This implicitly gave greater weight to predictions near 0 or 1 because such predictions can only occur when one of the counts is large. Later versions improved on this by transforming probabilities into the logistic domain, $\log(p/(1-p))$, followed by averaging. This allowed greater flexibility in modeling techniques. Logistic mixing was introduced in 2005, but it wasn't until later when Mattern (2012) proved that logistic mixing is optimal in the sense of minimizing [Kullback-Leibler divergence](#), or wasted coding space, of the input predictions from the output mix.

In most PAQ based algorithms, there is also a procedure for evaluating the accuracy of models and further adjusting the weights to favor the best ones. Early versions used fixed weights.

4.3.1. Linear Evidence Mixing

In PAQ6 (Mahoney, 2005a), a probability is expressed as a count of zeros and ones. Probabilities are combined by weighted addition of the counts. Weights are adjusted in the direction that minimizes coding cost in weight space. Let n_{0i} and n_{1i} be the counts of 0 and 1 bits for the i 'th model. The combined probabilities p_0 and p_1 that the next bit will be a 0 or 1 respectively, are computed as follows:

$$\begin{aligned} S_0 &= \varepsilon + \sum_i w_i n_{0i} = \text{evidence for 0} \\ S_1 &= \varepsilon + \sum_i w_i n_{1i} = \text{evidence for 1} \\ S &= S_0 + S_1 = \text{total evidence} \\ p_0 &= S_0/S = \text{probability that next bit is 0} \\ p_1 &= S_1/S = \text{probability that next bit is 1} \end{aligned}$$

where w_i is the non-negative weight of the i 'th model and ε is a small positive constant needed to prevent degenerate behavior when S is near 0.

The optimal weight update can be found by taking the partial derivative of the coding cost with respect to w_i . The coding cost of a 0 is $-\log p_1$. The coding cost of a 1 is $-\log p_0$. The result is that after coding bit y (0 or 1), the weights are updated by moving along the cost gradient in weight space:

$$w_i := \max[0, w_i + (y - p_i)(S n_{1i} - S_1 n_i) / S_0 S_1]$$

Counts are discounted to favor newer data over older. A pair of counts is represented as a bit history similar to the one described in section 4.1.3, but with more aggressive discounting. When a bit is observed and the count for the opposite bit is more than 2, the excess is halved. For example if the state is $(n_0, n_1) = (0, 10)$, then successive zero bits will result in the states $(1, 6)$, $(2, 4)$, $(3, 3)$, $(4, 2)$, $(5, 2)$, $(6, 2)$.

4.3.2. Logistic Mixing

PAQ7 introduced logistic mixing, which is now favored because it gives better compression. It is more general, since only a probability is needed as input. This allows the use of direct context models and a more flexible arrangement of different model types. It is used in the PAQ8, LPAQ, PAQ8HP series and in ZPAQ.

Given a set of predictions p_i that the next bit will be a 1, and a set of weights w_i , the combined prediction is:

$$p = \text{squash}(\sum_i w_i \text{stretch}(p_i))$$

where

$$\text{stretch}(p) = \ln(p/(1 - p))$$

$$\text{squash}(x) = \text{stretch}^{-1}(x) = 1/(1 + e^{-x})$$

The probability computation is essentially a neural network evaluation taking stretched probabilities as input. Again we find the optimal weight update by taking the partial derivative of the coding cost with respect to the weights. The result is that the update for bit y (0 or 1) is simpler than back propagation (which would minimize RMS error instead).

$$w_i := w_i + \lambda (y - p) \text{stretch}(p_i)$$

where λ is the learning rate, typically around 0.01, and $(y - p)$ is the prediction error. Unlike linear mixing, weights can be negative.

Compression can often be improved by using a set of weights selected by a small context, such as a bitwise order 0 context.

In PAQ and ZPAQ, `squash()` and `stretch()` are implemented using lookup tables. In PAQ, both output 12 bit fixed point numbers. A stretched probability has a resolution of 2^{-8} and range of -8 to 8. Squashed probabilities are multiples of 2^{-12} . ZPAQ represents stretched probabilities as 12 bits with a resolution of 2^{-6} and range -32 to 32. Squashed probabilities are 15 bits as an odd multiple of 2^{-16} . This representation was found to give slightly better compression than PAQ.

ZPAQ allows different components (models and mixers) to be connected in arbitrary ways. All components output a stretched probability, which simplifies the mixer implementation. ZPAQ has 3 types of mixers:

- AVG performs weighted averaging of two (stretched) predictions with fixed, user specified weights that add to 1.
- MIX2 is like AVG except that weights are updated with the constraint that they add to 1. The user specifies λ and a context to select a pair of weights.
- MIX is like MIX2 except that it takes any number of inputs and does not constrain the weights to add to 1. A 2 input MIX often gives better compression than a MIX2.

Mixer weights in PAQ are 16 bit signed values to facilitate vectorized implementation using MMX/SSE2 parallel instructions. In ZPAQ, 16 bits was found to be inadequate for best compression. Weights were expanded to 20 bit signed values with range -8 to 8 and precision 2^{-16} .

4.3.3. Secondary Symbol Estimation (SSE)

SSE (secondary symbol estimation) is implemented in all PAQ versions beginning with PAQ2. Like in `ppmonstr`, it inputs a prediction and a context and outputs a refined prediction. The prediction is quantized typically to 32 or 64 values on a nonlinear scale with finer resolution near 0 and 1 and sometimes interpolated between the two closest values. On update, one or both values are adjusted to reduce the prediction error, typically by about 1%. A typical place for SSE is to adjust the output of a mixer using a low order (0 or 1) context. SSE components may be chained in series with contexts typically in increasing order. Or they may be in parallel with independent contexts, and the results mixed or averaged together. The table is initialized so that the output prediction is equal to the input prediction for all contexts.

SSE was introduced to PAQ in PAQ2 in 2003 with 64 quantization levels and no interpolation. Later versions used 32 levels and interpolation with updates to the two nearest values above and below. In some versions of PAQ, SSE is also known as an APM (adaptive probability map).

ZPAQ allows a SSE to be placed anywhere in the prediction sequence with any context. Recall that ZPAQ probabilities are stretched by mapping to $\ln(p/(1-p))$ as a 12 bit fixed point number in the range -32 to +32 with resolution 1/64. The SSE input prediction is clamped and quantized to an odd multiple of 1/2 between -15.5 and 15.5. The low 6 bits serve as an interpolation weight. For example, if $\text{stretch}(p) = 2.7$, then the two table entries are selected by `below=2.5` and `above=3.5`, and the interpolation weight is 0.2. Then the output prediction is $\text{SSE}[\text{context}][\text{below}] * (1-w) + \text{SSE}[\text{context}][\text{above}] * w$. Upon update with bit y , the table entry nearest the input

prediction (SSE[context][below] in this example) is updated by reducing the prediction error ($y - \text{output}$) by a user specified fraction.

There are other possibilities. CCM, a context mixing compressor by Christian Martelock, uses a 2 dimensional SSE taking 2 quantized predictions as input.

4.3.4 Indirect SSE (ISSE)

ISSE (indirect secondary symbol estimation) is a technique introduced in [paq9a](#) in Dec. 2007 and is a component in ZPAQ. The idea is to use SSE as a direct prediction method rather than to refine an existing prediction. However, SSE does not work well with high order contexts because the large table size uses too much memory. More generally, a large model with lots of free parameters (each table entry is a free parameter) will overfit the training data and have no predictive power for future input. As a general rule, a model should not be larger than the input it is trained on.

ISSE does not use a 2-D table. Instead it first maps a context to a bit history as with an indirect context map. Then the 8-bit bit history is used as a context to select the pair of weights for a 2 input mixer taking the input prediction and a fixed constant as its two inputs. The weights are initialized to (1.0, 0.0) meaning that the initial output prediction is equal to the input.

PAQ9A and the default compression mode of ZPAQ both start with an order 0 model prediction and refine it using a chain of ISSE components in increasing order.

In ZPAQ, the weights are 20 bit signed, fixed point numbers with range -8 to 8 and precision 2^{-16} like in a MIX. The fixed input is 4.0 and the learning rate is fixed at $\lambda = 2^{-8}$.

4.3.5. Match Model

A match model finds the last occurrence of a high order context and predicts whatever symbol came next. The accuracy of the prediction depends on the length of the context match. Longer matches generally give more confidence to the prediction. Typically a match model of order 6-8 is mixed with lower order context models. A match model is faster and uses less memory than a corresponding context model but does not model well for low orders.

Match models are used in PAQ and ZPAQ. They consist of a rotating history buffer and a hash table mapping contexts to pointers into the buffer. In ZPAQ, a pointer to the match is maintained until a mismatching bit is found. The model will then look for a new match at the start of the next byte. On each byte boundary, the buffer is updated with the modeled byte and the hash table is updated so that the current context hash points to the end of the buffer. ZPAQ allows both the hash table size and buffer size to be user specified (as powers of 2). For best compression, the history buffer should be as large as the input (if this is practical) and the hash table size is typically 1/4 of this. Because each pointer is 4 bytes, both data structures use the same amount of memory.

Match models in PAQ maintain multiple context hashes of different orders and multiple pointers into the buffer. The prediction is indirect by mapping the match length to a prediction through a direct context model. ZPAQ uses a simpler match model with just one pointer and one hash, although it is possible to have multiple, independent match models. The prediction for a match of L bytes (or $8L$ bits) is that the next bit will be the same with probability $1 - 1/8L$.

The user may specify the context length by using a rolling hash that depends on the desired number of characters. If h is the context hash, c is the input byte, then the update:

$$h = h * ((2^{k+1}) \ll m) + c;$$

specifies an order $\text{ceil}(n/m)$ context hash for a hash table size of 2^n and any k . For example, the hash update " $h = h * 40 + c;$ " ($m = 3$) is an order 6 context hash for a table size of 2^{18} . Only the low 18 bits of h would be used to index the hash table of this size, and these bits depend only on the last 6 values of c .

4.3.6 PAQ Models

The high compression ratio (and slow speed) in [PAQ](#) comes from using many different context models for different types of data. These are described in historical order.

Schmidhuber and Heil (1996) developed an experimental neural network data compressor. It used a 3 layer network trained by back propagation to predict characters from an 80 character alphabet in text. It used separate training and prediction phases. Compressing 10 KB of text required several days of computation on an HP 700 workstation.

Mahoney (2000) made several improvements that made neural network compression practical.

- The neural network predicts one bit at a time instead of one character.
- The training is online, eliminating multiple passes.
- The first layer of the neural network is replaced by a hash function that selects one neuron per context for each of the orders 1 through 5.

These changes make the algorithm about 10^5 times faster, mainly because only a few neurons (out of millions) are active at any one time. To make the training online, it is necessary to add a pair of counters to each weight (to count 0 and 1 bits) so that the learning rate is initially large. The rate decreases in inverse proportion to the smaller of the two counts.

There are 3 versions: P5, P6, and P12. P5 uses 256 KB memory to represent 5 orders using 2^{16} input neurons (each representing a context hash) and one output (the bit prediction). P6 uses 2^{20} input neurons. P12 is the same size as P6 but adds a whole word model. This context hash depends only on alphabetic characters mapped to lower case, and is reset after a nonalphabetic character. It improves compression of text files.

In PAQ1 (Mahoney, 2002), it was realized that the counts alone could be used to make predictions, so the weights were eliminated. Predictions are combined by adding the 0 and 1 counts associated with each context. Each counter is 1 byte.

PAQ2 added SSE by Serge Osnach in May 2003. It uses 64 quantization levels without interpolation.

PAQ3 modified SSE to use 32 levels with interpolation in Sept. 2003.

PAQ3N by Serge Osnach in Oct. 2003 added sparse models: three order-2 models that skipped 1, 2, or 3 bytes of context between the two context bytes. This improves compression of some binary files.

PAQ4 (Nov. 2003) uses adaptive linear weighting of models as described in section 4.3.1. It also introduced a record model. It identifies structures that repeat at regular intervals, as found in spreadsheets, tables, databases, and images, and adds contexts of adjacent bytes in two dimensions.

PAQ5 (Dec. 2003) has some minor improvements over PAQ4, including word models for text, models for audio and images, an improved hash table, and modeling of run lengths within contexts. It uses two mixers with different contexts to select their weights. The two mixer outputs are averaged together. It uses about 186 MB of memory.

PAQ6 (Jan. 2004) adds models for x86 code (modeling jump/call addresses) and CCITT images and more aggressive discounting of opposing bit counts. It takes options allowing up to 1616 MB memory. It is the basis of a number of forks and dozens of versions. An early version won the Calgary challenge. Many other models and optimizations were added by Berto Destasio, Johan de Bock, David A. Scott, Fabio Buffoni, Jason Schmidt, Alexandar Rhatushnyak (PAQAR), Przemyslaw Skibinski (PASQDA, text preprocessing), Rudi Cilibrasi, Pavel Holoborodko, Bill Pettis, Serge Osnach, and Jan Ondrus.

PAQAR (v1.0 to 4.0, May-July 2004) by Alexander Rhatushnyak is a PAQ6 fork which is the basis of several winning submissions to the Calgary Challenge. The primary difference is a greatly increased number of mixers and SSE chains.

PAQ7 (Dec. 2005) was a complete rewrite. It uses logistic mixing rather than linear mixing, as described in section 4.3.2. It has models for color BMP, TIFF, and JPEG images. The BMP and TIFF models use adjacent

pixels as context. JPEG is already compressed. The model partially undoes the compression back to the DCT (discrete cosine transform) coefficients and uses these as context to predict the Huffman codes.

PAQ8A (Jan. 2006) adds a E8E9 preprocessor to improve compression of x86 (EXE and DLL) files. The preprocessor replaces relative CALL and JMP addresses with absolute addresses, which improves compression because an address may appear multiple times. Many other compressors use this technique.

PASQDA (v1.0-v4.4, Jan. 2005 to Jan. 2006) is a fork by Przemyslaw Skibinski. It adds an external dictionary to replace words in text files with 1 to 3 byte symbols. This technique was used successfully in the Hutter Prize and in the top ranked programs in the large text benchmark. PAQ8A2, PAQ8B, PAQ8C, PAQ8D, PAQ8E, and PAQ8G (Feb. to Mar. 2006) also use this technique, as does PAQAR 4.5 by Alexander Rhatushnyak.

PAQ8F (Feb. 2006) adds a more memory efficient context model and a new indirect model: The byte history within a low order context is modeled by another low order context.

PAQ8L (Mar. 2006) adds a DMC model. Its predictions are mixed with those of other models.

As of Feb. 2010, development remains active on the PAQ8 series. There have been hundreds of versions with improvements and additional models. The latest is PAQ8PX_V67. Most of the improvements have been for file types not included in the Calgary corpus such as x86, JPEG, BMP, TIFF, and WAV.

A benchmark for the Calgary corpus is given below for versions of PAQ from 2000 to Jan. 2010 showing major code changes. About 150 intermediate versions with minor improvements have been omitted. Older programs marked with * were benchmarked on slower machines such as a 750 MHz Duron and have been adjusted to show projected times on a 2.0 GHz T3200, assumed to be 5.21 times faster. Sizes marked with a D use an external English dictionary that must be present during decompression. The size shown does not include the dictionary, so it is artificially low. However, including it would give a size artificially high because the dictionary is not extracted from the test data. All versions of PAQ are archivers that compress in solid mode, exploiting similarities between files. Decompression time is about the same as compression time.

Compressor	Calgary	Seconds	Memory	Date	Author	Major changes
-----	-----	-----	-----	-----	-----	-----
P5	992,902	6.1*	256 KB	2000	Mahoney	64K x 1 neural network
P6	841,717	7.4*	16 MB	2000	Mahoney	1M neurons
P12	831,341	7.5*	16 MB	2000	Mahoney	Word context model
PAQ1	716,704	13*	48 MB	2002	Mahoney	Linear mixing with fixed weights
PAQ2	702,382	18*	48 MB	May 2003	Osnach	SSE
PAQ3	696,616	15*	48 MB	Sep 2003	Mahoney	Interpolated SSE
PAQ3N	684,580	30*	80 MB	Oct 2003	Osnach	Sparse models
PAQ4	672,134	43*	84 MB	Nov 2003	Mahoney	Adaptive mixer weights, record models
PAQ5	661,811	70*	186 MB	Dec 2003	Mahoney	Models for text, audio, images, runs, 2 mixers
PAQ6 -6	648,892	99*	202 MB	Jan 2004	Mahoney	Models for PIC, x86
PAQAR 4.0 -6	604,254	408*	230 MB	Jul 2004	Rhatushnyak	Many mixers and SSE chains
PAQ7 -5	611,684	142*	525 MB	Dec 2005	Mahoney	Logistic mixing, image models
PAQ8A -4	610,624	152*	115 MB	Jan 2006	Mahoney	E8E9 preprocessor
PASQDA 4.4 -7	D 571,011	283*	470 MB	Jan 2006	Skibinski	PAQ7 + external dictionary
PAQAR 4.5 -5	D 570,374	299*	191 MB	Feb 2006	Rhatushnyak	PAQAR + external dictionary
PAQ8F -6	605,650	161*	435 MB	Feb 2006	Mahoney	Byte-wise indirect model, memory tuning
PAQ8L -6	595,586	368	435 MB	Mar 2007	Mahoney	DMC model
PAQ8PX_V67 -6	598,969	469	421 MB	Jan 2010	Ondrus	Improved JPEG, TIFF, BMP, WAV models

Since 2007, [development has continued](#) on PAQ. In addition to PAQ8PX, there are 3 additional forks, no longer under active development.

- [PAQ8HP](#) by Alexander Rhatushnyak, a basis for the [Hutter Prize](#). The series was optimized for this data and the [large text benchmark](#). It uses a dictionary transform based on [XWRT](#) which replaces words from a dictionary with 1 to 3 byte codes. The dictionary for the paq8hp series is optimized for these benchmarks, which is allowed under the rules. There were 12 versions from Aug. 2006 through May 2007.
- [LPAQ](#) by Matt Mahoney with later versions by Alexander Rhatushnyak. This was a "lite" PAQ, faster but with less compression. Later versions were tuned for text. It includes a mix of contexts of different orders and a match model. There were 24 version from July 2007 through Feb. 2009.
- [PAQ9A](#) (Dec. 2007) was an experiment in LZP preprocessing followed by chained ISSE modeling. Compression is similar to early versions of LPAQ. The LZP preprocessor removes long redundant strings.

It includes sparse and text models but no match model because high order contexts were removed.

4.3.7. ZPAQ

Of the hundreds of PAQ versions, no program can decompress files compressed by any other version. The goal of the proposed [ZPAQ standard](#) is to solve this problem. It specifies a format in which a description of the compression algorithm is stored in the compressed archive. The specification includes a [reference decoder](#).

The specification does not describe the compression algorithm. However, several compression programs and models are available on the [ZPAQ page](#). There is a ZPAQ program that takes a configuration file to describe the compression algorithm, as well as other programs like ZPIPE that use a fixed compression algorithm. All of these produce files that can be decompressed with the reference decoder or by any of these programs. The standard was published in Mar. 2009 by Matt Mahoney.

ZPAQ describes an archive format, although it may be used for single file compression or memory to memory compression. A compressed stream consists of a sequence of blocks that are independent and can be reordered. Each block starts with a header that describes the decompression algorithm. A block consists of a sequence of compressed segments that must be decompressed in sequence from the start of the block. A segment may be a file or a part of a file. Each segment has an optional file name, an optional comment (file size, timestamp, etc.), and ends with an optional 20 byte [SHA-1](#) checksum. If the file name is omitted, then the decompressor must supply it.

An algorithm description consists of a network of components (each making a prediction), a program that computes the bitwise contexts for each component, and an optional program that post-processes the output. Both programs are written in a language called ZPAQL which is compiled or interpreted by the decompressor. If the post-processing program is present, then it is appended to the front of the first uncompressed segment and compressed along with its input data. If not, then the data is compressed directly with a one byte header to indicate its absence.

Up to 255 components are placed in an array. Each component in the model inputs a context hash and possibly the predictions of earlier components, and outputs a stretched prediction. The output of the last component is squashed and used to arithmetic code the data one bit at a time. The components are as follows:

- CONST - outputs a fixed prediction in the stretched range -16 to 16 in increments of 1/16.
- CM - a direct context model (section 4.1.2). Inputs a context and outputs a prediction, which is then updated to reduce the prediction error. The table size (and thus the context size) may be any power of 2. The count limit, which controls the learning rate, ranges from 4 to 1020 in increments of 4.
- ICM - an indirect context model (section 4.1.3). Maps a context to a bit history, which is mapped to a prediction. The size may be any power of 2 at least 64 bytes.
- MATCH - (section 4.3.5). Predicts the next bit from the last matching context in the history buffer. The hash table size and buffer size may be any power of 2.
- AVG - weighted average of any 2 predictions. The weight ranges from 0 to 1 in increments of 1/256.
- MIX - (section 4.3.2). Adaptively averages predictions using weights selected by a context. Specifies a context size, a range of inputs, a learning rate in increments of 1/4096 up to 1/16, and an 8 bit mask to turn on or off bits of the current byte in the order 0 context.
- MIX2 - (section 4.3.2). Like a MIX but with any 2 inputs, and the weights constrained to add to 1.
- SSE - (section 4.3.3). Maps a context and a prediction to a new prediction using a 2-D table. Specifies a context size, input prediction, and initial and maximum counts for the context model (0..255 and 4..1020 by 4 respectively) and a bit mask for the order 0 context like in a MIX or MIX2.
- ISSE - (section 4.3.4). Maps a context and a prediction to a new prediction using a bit history to select the weights for a 2 input mixer with the other input constant. Specifies a context size and input prediction.

Contexts are computed by a ZPAQL program that is called once per modeled byte with that byte as input. The program places the context hashes in an array H of 32 bit unsigned values. Each element of H is the input for one component, beginning with H[0] for the first one. Only the low bits of the hash are used, depending on the table size of each component. Because each bit must be modeled, the context hashes are combined with the previous bits of the current byte. This is done by expanding the previous bits to a 9 bit value (ICM or ISSE) or 8 bits (all others) and adding it to the bitwise context.

ZPAQL is designed for fast execution rather than ease of programming. It resembles an assembly language instruction set. A ZPAQL program runs on a virtual machine with the following state, all initialized to 0 at the start of a block:

- A 16 bit program counter.
- 32 bit registers A, B, C, D, R0 through R255. The accumulator A is used for input and the results of most computations.
- A 1 bit flag F to hold the result of comparisons.
- An array H of 32 bit values, indexed by D. The first 256 elements of H hold computed contexts.
- An array M of 8 bit values, indexed by B or C.

The sizes of H and M are specified as powers of 2 in the block header. Most instructions are either 1 byte, or 2 bytes including a numeric operand. The instruction set is as follows:

- Assignment, for example $A=B$ meaning copy B into A. The left and right operands may be A, B, C, D, *B, *C, or *D. *B and *C mean the element of M addressed by the low bits of B or C (depending on the size of M). *D means an element of H. The right operand may also be a constant from 0 through 255.
- Arithmetic, for example, $A+=B$ meaning add B to A. Operands are as above except that the left operand is always A. Operations are $+=$ $-=$ $*=$ $/=$ $\%=$ $\&=$ $|=$ $\wedge=$ $<<=$ $>>=$ $\&\sim$ with their usual meanings in C/C++. Division or mod by 0 is 0. $\&\sim$ means $\&\sim$.
- Comparison $<==$ $>$ only. The left operand is A. The result goes in F.
- Unary operations $++$ (increment) $--$ (decrement) $!$ (bit compliment) $<>A$ (swap with A), $=0$ (clear). The operand may be A, B, C, D, *B, *C, or *D, always written first as in B! to compliment the bits of B ($B = \sim B$ or $B = -B-1$); $A<>A$ is not valid.
- HASH meaning $A = (A + *B + 512) * 773$, a convenient hashing function.
- HASHD meaning $*D = (*D + A + 512) * 773$.
- OUT meaning output A (used only in postprocessing).
- HALT meaning end execution.
- Conditional jumps such as JT -4 meaning jump back 4 bytes (from the next instruction) if F is true. JF means jump if false. JMP is unconditional. The operand is -128 to 127.
- Long jump LJ with a 2 byte operand 0 through 65535 from the start of the program.
- Access to R0 through R255 for example $R=A\ 5$ meaning $R5=A$. Operations are $A=R$, $B=R$, $C=R$, $D=R$, $R=A$.

The post-processor (called PCOMP), if it is present, is called once per decoded byte with that byte in the A register. At the end of each segment, it is called once more with -1 in A. The decompressor output is whatever is output by the OUT instruction.

The context model (called HCOMP) is always present. It is called once per decoded byte. It puts its result in H. OUT has no effect. HCOMP sees as input the PCOMP code (if present) followed by a contiguous stream of segments with no separator symbols.

The ZPAQ program is a development environment for writing and debugging models. It allows programs to be single stepped or run separate from compression. It accepts control statements IF/IFNOT-ELSE-ENDIF and DO-WHILE/UNTIL/FOREVER and converts them to conditional jumps. It allows passing of numeric arguments and comments in parenthesis. If a C++ compiler is present, then ZPAQL code is compiled by converting it to C++ and then running it. Otherwise the code is interpreted. Compiling makes compression and decompression 2 to 4 times faster.

The default configuration for both ZPAQ and ZPIPE is described by the file mid.cfg below.

```
(zpaq 1.07+ config file tuned for average compression.
Uses 111 * 2^$1 MB memory, where $1 is the first argument.)
```

```
comp 3 3 0 0 8 (hh hm ph pm n)
0 icm 5 (order 0...5 chain)
1 isse 13 0
2 isse $1+17 1
3 isse $1+18 2
4 isse $1+18 3
5 isse $1+19 4
6 match $1+22 $1+24 (order 7)
```

```

7 mix 16 0 7 24 255 (order 1)
hcomp
c++ *c=a b=c a=0 (save in rotating buffer M)
d= 1 hash *d=a (orders 1...5 for isse)
b-- d++ hash *d=a
b-- d++ hash *d=a
b-- d++ hash *d=a
b-- d++ hash *d=a
b-- d++ hash b-- hash *d=a (order 7 for match)
d++ a=*c a<= 8 *d=a (order 1 for mix)
halt
post
0
end

```

The comment about \$1 means that the model can be run with additional memory to improve compression. For example:

```
zpaq ocmid.cfg archive files...
```

will compress with 111 MB memory, and

```
zpaq ocmid.cfg,3 archive files...
```

will compress with $111 * 2^3 = 888$ MB memory. Decompression requires the same amount. The effect of ",3" is to make substitutions like "\$1+17" with 20 throughout the configuration file. Up to 9 parameters (to \$9) are allowed.

The command "oc" means optimize (compile the ZPAQL into C++) and compress. If the "o" is dropped then no external C++ compiler is required, but compression and decompression takes twice as long.

The configuration file is divided into 3 sections. COMP describes the arrangement of components. HCOMP contains ZPAQL code that computes the contexts and puts them in H. POST 0 indicates that there is no post-processing.

COMP is followed by 5 arguments: hh, hm, ph, pm, n. hh and hm specify the sizes of H and M in HCOMP as powers of 2 ($2^3 = 8$ each). ph and pm are 0 because these arrays are not used. (Their size is actually 1). n = 8 is the number of components. They must be numbered 0 through n-1 in the COMP section. Except for the line numbers, each token compiles to one byte of ZPAQL. (Thus, ZPAQ requires "A= 10" be written exactly like this and not "A=10" or "A = 10" to indicate it is a 2 byte instruction).

The line

```
0 icm 5
```

describes an indirect context model with a table size of $64 * 2^5$ bytes. It takes its context from the low 15 bits of H[0]. The low 7 bits index a table of 16 byte arrays, and the next 8 bits are the checksum to detect collisions. Since this is an order 0 context, H[0] is left at 0 and only the bitwise context (a 9 bit value) is used. The line

```
1 isse 13 0
```

describes an indirect SSE using $64 * 2^{13}$ bytes taking its input from component number 0 and context from H[1]. The line

```
2 isse $1+17 1
```

describes an indirect SSE using $64 * 2^{\$1+17}$ bytes, where \$1 is the argument passed to mid.cfg. For example, if the argument is 3, it uses $64 * 2^{20}$ bytes. \$1 defaults to 0. It gets its input prediction from component 1 and its context from H[2]. The line

```
6 match $1+22 $1+24
```

specifies a match model with a hash table size of 2^{22} and history buffer of size 2^{24} (taking 16 MB each if \$1 is 0). Its context is H[6]. The line

```
7 mix 16 0 7 24 255
```

specifies a mixer with 2^{16} sets of weights selected by the low 16 bits of the context, taking as input predictions from the range of components 0 through 7-1, with a learning rate of 24/4096, and no masking (AND the bitwise context with 255). The context is H[7].

The HCOMP section computes the contexts and puts them in H. It puts order 0 through 5 context hashes in H[0] through H[5], an order 7 context in H[6] for the match model, and an unhashed order 1 context in bits 8..15 of H[7] for the mixer. It leaves bits 0..7 clear because the bitwise context will be added to this. This is not a concern for the other contexts because they are hashed.

HCOMP is called once after modeling each byte with the byte in the A register. All state information except A and PC (which is reset to the first instruction) is preserved between calls.

This program uses M as a rotating history buffer of 8 bytes (hm = 3) with the low 3 bits of C pointing to the last byte stored. It uses B as a working pointer to compute hashes and D as a pointer into H to store the result. The instructions

```
c++ *c=a b=c a=0
```

increment C, store the input byte in M[C], copy C to B, and clear A.

```
d= 1 hash *d=a
```

assigns 1 to D so that it points to H[1]. The hash instruction takes input from M[B] and combines it with A (0), so A now contains a hash of the last input byte. It is stored in H[D] = H[1] as an order 1 context for component 1. Subsequent instructions store order 2, 3, 4, 5, and 7 hashes in H[2] through H[6]. Note that the space in "d= 1" is required because it is a 2 byte instruction. "a=0" doesn't require this because there is a special 1 byte instruction for clearing a register.

```
d++ a=*c a<<= 8 *d=a
```

computes the mixer context by putting the input byte (saved in *C) into bits 8..15 of D[7] and leaving the other bits at 0. Execution ends at the halt instruction.

The following is the configuration max.cfg, which gets better compression but is slower.

```
(zpaq 1.07+ config file tuned for high compression (slow)
Uses 245 x 2^$1 MB memory, where $1 is the first argument.
```

```
comp 5 9 0 0 22 (hh hm ph pm n)
0 const 160
1 icm 5 (orders 0-6)
2 isse 13 1 (sizebits j)
3 isse $1+16 2
4 isse $1+18 3
5 isse $1+19 4
6 isse $1+19 5
7 isse $1+20 6
8 match $1+22 $1+24
9 icm $1+17 (order 0 word)
10 isse $1+19 9 (order 1 word)
11 icm 13 (sparse with gaps 1-3)
12 icm 13
13 icm 13
14 icm 14 (pic)
15 mix 16 0 15 24 255 (mix orders 1 and 0)
16 mix 8 0 16 10 255 (including last mixer)
17 mix2 0 15 16 24 0 (order -1)
18 sse 8 17 32 255 (order 0)
19 mix2 8 17 18 16 255 (order 0)
20 sse 16 19 32 255 (order 1)
```

```

21 mix2 0 19 20 16 0 (order -1)
hcomp
c++ *c=a b=c a=0 (save in rotating buffer)
d= 2 hash *d=a b-- (orders 1,2,3,4,5,7)
d++ hash *d=a b--
d++ hash *d=a b--
d++ hash *d=a b--
d++ hash *d=a b--
d++ hash b-- hash *d=a b--
d++ hash *d=a b-- (match, order 8)
d++ a=*c a&~ 32 (case insensitive words)
a> 64 if
  a< 91 if (if a-z)
    d++ hashd d-- (update order 1 word hash)
    *d<>a a+=*d a*= 20 *d=a (order 0 word hash)
    jmp 9
  endif
endif
  (else not a letter)
  a=*d a== 0 ifnot (move word order 0 to 1)
  d++ *d=a d--
  endif
  *d=0 (clear order 0 word hash)
(end else)
d++
d++ b=c b-- a=0 hash *d=a (sparse 2)
d++ b-- a=0 hash *d=a (sparse 3)
d++ b-- a=0 hash *d=a (sparse 4)
d++ a=b a-= 212 b=a a=0 hash
  *d=a b<>a a-= 216 b<>a a=*b a&= 60 hashd (pic)
d++ a=*c a<= 9 *d=a (mix)
d++
d++
d++ d++
d++ *d=a (sse)
halt
post
0
end

```

The COMP section begins with an ISSE chain of orders 0 through 6 like mid.cfg (with one extra ISSE). "const 160" provides a bias for the mixers that follow later. It outputs a fixed prediction of $(160-128)/16 = 2$ (stretched). The match model is order 8. As with mid.cfg, M is used as a rotating history buffer, but with a size of $2^{hm} = 2^9 = 512$. H is $2^{hh} = 32$ elements. There are $n = 22$ components.

Components 9 and 10 are an ISSE chain of word-oriented order 0 and 1 contexts for modeling text. These form a separate chain. Generally, the best compression is obtained when each ISSE context contains the lower order context of its input. Otherwise the models should be independent and mixed later. The context is formed by mapping upper and lower case letters together and discarding all other input. The order 0 context is a hash of all of the letters since the beginning of the word or 0 if the last byte was not a letter. The order 1 hash combines this with the previous word.

Components 11 through 13 are sparse order 1 contexts with a gap of 1 to 3 bytes between the context and the current byte. These are useful for modeling binary files.

Component 14 is a model for CCITT binary fax images (PIC in the Calgary corpus). The image width is 1728 pixels or 216 bytes, mapped one bit per pixel in MSB to LSB order (0=white, 1=black). The context is the 8 bits from the previous scan line and 2 additional bits from the second to last scan line.

Components 15 and 16 are order 1 and 0 mixers taking all earlier components as inputs. The second (order 0) mixer has a slower learning rate because it has fewer free parameters. Those two mixer outputs are mixed by the context free (size 0) MIX2 at 17. Its output is refined by the order 0 SSE at 18. The input and output of the SSE are mixed at 19. That prediction is refined by the order 1 SSE at 20. Finally the input and output of that SSE are mixed by the context free MIX2 at 21.

The code for computing the order 0 and 1 word contexts in H[9..10] starts at

```
d++ a=*c a&~ 32 (lowercase words)
```

This increments D to point to H[9] (the order 0 word model), retrieves the input byte saved in M[C], and clears bit 5 (meaning a $\&= \sim 32$) which converts lower case to upper case. Then

```
a> 64 if
a< 91 if (if a-z)
```

tests if A is in the range A..Z (ASCII 65..90). The "if" is converted to a conditional jump to the matching "else" or "endif". If the test passes then the two word hashes are updated. The instruction $*d<>a$ means swap H[D] with A. The hash in D[9] is a rolling hash but in D[10] is cumulative. JMP 9 skips 9 bytes past the commented "else" clause. If the input is not a letter then H[9] is moved to H[10] and H[9] is cleared.

The following results are for the Calgary corpus as a solid archive when supported. Compression is timed in seconds on a 2 GHz T3200.

Program	Size	Time	Memory (MB)
-----	-----	----	-----
zip -9	1,020,831	0.6	0.5
ppmd	804,992	0.6	7.5 (calgary.tar)
ppmd -m256 -o16	754,243	1.3	62 (calgary.tar)
ppmonstr	669,497	8	51 (calgary.tar)
lpaq1 6	682,512	8	195 (calgary.tar)
lpaq9m 6	686,161	8	198 (calgary.tar)
paq9a -6	676,914	12	209
zpaq ocmid.cfg	699,474	8	111
zpaq ocmax.cfg	644,433	20	246
zpaq ocmax.cfg,1	644,320	20	477
paq8l -6	595,474	368	435
paq8px_v67 -6	598,969	469	421

4.3.8. Crinkler

[Crinkler](#) is a compressing linker written by Rune L. H. Stubbe and Ashe Simon Christensen in 2010 for producing small, self extracting x86 executable programs. It is optimized for producing demos, or programs limited to 4096 bytes. It works like [UPX](#), extracting to memory and then running itself. But for this application, it is important to use a much better compression algorithm and that the decompression code be very small, in this case 210 to 220 bytes of size-optimized x86-32 assembler. Compression speed is much less important.

As [explained](#) by Alexandre Mutel, Crinkler is based on PAQ1 ([linear mixing](#)) with sparse models and static weights tuned by the compressor. The model is transmitted as a string of bytes, each representing one context over the last 8 bytes, for example, 11100000b = 0xE0 for an order 3 context, or 01100000b = 0x60 for a sparse order 2 context with a 1 byte gap. Model weights are powers of 2. Up to 32 model weights are delta coded and packed into 4 bytes, such that a 1 bit means that the weight is twice the previous weight and a 0 means it is the same.

Contexts are mapped to a pair of 8-bit counters in a hash table. To make the decompressor as small as possible, there is no hash table collision detection. To compensate, the hash table is made very large, usually 256 MB or 512 MB. Also, there is no counter overflow detection. This saves 6 bytes of code.

The compressor searches model space to find the best compression. It tries adding one model at a time, iterating in bit mirrored order (00000000, 10000000, 01000000...) and keeping a model if it improves compression. Each time a model is added, it tries removing existing models one at a time to see if compression improves further.

Crinkler uses additional tricks to save space prior to compression. It is a size-optimized linker. Thus, it can reorder code and data segments to group similar segments together. It eliminates unnecessary data that some linkers include, such as debugging symbols, and it replaces the DLL symbol table with a small piece of custom code. It compresses the code and data segments separately and applies a [E8E9](#) transform (replacing relative JMP and CALL operands with absolute addresses) to the code segment.

5. Transforms

A transform converts data into a sequence of symbols which can be compressed with a simpler or faster model, or one using less memory, such as an order 0 context model. Those symbols still need to be modeled and coded as described in sections 4 and 3 respectively.

A transform should ideally be a [bijection](#). For each input, there should be exactly one possible representation. More often, the transform is an injection, and its inverse a surjection. An input may have more than one valid representation, either of which is transformed back to the original data during decompression. This increases the information content of the transformed data because the arbitrary choice of representation has to be modeled and coded, taking up space.

5.1. Run Length Encoding

A [run length code](#) replaces a long repeating sequence of identical symbols with two symbols representing a count and the value to be repeated. For example, the string "AAAAA" would be coded as (5,"A").

5.2. LZ77 and Deduplication

In [LZ77](#) compression, duplicate strings in the input are replaced by pointers to previous occurrences. LZ77 is not a bijection. For example, given the string:

AB . . BC . . ABC

"ABC" could be coded as:

- a pointer to AB, and literal C,
- a literal A and pointer to BC,
- or 3 literals, A, B, C.

A pointer consists of an offset and a length. It is allowed for the copied string to overlap the output. For example "AAAAA" could be coded as a A,(-1,4) meaning write a literal "A" and then go back 1 and copy the next 4 bytes. Thus, LZ77 may also encode run lengths.

LZ77 decompression is extremely fast, faster than compression. The compressor must search for matching strings, typically using a hash table or tree. The decompressor only needs to maintain an output buffer from which to copy repeated strings, and then write a copy of its output to the buffer.

The name "LZ77" comes from Lempel and Ziv, who described it in a 1977 paper (Ziv and Lempel, 1977).

5.2.1 LZSS

[LZSS](#) (Lempel-Ziv-Storer-Szymanski, 1982) uses 1 bit flags to mark whether the next symbol is a literal or a pointer. LZSS is used in [NTFS](#) file compression in Windows when the folder property is set to store files compressed. Its primary goal is to be extremely fast (faster than disk access) rather than provide good compression. The exact format was not published. Rather, it was [reverse engineered](#) (in Russian) in 1998. 16 literal/pointer flags are packed into 2 bytes. This is followed by 16 symbols which are either 1 byte literals or 2 byte pointers. The offset is variable length with a maximum of 12 bits. Any remaining bits are allocated to the length, which has a minimum value of 3. Thus, after 2K of input, each pointer is a 12 bit offset and a 4 bit length ranging from 3 to 18.

Windows indicates that a compressed folder containing the Calgary corpus occupies 1,916,928 bytes. On the large text benchmark, the 1 GB text file enwik9 compresses to 636 MB, slightly larger than an order 0 coder and about twice the size of zip. Copying enwik9 between 2 uncompressed folders takes 41 seconds on the test machine (a laptop with a 2.0 GHz T3200). Copying from a compressed folder to an uncompressed folder takes 35 seconds, i.e. decompression is faster than copying. Copying from an uncompressed folder to a compressed folder takes 51 seconds. This is equivalent to compressing the Calgary corpus in 0.03 seconds over the time to copy it.

The NTFS implementation of LZSS is very similar to [lzw1-a](#) implemented by Ross Williams in 1991. lzw1-a uses a fixed 12 bit offset and 4 bit length.

5.2.2. Deflate

The widely popular [deflate](#) format is used in zip and gzip and is supported by many other archivers. It is used internally in PNG images, PDF documents, and Java JAR archives. The format is documented in [RFC 1951](#) (1996) and supported by the open source [zlib](#) library.

In the deflate format, pointer offsets range from 1 to 32768 and length from 3 to 258. Literals and lengths are coded in a 286 symbol alphabet which is Huffman coded followed by up to 5 extra uncompressed bits of the length. A length code is followed by an offset from a 30 symbol Huffman coded alphabet followed by up to 13 extra uncompressed bits. Specifically the alphabets are as follows:

```
0..255 = literal byte
256 = end of data
257..264 = lengths 3..10
265..268 = lengths 11..18 followed by 1 extra bit
269..272 = lengths 19..34, 2 extra bits
273..276 = lengths 35..66, 3 extra bits
277..280 = lengths 67..130, 4 extra bits
281..284 = lengths 131..257, 5 extra bits
285 = length 258
```

Lengths are followed by an offset coded from a 30 symbol alphabet:

```
0..3 = offset 1..4
4..5 = offset 5..8 followed by 1 extra bit
6..7 = offset 9..16, 2 extra bits
8..9 = offset 17..32, 3 extra bits
...
28..29 = offset 16385..32768, 13 extra bits
```

The format allows either a default or a custom Huffman code. The default code lengths are as follows:

```
Literal/length
0..143 = 8 bits
144..255 = 9 bits
256..279 = 7 bits
280..287 = 8 bits

Offset
0..29 = 5 bits
```

If a custom Huffman table is used, then the table is transmitted as a sequence of code lengths. That sequence is itself compressed by run length encoding using another Huffman code to encode the literals and run lengths. It uses a 19 symbol alphabet:

```
0..15 = code lengths of 0..15
16 = copy the previous code 3..6 times, followed by 2 extra bits
17 = copy 3..10 times, 3 extra bits
18 = copy 11..138 times, 7 extra bits
```

The Huffman table for these codes are sent as a sequence of up to 19 3-bit numbers. This sequence is further compressed by reordering the sequence so that the values most likely to be 0 (not used) are at the end, and sending the sequence only up to the last nonzero value. A 4 bit number indicates the sequence length. The order is: 16, 17, 18, 0, 8, 7, 9, 6, 10, 5, 11, 4, 12, 3, 13, 2, 14, 1, 15. All Huffman codes are packed in LSB to MSB order.

zip and gzip take an option -1 through -9 to select compression level at the expense of speed. All options produce compressed data in deflate format which decompresses at the same speed (much faster than compression) with the same algorithm. The difference is that with the higher options, the compressor spends more time looking for encodings that compress better. A typical implementation will keep a list of 3 byte matches (the shortest allowed) in a hash table and test the following data to find the longest match. With a higher option, the compressor will spend more time searching. It is also sometimes possible to improve compression by encoding a literal even if a match is found, if it results in a longer match starting at the next byte. Such testing

also increases compression time. [kzip](#) performs an extreme level of optimizations like this. Compressed sizes and compression times on a 2.0 GHz T3200 are shown below for the 14 file Calgary corpus.

Program	Size	Time
-----	-----	-----
zip -1	1,194,720	.17 sec.
zip -2	1,151,711	.23
zip -3	1,115,078	.25
zip -4	1,072,909	.25
zip -5	1,041,083	.33
zip -6	1,028,171	.40 (default)
zip -7	1,025,244	.42
zip -8	1,021,607	.50
zip -9	1,020,831	.67
kzip	978,707	24.21
unzip		.10

5.2.3. LZMA

[LZMA](#) (Lempel Ziv Markov Algorithm) is the native compression mode of [7-zip](#). Although 7-zip is open source, the algorithm was never well documented or described until it was [analyzed](#) by Charles Bloom in Aug. 2010. Compression is improved over LZ77 by using a longer history buffer (up to 4 GB), optimal parsing, shorter codes for recently repeated matches (like LZX), literal exclusion after matches, and arithmetic coding.

Optimal parsing. Simple LZ77 uses greedy parsing. The next byte is always coded using the longest match found in the history buffer. 7-zip uses look-ahead to consider a shorter match or a literal if it allows a longer match to be coded afterward with a shorter total code length.

Matias and Sahinalp (1999) proved that for any dictionary encoding method with the prefix property, that there is a fast ($O(n)$) algorithm that optimally parses the input into the minimum number of phrases by looking only at the lengths of the next two phrases considered. A dictionary has the prefix property if for any phrase in it, all of its prefixes are also in the dictionary. LZ77 and LZW both have this property. However, optimal parsing is more difficult to achieve in practice because phrases may have different code lengths, so that the minimum number of phrases does not necessarily give the best compression.

Repeated matches. 7-zip uses shorter codes (like LZX) to code the 3 most recent matches in place of the LZ77 offset. In the case of the most recent match, the match length can be as small as 1.

Literal exclusion after matches. 7-zip uses a "delta literal" to code the first literal after the match. The idea is that the literal to be coded is known to be different than the predicted byte following the match in the history buffer. An ordinary context model would incorrectly assign a high probability to the predicted byte, which would waste code space. To prevent this, the coded literal is exclusive-ored with the predicted byte.

Coding. 7-zip uses a binary arithmetic coder. It first codes one of 4 possible choices: literal, match, rep (repeat of one of the 3 last matches) or short rep (repeat the last match with a length of 1). A literal code is followed by the 8 bits of the coded byte in MSB to LSB order. A match is coded as a length followed by the offset, each using a variable length code in which the length of the binary value is transmitted followed by the value. A rep is followed by a length and a 1 or 2 byte code indicating the position in a move-to-front queue of the last 3 offsets. A short rep (SRep) code is complete.

Context modeling. Literals are coded using an order 1 model plus the low bits of the current position. The position is useful for modeling binary files that have a repeating structure that is a power of 2, for example, an array of integers or floating point numbers. An order 1 model means that the context for each bit is the 8 bits of the previous byte and the 0 to 7 previously coded bits of the current byte.

Match offsets are coded by transmitting the binary length of the offset and then the offset in LSB to MSB order. The length of the offset is coded using the match length as context. The context for the 4 low bits is the 0 to 3 previously coded bits. The higher order bits are not compressed.

The match/literal flags (literal, match, rep, short rep) are coded using the low bits of the position as context (as with literals) plus the encoding state. The state is one of 12 values and is updated by the following state table:

Next state					Meaning of current state
State	Lit	Mat	Rep	SRep	
0	0	7	8	9	literal after 2 or more literals
1	0	7	8	9	literal after match, literal
2	0	7	8	9	literal after rep, literal or after non-literal, short rep, literal
3	0	7	8	9	literal after literal, short rep, literal
4	1	7	8	9	(delta) literal after match
5	2	7	8	9	(delta) literal after rep or after non-literal, short rep
6	3	7	8	9	(delta) literal after literal, short rep
7	4	10	11	11	match after literal
8	5	10	11	11	rep after literal
9	6	10	11	11	short rep after literal
10	4	10	11	11	match after non-literal
11	5	10	11	11	rep or short rep after non-literal

5.2.4. LZX

[LZX](#) is an LZ77 variant used in Microsoft [CAB](#) files and compressed help (CHM) files. It uses a history buffer of up to 2 MB and Huffman coding. To improve compression, it uses shorter codes to code the 3 most recent matches as with LZMA.

5.2.5. ROLZ and LZF

The idea of using shorter codes for recent matches can be extended. The compressor for [lzw3](#) builds a dictionary (a hash table) of pointers into the history buffer as usual to find matching strings, but instead of coding the offset, it codes the index into the table. The decompressor builds an identical hash table from the data it has already decompressed, then uses the index to find the match. The length is coded as usual.

[ROLZ](#) (reduced offset LZ) extends this idea further by replacing the large hash table with many smaller hash tables selected by a low order context. This reduces the size of the offset, although it can sometimes cause the best match to be missed. ROLZ was implemented in WinRK.

The extreme case of ROLZ is to use one element per hash table. In this case, only a literal or length must be coded. This algorithm is called [LZF](#). It was first described by Charles Bloom in 1995. LZF works best with a high order context. Thus, it is often used as a preprocessor to a low or moderate order context model, rather than a fast order 0 model like LZ77.

5.2.6. Snappy

[snappy](#) is a free, open source (Apache) compression library from Google in 2011. It uses byte aligned LZ77 intended for high speed (especially on x86-64 hardware) rather than good compression. Google uses snappy internally to compress its data structures for its search engine.

The compressed data contains tag bytes such that the low 2 bits indicate literals and matches as follows:

```
00 = literal
01 = 1 byte match
10 = 2 byte match
11 = 4 byte match (not used)
```

A literal of length 1 to 60 is encoded by storing the length - 1 in the upper 6 bits. Longer literals are coded by storing 60.63 in the upper 6 bits to indicate that the length - 1 is encoded in the next 1 to 4 bytes in little-endian (LSB first) format. This is followed by the uncompressed literals.

Matches of length 4 to 11 with offsets of 1 to 2047 are encoded using a 1 byte match. The match length - 4 is stored in the middle 3 bits of the tag byte. The most significant 3 bits of the offset are stored in the upper 3 bits of the tag byte. The lower 8 bits of the offset are stored in the next byte. A match may overlap the area to be copied. Thus, the string "abababa" could be written using a literal "ab" and a match with an offset of 2 and length of 5. This would be encoded as:

```

000001 00 (literal of length 2)
01100001 (literal 'a')
01100010 (literal 'b')
000 001 01 (high bits of offset, match of length 5)
00000010 (low 8 bits of offset)

```

Matches of length 1 to 64 with offsets of 1 through 65535 are encoded using a 2 byte match. The length - 1 is encoded in the high 6 bits of the tag byte. The offset is stored in the next 2 bytes with the least significant byte first. Longer matches are broken into matches of length 64 or less.

A 4 byte match allows offsets up to $2^{32} - 1$ to be encoded as with a 2 byte match. The decompressor will decode them but the compressor does not produce them because the input is compressed in 32K blocks such that a match does not span a block boundary.

The entire sequence of matches and literals is preceded by the uncompressed length up to $2^{32} - 1$ written in base 128, LSB first, using 1 to 5 digits in the low 7 bits. The high bit is 1 to indicate that more digits follow.

Compression searches for matches by comparing a hash of the 4 current bytes with previous occurrences of the same hash earlier in the 32K block. The hash function interprets the 4 bytes as a 32 bit value, LSB first, multiplies by 0x1e35a7bd, and shifts out the low bits. The hash table size is the smallest power of 2 in the range 256 to 16384 that is at least as large as the input string. This is a speed optimization to reduce the time to clear the hash table for small inputs.

As an optimization for hard to compress data, after 32 failures to find a match, the compressor checks only every second location in the input for the next 32 tests, then every third for the next 32 tests, and so on. When it finds a match, it goes back to testing every location.

As another optimization for the x86-64 architecture, copies of 16 bytes or less are done using two 64-byte assignments rather than memcpy(). To support this, if 15 or fewer bytes remain after a match then they are encoded as literals with no further search.

5.2.7. Deduplication

Deduplication is the application of LZ77 to a file system rather than a data stream. The idea is to find duplicate files or files containing large blocks of data duplicated elsewhere, and replace them with pointers. Deduplication can save disk space, but is more commonly used for incremental backups where only data that has been changed is saved.

Whole file deduplication is easy to implement. If a file system enforces a "last modified" timestamp or an "archive" flag, then files that have not been modified since the last backup need not be saved. Otherwise, or for on-disk deduplication, a secure hash such as [SHA-1](#) of each file can be computed. If two files have the same 160 bit hash, then it is safe to assume that the contents are identical. It is not impossible for different files to have the same hash, but the probability of a collision, 2^{-160} is far smaller than the 2^{-60} probability of a disk I/O error. Furthermore, it is believed to be computationally infeasible to deliberately produce a collision.

Compression can be improved by dividing large files into smaller blocks and searching for duplicates among them. This can be done efficiently by computing a hash of each block and saving pointers to the blocks in a hash table. To save space, it might be desirable to use a smaller but insecure hash and check for potential collisions by comparing the blocks.

Dividing large files into fixed sized blocks (typically 8 KB to 64 KB) will not find duplicates between versions in which a small amount of data was inserted or deleted near the beginning because the data afterward would have moved relative to the block boundaries. This problem can be solved by selecting boundaries that depend on a [rolling hash function](#) over a sliding window. We set a boundary whenever n bits of the hash are equal to a fixed value. The average block size will be 2^n . For example, the following rolling hash function uses a sliding window of 32 bytes and marks a boundary on average every 16 KB when the high 14 bits of the hash are set to 1.

```

uint32_t h = 0; // rolling hash value
int c; // input byte
const int K = 876543210; // a random 32 bit even number not a multiple of 4

```

```

while ((c = getchar()) != EOF) {
    h = (h + c + 1) * K;
    if (h >= 0xffffc0000)
        mark_boundary();
}

```

5.3. LZW and Dictionary Encoding

Dictionary methods substitute codes for common strings from a table or dictionary. A dictionary code may be, fixed, static or dynamic. In the fixed case, the dictionary is specified as part of the algorithm. In the static case, the compressor analyzes the input, constructs a dictionary, and transmits it to the decompressor. In the dynamic case, both the compressor and decompressor construct identical dictionaries from past data using identical algorithms.

5.3.1. LZW

[LZW](#) (Lempel-Ziv-Welch) is a dynamic dictionary method. It is used by the UNIX [compress](#) program, [GIF](#) images, and is one of the compressed modes in [TIFF](#) images. The algorithm was patented by both Sperry (later Unisys) in 1981 and by IBM in 1983 when the USPTO did not realize that they were the same algorithm. Unisys was criticized for waiting until GIF became an established standard before demanding royalties from makers of software that could read or write GIF images. Both patents are now expired.

LZW starts with a dictionary of 256 1-byte symbols. It parses the input into the longest possible strings that match a dictionary entry, then replaces the string with its index. After each encoding, that string plus the byte that follows it is added to the dictionary. For example, if the input is ABCABCABCABC then the encoding is as follows:

```

65 = A   (add AB to dictionary as code 256)
66 = B   (add BC as 257)
67 = C   (add CA as 258)
256 = AB (add ABC as 259)
258 = CA (add CAB as 260)
257 = BC (add BCA 261)
259 = ABC (end of input)

```

Dictionary codes grow in length as it becomes larger. When the size is 257 to 512, each code has 9 bits. When it is 513 to 1024, each code is 10 bits, and so on. When the dictionary is full (64K = 16 bits), it is discarded and re-initialized.

A Windows version of the UNIX compress program compresses the Calgary corpus to 14 files totaling 1,272,722 bytes in 0.34 seconds and decompresses in 0.23 seconds.

Other variants of LZW may use larger dictionaries, or may use other replacement strategies like LRU (least recently used), or other strategies for adding new symbols such as concatenating the last two coded symbols instead of a symbol plus the next byte.

5.3.2. Dictionary Encoding

Dictionary encoding improves the compression of text files by replacing whole words with symbols ranging from 1 to 3 bytes. Fixed English dictionaries are used in WinRK, durilca, and in some versions of PAQ such as PAsQDacc 4.3c -7, which compresses the Calgary corpus to 567,668 using a dictionary extracted from the corpus itself, but not included in the compressed size. It is, of course, possible to compress to arbitrarily small sizes using this technique. The extreme case is [barf](#). It has a built in 14 word dictionary, one for each file of the Calgary corpus. When the compressor detects a match, it "compresses" the file to 1 byte, which the decompressor correctly expands.

For this reason, the [large text benchmark](#) and contests like the [Calgary challenge](#) and [Hutter prize](#) include the size of the decompression program and all other files needed to decompress. Still, it may be useful to use a dictionary for one or more languages if the input is expected to contain text in those languages.

Of more interest are static dictionaries. These are used by the top 3 programs on the large text benchmark (as of Feb. 2010), and in all of the Hutter prize winners. Some of the later Calgary challenge winners also use small dictionaries.

5.3.2.1. Modeling Text

Recall from section 1.4 that text compression is an AI problem. This can be seen by playing Shannon's character guessing game which he used to estimate that the entropy of written English is about 1 bpc (Shannon, 1950). Try partially covering some text with your hand and guessing what letters come next from the earlier context, for example: "the cat caught a mo____". Humans can beat computers at the game because the prediction problem requires vast understanding of English and of the world. Nevertheless, some of the constraints of natural language can be modeled. These rules are categorized as follows:

- lexical: "moqse" is wrong because it is not a word.
- semantic: "moose" is wrong because it is not associated with things that a cat would chase.
- syntactic or grammatical: "moves" is wrong because we expect "a" to be followed by a singular noun phrase.

While playing the game, you will notice that useful contexts start on word boundaries. Thus, "a mo_" and "caught a mo_" are useful contexts, but "ght a mo_" is no more useful than the lower order "a mo_". Thus, text models in PAQ construct contexts that start on word boundaries.

It should be irrelevant if a context spans a line break. Thus, word contexts in PAQ discard the characters between words. Furthermore, it should be irrelevant if the context is upper or lower case, because it does not change the meaning. Thus, there are word models that ignore case.

Semantics can be modeled by associating each pair of words like (cat, mouse) with a co-occurrence frequency over a small window. Words that frequently occur near each other tend to have related meanings. This can be modeled with a sparse order-1 word model, skipping one or more words in between the context and the predicted word. Many PAQ versions have sparse word models with small gaps of 1 to 3 words.

For modeling semantics, it is useful to split text into "meaningful" units or [morphemes](#) if possible. For example, "moves" really has two independent parts, the stem "move" and suffix "s". Ideally these should be modeled as separate words.

Grammar constrains text to make certain sequences more likely, such as (the NOUN) or (a ADJ NOUN). It is possible to learn the parts of speech by observing when words occur in similar contexts and grouping them. For example "the dog", "the cat", and "a dog" could be used to predict the unseen sequence "a cat". This works by grouping "the" and "a" into one semantic category and "dog" and "cat" into another category.

A dictionary transform works by replacing the input text with a sequence of highly predictive symbols corresponding to morphemes, independent of capitalization and punctuation. This improves compression both by allowing simpler models to be used, and by reducing the size of the input, which improves speed and reduces pressure on memory. Compression can be improved further by grouping semantically or grammatically related words so that the compressor can merge them into single contexts by ignoring the low bits of dictionary codes. Care should be taken not to remove useful context, which can happen if a dictionary is too large or divides words in the wrong places. For example, coding "move" and "remove" as unrelated symbols would not help compression.

5.3.2.2. Capitalization Modeling

A capitalization model replaces upper case letters with a special symbol followed by a lower case letter. For example, "The" could be coded as "Athe" where "A" directs the decompressor to capitalize the next letter. Alternatively, a more sophisticated model might automatically capitalize the first letter after a period, and insert a special symbol to encode exceptions.

5.3.2.3. Newline Modeling

Because newlines are semantically equivalent to spaces, it is sometimes useful to replace them with spaces, and encode in a separate stream the information to put them back. A simple transform is space stuffing, where a space is inserted after every newline. For example, this has the effect of merging the order 4 contexts " the" and "\nthe" (where \n is a linefeed) by replacing the latter with "\n the". Space stuffing does not help with multi-word contexts that span lines. However the alternative is to remove context that could predict newlines, such as periods at the end of a paragraph.

5.3.2.4. Word Encoding

Word encoding is done in two passes. In the first pass, the text is parsed into words (sequences of A-Z, a-z, and possibly UTF-8 characters in non-English alphabets) and counted. Words with counts below a threshold are discarded. In the second pass, words found in the dictionary are replaced with 1 or 2 byte codes (or 3 bytes for large dictionaries). The dictionary is listed at the beginning of the output, followed by the encoded data. Words not found in the dictionary and non-letters are passed unchanged.

Words are typically encoded with bytes from the part of the ASCII set that does not normally appear in text, namely 0..8, 11..12, 14..31, and 127..255. If capitalization modeling was done, then 65..90 (A-Z) may also be used. If such bytes do appear, they must be preceded by an escape byte, designated as one of the above. The remainder of the alphabet may be used to encode words.

[XWRT](#) (XML Word Reducing Transform) by Przemyslaw Skibinski in Oct. 2007 performs dictionary encoding. The dictionary is appended as a header to the output with one word per line in decreasing order of frequency. The most frequent words are encoded with one byte.

5.3.2.5. Results

The following table shows the effect of simple capitalization modeling (using "A" followed by lowercase), space stuffing, and word encoding using XWRT on book1 from the Calgary corpus on 4 compressors. The -f option selects the minimum word frequency for inclusion in the dictionary. The number in parenthesis shows the dictionary size that results. (The exact options are -o -l0 -c -s -n -w -m256 for xwrt 3.2 to turn off other transform options. There is no space or capitalization modeling).

	book1	zip -9	sr2	bzip2 -9	ppmonstr
No transform	768,771	312,502	276,364	232,598	203,247
Capitalization	785,101	<u>311,696</u>	275,124	231,594	202,650
Space stuffing	785,393	313,640	275,161	229,988	202,274
Both	801,723	312,856	<u>273,864</u>	<u>229,861</u>	<u>201,706</u>
xwrt -f3 (4307)	366,323	<u>265,721</u>	246,897	233,928	211,382
xwrt -f6 (2806)	378,289	267,522	<u>246,760</u>	231,675	208,801
xwrt -f20 (789)	449,233	278,639	254,227	230,243	204,897
xwrt -f100 (174)	542,670	290,268	262,575	<u>228,904</u>	<u>202,832</u>

The table shows that space stuffing and capitalization usually help, but that word encoding becomes less effective as the compression improves. It is nevertheless useful for compressing the large text benchmark where memory constraints are severe because it reduces the size of the input. The top 3 programs use it. Capitalization modeling is also useful, but space stuffing is not because line breaks are only used to separate paragraphs.

XWRT is ranked sixth (as of Feb. 2010) on the large text benchmark when used with its built in LPAQ6 compressor. The optimal setting in this case is -f200 to select a dictionary size of about 40,000 words. It does slightly better (ranking fifth) as a preprocessor to ppmonstr with option -f64.

paq8hp12 and drtlpaq9m, both by Alexander Rhatushnyak, are ranked second and third on the large text benchmark and are the basis of winning entries for the Hutter prize. These both use a custom dictionary of about 44,000 words. The higher frequency words are grouped semantically, such as "son" with "daughter" and "monday" with "tuesday". Among the lower frequency words, they are grouped by common suffix (alphabetical order when reversed) to make the dictionary compress smaller.

durilca_kingsize is the top ranked program on the large text benchmark, but only because it uses 13 GB of memory, vs. 2 GB. It uses a dictionary of about 124,000 words. These are also grouped semantically, but by an automated process that clustered words in context space. The algorithm was not documented, but the idea is roughly to group words together if they are likely to appear in the same context.

5.4. Symbol Ranking

Symbol ranking, or move-to-front (MTF), is a transform in which the alphabet is maintained in a queue from newest to oldest and coded by its position. The idea is that the most recently seen symbol is the most likely to occur in the future.

[srank](#) is a symbol ranking compressor by Peter Fenwick in 1996. An order 3 context hash without collision detection is mapped to a queue of length 3 representing the 3 most frequently seen bytes in that context. These are Huffman coded with 1, 3, or 4 bits respectively. Long runs of first place bytes are run length encoded with 12 bits to encode the run length. Bytes not seen in the queue are modeled in an order 0 pseudo-MTF queue using 7 bits for the first 32 positions and 12 bits for the other 224. It is called "pseudo-MTF" because when a byte is observed, it is moved only about half way to the front with some random dithering. This is an optimization for speed. It allows a fast update of an index into the queue. The order 3 hash table maximum size is 2^{18} queues (1 MB memory).

[sr2](#) is an improved (but slower) symbol ranking compressor by Matt Mahoney in Aug. 2007. An order 4 context hash is mapped to a table of 2^{20} 3 byte MTF queues and a counter for consecutive first place hits ranging from 0 to 63. When the first place byte is observed, the counter is incremented. For all other values, the counter is reset to 1 if in the queue or 0 if not. The new value is pushed to the front of the queue and the others pushed back. For example, the sequence ABCCC would result in the queue (C,B,A,3) with C at the front. A subsequent B would result in (B,C,A,1). A subsequent D would result in (D,B,C,0).

A byte is first Huffman coded and then arithmetic coded. The point of the Huffman code is to reduce the number of arithmetic coding operations for better speed. Suppose the queue contains (c1,c2,c3,n). The coding and next state is as follows:

Input	Code	(c1 c2 c3 n) next state
Initial	----	-----
		(0, 0, 0, 0)
c1	0	(c1, c2, c3, min(n+1, 63))
c2	110	(c2, c1, c3, 1)
c3	111	(c3, c1, c2, 1)
other c	10cccccccc	(c, c1, c2, 0)

The bits are coded using a direct context model with a count ranging from 2 to 128 (section 4.1.2). For $n \geq 4$, the context is order 1 plus n plus the previous bits of the current symbol. For $n < 4$, the model is the same except order 2.

The following comparison is for the Calgary corpus as 14 files compressed separately.

Program	Size	Compr	Decompression
-----	-----	-----	-----
srank -C8	1,281,984	0.20	0.20 sec.
sr2	975,208	0.48	0.50 sec.

During development, it was observed that an order 3 context sometimes compressed better on smaller files, but order 4 works better on larger files. Increasing the hash table beyond 2^{20} did not help much, in spite of the fact that more memory almost always helps any algorithm.

Arithmetic decoding is slightly slower than encoding. Recall that the steps to compress are:

1. predict next bit
2. code the bit
3. update the model

For decompression:

1. predict next bit
2. decode the bit
3. update the model

Modern processes can reorder instructions and execute them in parallel. During compression, steps 2 and 3 are independent, so they can overlap. During decompression, the model cannot be updated until the bit has finished being decoded.

5.5. Context Sorting (BWT)

A [Burrows-Wheeler Transform](#) sorts the input by its right context. By bringing together characters with similar contexts, the transformed data can be more easily compressed with a fast adapting order 0 model. Shown below is a portion of the Burrows Wheeler transform of book1 from the Calgary corpus (with newlines converted to spaces for clarity). The column in bold is the BWT.

BWT block --+ +--- Sorted on this column

\ /
VV

```
ing. Her culpability lay in her m
e of the instability of a woman ?
hat the desirability of her exist
tion, from inability to guide inc
husband's inability to meet the
nervous excitability, he returned
stimony, probability, induction -
le of respectability, were now si
of a ship's cabin, with wood slid
new riding-habit -- the most ele
mostly her habit hen excited, he
's virtuous habit of entering the
new riding-habit of myrtle-green
aracter and habit, and seemed so
e no riding-habit, looked around
besides the habitable inn itself,
ceived no inhabitant for the spac
' by the inhabitants of Caster+
ttle <P 61> habitation, and here
every human habitation, and the h
those old-inhabited walls. It was
ntly to old habits and usages, si
o imply his habitual reception of
rgrass, who habitually spoke on a
at everybody abjured her -- for w
osed all the able-bodied men upon
en the favourable-con+ junction s
ame to the stable-door and looked
d been answerable .! ' " We must
```

The BWT is the column in **bold**. It is "...ptrnntbtchhhhhhhhhhhhhh rtr...".

BWT is best suited for stationary sources. For example, a sorted list of words would be compressed poorly because local rules (newline is followed by "A", later changing to "B") become spread throughout the transform. For these cases, separating different data types into independently compressed blocks can improve compression. Otherwise, the largest possible block size should be used.

BWT compression depends on the alphabet order. Best compression is obtained when related symbols such as letters or digits that make similar predictions are grouped together. The ASCII character set already has this property, but is not optimal.

5.5.1. Forward transform

Compressors that use BWT are called context sorting or block sorting algorithms. A typical implementation is to divide the input into fixed sized blocks (as large as memory allows) and sort an array (of the same size) of pointers into the block. The sort order is the lexicographical order of the string to which it points, wrapping

around to the beginning of the block if necessary. Wrapping around can be avoided by appending a virtual "end of block" symbol which is lexicographically less than any of the symbols in the alphabet (e.g. -1).

A straightforward sort using a fast sorting algorithm like a [radix sort](#) or [quicksort](#) can be very slow on highly redundant or compressible data. Quicksort requires on average $O(n \log n)$ string comparisons, but worst case string comparison is $O(n)$, as in the case of a block of identical bytes. This would result in $O(n^2 \log n)$ run time. There are several ways to avoid the problem of long common prefixes.

- Flip a small fraction of the bits using a pseudo-random algorithm to break up long strings. This hurts compression slightly. Stuffed (sitx format) uses this technique.
- Preprocess the data using a high order LZP or LZ77 preprocessor.
- Use a [Schindler transform](#), a variation in which the stable sort order is based on a truncated string comparison. This method is faster but hurts compression because sorting occurs on a bounded context. Also, this method is protected by [patent 6,199,064](#) in the U.S. The patent expires Nov. 14, 2017.
- Use one of many suffix array sorting algorithms, some of which are as fast as $O(n)$ worst case.


There are fast suffix sorting algorithms that use 5 or 6 times the block size ($5n$ to $6n$) in memory. One of the earliest is a suffix tree sorting algorithm developed by Ukkonen (1995), which is $O(n)$ but requires $13n$ - $15n$ memory. The idea is to build a tree in which each path from root node to a leaf node represents one of the n suffixes of the block. Each edge is labeled with a string of one or more bytes (represented as a pair of pointers into the block) Each of the n leaf nodes is a pointer to the start of the suffix. Each internal node except possibly the root node is required to have two or more children (in lexicographical order from left to right). This limits the number of nodes to at most $2n$. The algorithm is to build the tree and then output the sorted suffix pointers by a recursive inorder traversal from the root node:

```

traverse(node):
  if the node is a leaf then
    output block[leaf-1]
  else
    for each child from left to right
      traverse(child).

```

A suffix tree for the string "BANANA" with end of block symbol "\$" is shown below. The downward arrows are called transitions. Each transition is labeled with a string of one or more bytes. The dotted lines are called suffix links. They form a path from the longest suffix represented by a nonterminal node (having a child \$) to successively shorter suffixes back to the root (which represents the empty suffix). The suffix links are maintained to make the tree building algorithm efficient.



A suffix tree of "BANANA\$" (from [Wikipedia](#)).

Ukkonen's algorithm is to start with just a root node and scan the block from left to right in a single pass. For each byte, the suffix pointers form a linked list of nodes that need to be updated. An update consists of appending a child node, splitting edges as needed, and updating the suffix pointers. The suffix pointer traverse stops when there is already a child node for the character being added.

A [suffix array](#) represents a suffix tree using just the block, the sorted list of pointers, and the longest common prefix (LCP), which is the number of initial bytes shared with the previous suffix after sorting. The suffix array of "BANANA\$" is shown:

Index	Suffix	LCP
6	\$	0
5	A\$	0
3	ANA\$	1
1	ANANA\$	3
0	BANANA\$	0
4	NA\$	0
2	NANA\$	2

Yuta Mori has [benchmarked](#) several suffix array sorting algorithms. The fastest are his own [libdivsufsort](#), [MSufSort](#), used in [M99 and M03](#) by Michael Maniscalco, and [archon4r0](#) used in the [dark](#) archiver. All of the sorting algorithms are open source.

5.5.2. Inverse transform

It is rather surprising that a BWT block can be inverted to recover the original data. The only other information needed is the new position of the original first byte. We are given a BWT string $BWT[0..n-1]$ of length n , and the location p ($0 \leq p < n$) of the position of the first byte $BWT[p]$. The algorithm to output the original string is:

```
T = sort(BWT)
Repeat n times:
    output BWT[p]
    move p from the i'th location of c in T to the i'th location of c in BWT
```

For example, suppose that $BWT[0..5] = "NNBAAA"$ is the BWT of "BANANA". as shown. (To simplify the explanation, we omit the \$, but it works the same either way).

```
BWT Sorted context
  \ /
  NABANA
  NANABA
  BANANA p=2
  ABANAN
  ANABAN
  ANANAB
```

Create $T[0..5] = "AAABNN"$. We now have:

```
      p
    0 1 2 3 4 5
BWT = N N B A A A
T    = A A A B N N
```

The steps are:

```
output BWT[2] = B
p is the third A in T. Move to p=5, the third A in BWT.
output BWT[5] = A
p is the second N in T. Move to p=1, the second N in BWT.
output BWT[1] = N
p is the second A in T. Move to p=4, the second A in BWT.
output BWT[4] = A
p is the first N in T. Move to p=0, the first N in BWT.
output BWT[0] = N
p is the first A in T. Move to p=3, the first A in BWT.
output BWT[3] = A.
```

As an optimization, we may represent the sorted array T solely by the starting position of each of the 256 sequences of byte values, for example $A=0$, $B=3$, $C=3, \dots$, $N=4$, $O=6$. Furthermore, we can construct in advance a list $NEXT[0..n-1]$ such that $NEXT[p]$ is the next move for p . For example, $NEXT[2] = 5$ would be the first move. To build this list we scan BWT and use T to count the occurrence of each byte value.

```
for i in 0..n-1 do
    NEXT[T[BWT[i]]++] = i
```

In C++, the inverse BWT looks like this:

```
// Invert and output the BWT in bwt[0...n-1] starting at p
void invert_BWT(unsigned char *bwt, int n, int p) {

    // Collect cumulative counts of bwt:
    // t[i] = number of bytes < i
    int t[257]={0}; // cumulative counts
    for (int i=0; i<n; ++i)
        ++t[bwt[i]+1];
    for (int i=1; i<257; ++i)
        t[i]+=t[i-1];
    assert(t[256]==n);

    // Build linked list
    int *next=calloc(n, sizeof(int)); // linked list
    assert(next); // out of memory?
    for (int i=0; i<n; ++i)
        next[t[bwt[i]]++]=i;
    assert(t[255]==n);

    // Traverse and output list
    for (int i=0; i<n; ++i) {
        assert(p>=0 && p<n);
        putc(bwt[p], out);
        p=next[p];
    }
    free(next);
}
```

5.5.3. bzip2

[bzip2](#) is a popular open source BWT based file compressor developed in 1996 by Julian Seward. It takes an option -1 through -9 to select a block size of 100 KB to 900 KB. -9 generally gives the best compression. The compression algorithm is as follows:

1. The input is run length encoded to remove some (not all) high order redundancy. Sequences of 4 to 255 repeated bytes are coded as the first 4 bytes followed by one byte (0..251) representing the remaining count. For example, "AAAAA" is coded as "AAAA",1.
2. BWT.
3. Move to front (section 5.4). Each byte is coded as its position in a queue, then moved to the front of the queue. Runs of identical characters thus become runs of zeros.
4. Run length encoding of zeros. The run length is coded in binary in LSB to MSB order by two symbols (RUNA, RUNB) that have values 1 and 2 (instead of 0 and 1). Runs of length 1 through 10 would be coded as 1, 2, 11, 21, 12, 22, 111, 211, 121, 221.
5. The symbols RUNA, RUNB, queue positions 1..255 (0 is always run length encoded) and end of data symbol are Huffman coded.

bzip2 uses 2 to 6 Huffman tables, which are selected every 50 symbols to make the code adaptive. The tables are kept in a MTF queue. The selection code is unary coded. A unary code for a number n is n 1 bits and a 0. For example, 4 = 11110.

The Huffman tables are coded as a sequence of lengths. The lengths are delta coded, i.e. as the difference from the previous length. A difference is coded as 0 = 0, 10 = -1, 11 = +1, repeating as needed. A bitmap is used to mark unused queue selection codes, which are omitted from the sequence. The bitmap is divided into 16 16-bit words, where a 0 bit means the code is not used. If all 16 bits are 0, then the word is omitted. One additional 16 bit word is used to mark which words are omitted.

The original bzip was arithmetic coded, which is better suited for a fast adapting model. It was replaced with a Huffman code due to patents (now expired) on arithmetic coding.

5.5.4. BBB

[BBB](#) (big block BWT) is an open source file compressor written by Matt Mahoney in Aug. 2006. It has two innovations: a "slow" mode that allows blocks up to 80% of available memory, and a context mixing model of the BWT sequence. When released, it was top ranked on the large text benchmark among BWT compressors because it was the only program that could fit the 1 GB test file into a single block on a 2 GB machine.

5.5.4.1. Slow Mode BWT

To context sort a large block, it is first divided into 16 smaller blocks of 4-byte pointers which are sorted with wraparound using `std::stable_sort()` (normally a generic comparison sort such as [quicksort](#)). The pointers are then written to 16 temporary files, which are merged to produce the final result.


The inverse transform does not build a linked list, because this takes 4 times as much memory as the block size. Recall that the inverse transform first sorts the bytes in the block into an array *T* (represented by 256 cumulative counts), and that *p* points to the next output byte in the block. The steps to be iterated are:

1. output BWT[*p*]
2. if *T*[*p*] is the *i*'th occurrence of *c* in *T*, then set *p* to the *i*'th occurrence of *c* in BWT.

Normally, step 2 is done by traversing a link in the list NEXT. Instead, BBB searches the block for the *i*'th occurrence of *c*. To do this efficiently, it first consults an index that locates every 16'th occurrence of *c* in BWT to get the approximate location, and searches linearly from there. This index takes 1/4 as much memory as the BWT block.

5.5.4.2. Modeling

BBB uses an order 0 indirect context model (section 4.1.3) followed by 6 SSE stages and bitwise arithmetic coding. The model uses 5 MB of memory. It looks like this:



The ICM takes a bitwise order 0 context, i.e. just the previous bits of the current byte. Recall that an ICM maps a context to a bit history (an 8 bit state), which is mapped to a slow adapting probability table.

Each of the SSE (section 4.3.3) maps a context and a probability (stretched and quantized to 32 levels) to a new probability interpolated between the two nearest quantized outputs. SSE1 and SSE2 are both order 0, but SSE1 is fast adapting (learning rate 1/32) and SSE2 is slow adapting (learning rate 1/512). The two predictions are averaged in the linear domain. The SSE in BBB update both quantized table entries above and below the input probability, unlike ZPAQ which updates only the nearest.

SSE3 takes a bitwise order 1 context. SSE4 takes the previous but not the current byte as context, plus the run length quantized to 4 levels (0, 1, 2..3, 4+).

SSE5 takes a sparse order 1 context of just the low 5 bits and a gap of 1 byte, i.e. 5 of the last 16 bits, plus the current byte: ...xxxxx The output is averaged linearly with the input with weight 3/4 to the output.

SSE6 takes a 14 bit hash of the order 3 context. It is averaged linearly with its input with weight 1/2 each.

The table below compares the models for bzip2 and BBB on the Calgary corpus with each file compressed separately. In both cases the block size is 900 KB, which is large enough to hold each file in a single block. BBB is run in both fast and slow modes. About half of the compression time in both cases is due to PIC, which has long runs of 0 bytes. Unlike bzip2, BBB has no protection against long string comparisons while sorting.

Note also that compressing all of the data together as a tar file makes compression worse. As mentioned, BWT is poorly suited for mixed data types.

Program	Calgary	Compr	Decomp (seconds)	calgary.tar
-----	-----	-----	-----	-----
bzip2 -9	828,347	0.68	0.42	860,097
bbb cfk900	785,672	10.33	1.46 (fast mode)	800,762
bbb ck900	785,672	13.74	5.54 (slow mode)	

5.5.5. MSufSort v2

Most modern BWT based compressors use some type of suffix array sorting algorithm for the forward BWT transform. MSufSort version 2 is a fast and memory efficient (about 6n block size) algorithm developed by Michael Maniscalco (Puglisi, 2005; Puglisi, Maniscalco, 2006a). Later an improved version, MSufSort v3, used a different algorithm.

Recall that a suffix array (SA) is a list of pointers to the suffixes of a string in lexicographical order. Thus, the SA for "BANANA\$" is (6,5,3,1,0,4,2), where "\$" precedes all other characters. The BWT is calculated by outputting the characters to the left of each of the pointers, e.g. "ANNB\$AA".

MSufSort v2 actually calculates the inverse suffix array (ISA). The ISA lists the lexicographical rank of each character in the string. For example, the ISA of "BANANA\$" is (4,3,6,2,5,1,0). There is a simple algorithm for converting an ISA to an SA:

```
for each i in 0..n-1 do
  SA[ISA[i]] = i
```

However, because SA and ISA each require 4n memory, the inversion is typically done in place. This can be achieved by using one bit of each element (such as the sign bit) to mark those that have already been moved (Puglisi, 2005).

```
for i from 0 to n-1 do
  if ISA[i] > 0
    start := i
    suff := i
    rank := ISA[i] - 1
    repeat
      tmp := ISA[rank]
      ISA[rank] = -suff (use sign bit to mark as moved)
      suff := rank
      rank := tmp
    until rank = start
    ISA[rank] := -suff
```

SA sorting based on comparison sorts have worst case run times of $O(n^2)$ or worse because of the $O(n)$ time required to compare strings when the average LCP is large, as in the case of highly redundant data. A common feature of efficient algorithms, such as the 15 described by (Puglisi, Smyth, Turpin, 2007), is that they exploit the fact that if two suffixes have a common prefix of m characters and their order is known, then we also know the order of their first m suffixes. For example, once we know that "ANA\$" comes before "ANANA\$", then we also know the order of their next m = 3 suffixes, specifically that "NA\$" comes before "NANA\$", "A\$" comes before "ANA\$", and "\$" comes before "NA\$". Suffix trees implicitly exploit this feature by merging the m common characters into a single edge, thus achieving $O(n)$ sort times, although at a high cost in memory.

MSufSort v2 is faster in practice than a suffix tree, and uses less memory, which allows for better compression because blocks can be larger. However its run time is difficult to analyze because of its complexity.

MSufSort v2 begins by partially sorting the suffixes by their first two characters using a counting sort and storing the result as a stack of 64K linked lists reading right to left in the input string. To save space, the list pointers are overlayed on the ISA, and the head of each list is stored in a separate stack which is popped in lexicographical order. If a list contains only one element, then it is replaced by its rank in the ISA. The sign bit is used to distinguish ranks from pointers. If a list has more than one element, then that list is split into 64K new lists by a counting sort keyed on the next two characters, and the head of each list is pushed on the stack maintaining

lexicographical order. Each stack entry has a head pointer and a suffix length. ~ denotes the last element in a list. For example:

i	0	1	2	3	4	5	6		Stack
	B	A	N	A	N	A	\$	Input string	
ISA[i]	~	~	~	1	2	~		Initial chains	$(5,2)_{A\$}, (3,2)_{AN}, (0,2)_{BA}, (4,2)_{NA}$
ISA[i]	~	~	~	1	2	-1		Rank singleton A\$	$(3,2)_{AN}, (0,2)_{BA}, (4,2)_{NA}$
ISA[i]	~	~	~	~	2	-1		Split AN	$(3,4)_{ANA\$}, (3,4)_{ANAN}, (0,2)_{BA}, (4,2)_{NA}$
ISA[i]	~	~	~	-2	2	-1		Rank singleton ANA\$	$(3,4)_{ANAN}, (0,2)_{BA}, (4,2)_{NA}$
ISA[i]	~	-3	~	-2	2	-1		Rank singleton ANAN	$(0,2)_{BA}, (4,2)_{NA}$
ISA[i]	-4	-3	~	-2	2	-1		Rank singleton BA	$(4,2)_{NA}$
ISA[i]	-4	-3	~	-2	~	-1		Split NA	$(4,4)_{NA\$}, (2,4)_{NANA}$
ISA[i]	-4	-3	~	-2	-5	-1		Rank singleton NA\$	$(2,4)_{NANA}$
ISA[i]	-4	-3	-6	-2	-5	-1		Rank singleton NANA, stack is empty, so done.	

There are two important optimizations. First, we note that during the chain splitting, if the rank of a suffix one or two characters to the right is known, then it must precede other suffixes where this is not the case. Those suffixes can also be sorted using those ranks as keys, and don't have to be pushed on the stack. For example, at the point where the chain NA at locations 2 and 4 is sorted, the positions at 3 and 5 have already been ranked (-2 and -1 respectively). Thus, we already know from those ranks that position 5 must come first. This optimization is known as *induced sorting*.

This optimization is insufficient to prevent worst case $O(n^2)$ behavior on data with high average LCP. To handle this, MSufSort v2 detects repeating strings (like "ANAN") and applies the observation that if the suffix order of the last repeat is known, then the same order can be applied to all of the remaining prefixes and they can be written into the ISA. Once those ranks are applied, other rotations of the repeating string (like "NANA") are ranked by induced sorting.

5.5.6. Itoh-Tanaka

The open source [dark](#) compressor by Dmitry Malyshev uses Itoh-Tanaka suffix array sorting (Puglisi, Smyth, Turpin, 2007, sec. 3.3) for the forward BWT transform. Itoh-Tanaka has $O(n^2 \log n)$ worst case complexity, so additional steps such as LZP preprocessing are needed to avoid long common prefixes.

Itoh-Tanaka classifies each input byte by whether or not it is larger than the byte following it. For example:

	0	1	2	3	4	5	6
Input:	B	A	N	A	N	A	\$
Type:	L	S	L	S	L	L	S

We denote a character as type L if it is lexicographically larger than the byte after it and S (smaller) if it is smaller or equal. Itoh-Tanaka sorts the S characters using ordinary string comparison, then induces the sort order of the L bytes in a single pass. The first step is to put the unsorted type S characters into the suffix array, SA. We note that SA has 256 buckets plus the terminating sentinel, and furthermore each bucket is divided into two parts with all of the L bytes on the left and S on the right. Thus:

SA: $\$=\{6\}$ A={- 1 3} B={-} N={- -}

Next we sort the right halves of all of the buckets that have two or more type S symbols. Itoh and Tanaka use a [multikey quicksort](#) although any string sorting algorithm could be used. Recall that the quicksort algorithm picks one element at random, called a pivot, then moves all elements smaller than the pivot to the left side of the array and all others to the right side, with the pivot in between. Then it recursively quicksorts the two halves. In a multikey quicksort, the array is divided into three parts (less, equal, greater) based on the first character of the string. The left and right parts are recursively sorted, and the middle part is recursively sorted using the next character of the string as the sort key, similar to a radix sort. After sorting we have:

SA: $\$=\{6\}$ A={- 3 1} B={-} N={- -}

We now have enough information in SA to induce the rest of the elements. We can use the fact that when the order of two suffixes in SA is known and they are both preceded by the same character, then those two longer

suffixes must also be in the same order. For example, if we know that suffix 6 (\$) comes before 4 (NA\$) and both are preceded by "A", then it must be the case that suffix 5 (A\$) comes before 3 (ANA\$).

The procedure for the induced sort is to scan SA left to right. For each pointer already filled in, we look at the character to its left. If it is type L then we place it in the leftmost available position in the appropriate bucket. For example:

```
SA: $={6} A={- 3 1} B={-} N={- -}
  Scan 6. Input[5] = 'A' is type L, so we put 5 in bucket A.
SA: $={6} A={5 3 1} B={-} N={- -}
  Scan 5. Input[4] = 'N' is type L, so we put 4 at the left of bucket N.
SA: $={6} A={5 3 1} B={-} N={4 -}
  Scan 3. Input[2] = 'N' is type L, so we put 2 at the next open spot in bucket N.
SA: $={6} A={5 3 1} B={-} N={4 2}
  Scan 1. Input[0] = 'B' is type L, so we put 0 at the first open spot in bucket B.
SA: $={6} A={5 3 1} B={0} N={4 2}
  Scan 0. There is no character to the left, so continue.
  Scan 4. Input[3] is type S so continue.
  Scan 2. Input[1] is type S so continue.
```

We now have the complete suffix array, (6,5,3,1,0,4,2). Besides the input array and SA, the only additional memory needed is an array of 256 pointers to the first available position in each bucket.

5.5.7. DivSufSort

[LibDivSufSort](#) is an open source library for suffix array sorting by Yuta Mori. It is used for the forward BWT in [bsc](#) (Ilya Grebnov), [mcomp](#), (the engine for WinRK by Malcolm Taylor), and was used in early versions of [bcm](#) (Ilya Muraviev).

DivSufSort [improves](#) on Ikoh-Tanaka by directly sorting only the last S type suffix in each run of consecutive type S. Then it induces the rest of the S's in a right to left scan, filling each bucket in from the right side. For each element SA[i], if input[i-1] is type S, then its index is put at the rightmost open position in the appropriate bucket (but overwriting the original rightmost S elements). This process is just like the induced sort of Itoh-Kanaka, but in the opposite direction. After this step, the type L suffixes are sorted as in Itoh-Kanaka. This 3 step method was first described by (Ko, Alura, 2003).

DivSufSort would still have $O(n^2 \log n)$ complexity due to the string sorting in the first step. To avoid this problem, it detects and induces tandem repeats like with MSufSort v2.

5.5.8. MSufSort v3

MSufSort version 3 uses a completely different (and faster) suffix array sorting algorithm than version 2 (Puglisi, Maniscalco, 2006b). It uses slightly more than 5n memory. It is used in the [M03](#) compressor.

MSufSort v3 builds the SA rather than the ISA. It first uses a counting sort on the first two bytes to sort into 64K buckets. Then each bucket is sorted using the DivSufSort algorithm with detection and induction of tandem repeats as in MSufSort v2 to avoid delays caused by long common prefixes. Typically about 30% of the suffixes need to be sorted and the rest are induced. Strings are sorted using multikey quicksort (in increasing lexicographical order to support tandem repeat induction). As an optimization, groups smaller than 64 elements are sorted using [insertion sort](#). Also, to control the size of the recursion stack, the sorting algorithm is changed to a [heapsort](#) after a recursion depth of 48.

5.5.9. Bijective BWT

A Burrows-Wheeler transform of a string consists of the characters of that string sorted by context, plus the index of the new position of the original start of the string. This index is needed because the context sorted string is not unique. For example, the BWT of BANANA (without a sentinel) is (NNBAAA,2), and the BWT of its rotation ANANAB is (NNBAAA,5). Thus, the BWT is not a bijection.

Note also that adding the sentinel does not make BWT a bijection either. Although the sentinel makes the context sort unique, it increases the size of the output string.

David A. Scott discovered in 2007 a bijective variant of the BWT (called BWTS for BWT Scottified) which omits the need to transmit the index of the starting string. Because the output of the BWTS is the same size as the input, all strings also have a valid inverse BWTS. The algorithm is described in (Kufleitner, 2009), who also extended the idea to the sort (Schindler) transform.

The idea is to divide the input string into a set of smaller strings such that the starting index of each substring is at a known, fixed location, such as at the first character. Specifically, the input is partitioned into a lexicographically nonincreasing sequence of [Lyndon words](#). A Lyndon word is a string that lexicographically precedes any of its rotations. For example, the Lyndon factorization of BANANA is (B)(AN)(AN)(A). A Lyndon factorization is unique and can be calculated in $O(n)$ time.

Recall that the inverse BWT is to use a counting sort to calculate $T = \text{sort}(\text{BWT})$ and then to build and traverse a linked list from the i 'th occurrence of c in T to the i 'th occurrence of c in BWT. In a BWT, this linked list usually forms a single loop that returns to the initially transmitted index when complete. (A repeating string will have multiple loops but will decode the same way). In the BWTS, the list forms one loop for each Lyndon word. The end of each word is detected by traversing the list back to its start. The next loop starts at the first unused element. The Lyndon words are restored right to left. For example, we compute the BWTS of "BANANA" by sorting the rotations of its 4 Lyndon words, and taking the last character of the sorted column as the BWTS:

	B		B			
B	AN		AN	A		A
A Lyndon		rotations	NA	sort	AN	last character
N	----->	----->	----->	AN	----->	N
A	AN		AN	B		B
N			NA	NA		A
A			NA	NA		A
	A		A			

(When sorting, string comparison wraps around. Thus, B would fall between BA and BC).

To compute the inverse BWTS, we construct a linked list just like with a BWT.

i	T	BWTS	Next
-	-	----	----
0	A	A	0
1	A	N	4
2	A	N	5
3	B	B	3
4	N	A	1
5	N	A	2

The linked list has 4 cycles: (0), (1,4), (2,5), (3). Reversing the order of the cycles and concatenating, we get (3,2,5,1,4,0). The corresponding elements of T spell out "BANANA".

A BWTS is 4 bytes smaller than a BWT. Experimentally, the effect on compression is small. The following table shows the effect of a BWT and BWTS followed by encoding using an adaptive order 0 coder (fpaq0p, learning rate 1/32), and an indirect order 0 model (fpaq0f2). For BWT, MSufSort 3.1.1 was used. It adds a 4 byte starting index to the beginning. For BWTS, the [BWTS](#) program by Scott was used with BLOCKSIZE set to 4000000 in the source code. In both cases, the input file is sorted in a single block. The test files are calgary.tar (mixed data types) and book1 from the Calgary corpus (uniform data).

File	Encoder	BWT	BWTS
-----	-----	-----	-----
book1	fpaq0p	244,106	244,096
calgary.tar	fpaq0p	985,906	985,885
book1	fpaq0f2	220,338	220,344
calgary.tar	fpaq0f2	845,526	845,549

One possible benefit of a bijective compression algorithm is for encryption. Because all possible decryptions are valid inputs to the decompressor, it eliminates one possible test that an attacker could use to check whether a

guessed key is correct. (However, one should not assume that there are not other tests, and should still choose strong keys). Scott has also written a bijective arithmetic coder, [arb255](#). Thus, it is possible to make the entire compression algorithm a bijection.

5.6. Predictive Filtering

A predictive filter is a transform which can be used to compress numeric data such as audio, images, or video. The idea is to predict the next sample, and then encode the difference (the error) with an order 0 model. The decompressor makes the same sequence of predictions and adds them to the decoded prediction errors. Better predictions lead to smaller errors, which generally compress better.

5.6.1. Delta Coding

The simplest predictive filter is a [delta code](#). The predicted value is just the previous value. For example, the sequence (5,6,7,9,8) would be delta coded as (5,1,1,2,-1). A second pass would result in (5,-4,0,1,-3).

Delta coding works well on sampled waveforms containing only low frequencies (relative to the sampling rate), such as blurry images or low sounds. Delta coding computes a discrete derivative. Consider what happens in the frequency domain. A [discrete Fourier transform](#) represents the data as a sum of sine waves of different frequencies and phases. In the case of a sine wave with frequency ω radians per sample and amplitude A , the derivative is another sine wave with the same frequency and amplitude ωA . From the [Nyquist theorem](#), the highest frequency that can be represented by a sampled waveform is π or half the sampling rate. Frequencies above 1 radian per sample will increase in amplitude after delta coding, and lower frequencies will decrease. Thus, if high frequencies are absent, it should be possible in theory to reduce the amplitude to arbitrarily small values by repeated delta coding.

Eventually this fails because any noise (which is not compressible) in the prediction is added to noise in the sample with each pass. (Noise has a uniformly distributed spectrum, so its high frequency components are amplified by delta coding). Noise can come either from the original data or from quantization (rounding) errors during sampling. These are opposing sources. Decreasing the number of quantization levels removes noise from the original data but adds quantization noise.


The images below show the effects of 3 passes of delta coding horizontally and vertically of the image [lena.bmp](#) (a widely used benchmark image). The original image is in BMP format, which consists of a 54 byte header and a 512 by 512 array of pixels, scanned in rows starting at the bottom left. Each pixel is 3 bytes with the numbers 0..255 representing the brightness of the blue, green, and red components. The image is delta coded by subtracting the pixel value to the left of the same color, and again on the result by subtracting the pixel value below. (The order of the two encodings does not matter). To show the effects better, 128 is added to all pixel values (which does not affect compression). Thus, a pixel equal to its neighbors appears medium gray.




Original image



Delta coded once horizontally and vertically.



Delta coded twice.



Delta coded 3 times.

The original image is 786,486 bytes (with or without delta coding). The following table shows the compressed sizes when compressed with an order 0 indirect context model (ICM-0), with each of the 3 colors compressed in a separate stream.

	ICM-0

lena.bmp	569,299
delta 1	511,316
delta 2	645,634
delta 3	768,154

For comparison, [Image Magick](#) compresses to PNG format with size 474,573, and the top ranked paq8px_v67 -6 to 412,641 bytes.

Details: The ICM-0 model was implemented in ZPAQ 1.10 using the following configuration:

```
comp 0 0 0 0 1
  0 icm 7 (indirect context model using 27+6 bytes)
hcomp
  b++ a=b a== 3 if (b is 0,1,2 depending on color)
    a=0 b=0
  endif
  a<= 9 *d=a halt (context is color in bits 10..9 + order 0)
post
  0
end
```

5.6.2. Color Transform

lena.bmp can be compressed to 499,139 bytes by performing the color transform (red, green, blue) to (red-green/4, green, blue-green*3/4), then delta coding and modeling with ICM-0. The transform was tuned to this image, but is not optimal for all images. For many others, the transform (red-green, green, blue-green) works well. The transform works because when one color is brighter, the others tend to be too. Thus, one color can predict the others. The ideal transform depend on the average color of the image.

5.6.3. Linear Filtering

A linear filter is a [finite impulse response](#) filter with n taps that predicts a sample x_i in the sequence $x_1x_2\dots$ as follows:

$$\text{prediction} = \sum_{j=1..n} w_j x_{i-j}$$

where w_j is called the j 'th coefficient. A delta filter is the special case of the coefficient array $n = 1$, $\mathbf{w} = (1)$. Two passes of a delta filter is equivalent to $n = 2$, $\mathbf{w} = (2, -1)$, and 3 passes is equivalent to $n = 3$, $\mathbf{w} = (3, -3, 1)$.

An [adaptive filter](#) is a linear filter whose coefficients are adjusted to reduce prediction errors. A simple update rule is:

$$w_j := w_j + \lambda x_{j-i}(x_i - \text{prediction}), j = 1..n$$

where λ is the learning rate. The update rule is unstable because of a positive feedback loop: when the error is large, it can lead to large updates which could increase the error even more. An adaptive filter must compensate by limiting the magnitudes of the weights and updates.

5.7. Specialized Transforms

It is often possible to find transforms that improve compression for specialized data types. We mention two.

5.7.1. E8E9

The E8E9 transform is used to compress x86 executable code (EXE or DLL files). In x86, a CALL or JMP instruction (E8 or E9 hex) is followed by a 4 byte address (LSB first) relative to the program counter. Compression can be improved by converting to an absolute address, because code often contains many calls or jumps to the same address. The transform consists of searching for a byte with the value E8 or E9 hex (232 or 233), interpreting the next 4 bytes as a 32 bit number, and adding the offset from the beginning of the input file. The decompressor does the same, except that it subtracts the offset. E8E9 is used in CAB format (for CALL instructions only) and in many top end compressors. Recent versions of PAQ8PX by Jan Ondrus also transform conditional branch addresses.

In x86-64, all references to static memory (not just JMP and CALL) are relative addresses. Currently transforms for x86-64 are not yet widely used.

5.7.2. Precomp

[Precomp](#) is a program by Christian Schnaader that searches its input for segments of data compressed in deflate (zip) format and uncompresses them. This can improve compression if the (now larger) data is compressed with a better algorithm. Many applications and formats use deflate compression internally, including PDF, PNG, JAR (Java archive), ODT (OpenOffice) and SWF (Shockwave Flash).

To make the inverse transform bitwise identical, precomp tests by recompressing the data with [zlib](#) and comparing it. Recall that LZ77 is not a bijection. There are many different ways to compress a string that will decompress the same way. Precomp relies on the fact that most applications use zlib rather than write their own implementation. Still, precomp must test 81 combinations of options to find one that compresses to exactly the original data, and then stores those options in the output. If it fails to find a match (even in valid deflate data), then it must insert additional data.

Precomp can be used by itself. It is also built into two compressors, lprepaq (precomp+lpaq6) and paq808pre (precomp+paq808). paq808pre -7 compresses [flashmx.pdf](#) from the Maximum Compression corpus to 1,821,939 bytes in 692 seconds. As of Feb. 2010 the program has not yet been benchmarked and the best result is 3,549,197 bytes by WinRK 3.1.2. The improvement is obtained partly by unzipping many embedded BMP images and compressing them with paq808's specialized BMP model (which is also top ranked on [rafale.bmp](#)).

5.8. Huffman coding

The open source file compressor [M1x2 v0.6](#) by Christopher Mattern in Feb. 2010 uses order 1 Huffman coding as a preprocessor to a context mixing model. The idea is to reduce the size of the input to make compression faster. The model contexts are aligned on Huffman code boundaries instead of byte boundaries. The order 1 coder is actually 256 tables selected by the previous byte. Huffman codes are limited to 12 bits in length to simplify the implementation.

6. Lossy Compression

Lossy compression refers to discarding unimportant information. Generally this means compressing images, video, or audio by discarding data that the human perceptual system cannot see or hear.

Lossy compression is a hard AI problem. To illustrate, speech could theoretically be compressed by transcribing it into text and compressing it with standard techniques to about 10 bits per second. We are nowhere near that.

Even worse, we could imagine a lossy video compressor translating a movie into a script, and the decompressor reading the script and creating a new movie with different details but close enough so that the average person watching both movies one after the other would not notice any differences. We may use a result by [Landauer](#) (1986) to estimate just how tiny this script could be. He tested people's memory (over a period of days) over a wide range of formats such as words, numbers, pictures and music, and concluded that the human brain writes to long term memory at a fairly constant rate of about 2 bits per second. Currently we need 10^7 bits per second to store DVD quality MPEG-2 video.

The state of the art is to apply lossy compression only at a very low level of human sensory modeling, where the model is well understood.

6.1. Image Compression

All image formats, even BMP, may be regarded as a form of lossy image compression. An uncompressed image is normally a 2 dimensional array of pixels, where each pixel has 3 color components (red, green, blue) represented as an integer with a fixed range and resolution. A pixel array is an approximation of a 2 dimensional continuous field where the light intensity at any point would be properly described as a continuous spectrum. Note how lossy compression is applied:

- The eye can't see detail much smaller than 0.1 mm, so there is no need for an image to have more than a few thousand pixels in each dimension.
- The eye can't detect differences in brightness of less than about 1%, so there is no need to quantize brightness to more than a few hundred levels.
- The eye has 3 types of cones sensitive to red, green, and blue. Combinations of these colors can reproduce every color that the eye can see. There is no need to distinguish pure spectral yellow emitted by a rainbow from the apparent yellow from a monitor produced from a mixture of red and green light, even though there are instruments such as a spectrograph that can make such distinctions.
- The eye detects brightness on a logarithmic scale, so there is no need to use more bits to represent brighter lights. Sunlight is 1000 times brighter than room light, but doesn't look like it.

6.1.1. BMP

A [BMP](#) image uses 8 bits per pixel per color, which matches the resolution of most monitors. Each value is an 8 bit number ranging from 0 (darkest) to 255 (brightest). The values are proportional to apparent light intensity, not

actual intensity. The actual intensity is [gamma corrected](#) by the monitor by raising it to the power of $\gamma = 2.2$. Thus, a pixel value of 200 is a little over 4 times as bright as a pixel value of 100, although it only looks twice as bright.

6.1.2. GIF

The [GIF](#) image format is lossless except that it uses a color palette of up to 256 colors. The format is an array of 8 bit indexes into the palette. The limited number of colors noticeably reduces the quality of color photographs, although it is sufficient for grayscale or diagrams. A GIF file may contain multiple images for animations.

GIF uses LZW compression (section 5.3.1) with a maximum dictionary size of 4K. When the table is full, it is discarded and re-initialized. It reserves two codes to initialize the table and to mark the end of data.

Use of GIF was discouraged due to a patent on LZW, which is now expired.

6.1.3. PNG

[PNG](#) is a lossless image format. Images are normally 8 bits per pixel but can be more. A pixel has 3 color components and an optional fourth component for an alpha channel to indicate transparency.

PNG is compressed by predictive filtering (section 5.6) followed by deflate (section 5.2.2). There are 5 filters which can be selected for each scan line. The image is scanned left to right starting at the top. Let A, B, and C be the previously coded neighboring pixels of the predicted pixel x:

```

C B
A x

```

The 5 possible predicted values are 0, A, B, $(A+B)/2$, or Paeth. The Paeth filter is to predict A, B, or C, whichever is closest to $A + B - C$. The Paeth filter usually gives the best compression.

6.1.4. TIFF

[TIFF](#) is an image container format. Most commonly it is used for uncompressed images when BMP cannot be used because more than 8 bits per pixel or more than 3 color components are needed. TIFF supports several lossless compression modes. The most common is delta coding (subtracting the pixel to the left) followed by LZW.

6.1.5. JPEG

[JPEG](#) is the most widely used representation for photographic images. It uses lossy compression. It exploits two limitation of human visual perception. First, the eye has different degrees of sensitivity to brightness variation depending on spatial frequency, peaking at 0.1 to 0.2 degrees (a few pixels). Second, the eye is much less sensitive to color variation at high spatial frequencies. The compression steps for baseline JPEG are as follows:

- Color transform from RGB to [YCbCr](#).
- Downsampling the two chroma components Cb and Cr.
- 8 by 8 discrete cosine transform (DCT).
- Variable quantization depending on color and spatial frequency.
- Delta coding the DC coefficient.
- Reordering the coefficients in zigzag order from low to high frequency.
- Huffman coding with run length encoding of zeros.

The color transform from RGB (red, green, blue) to YCbCr is:

$$\begin{aligned}
 Y &= 0.299 R + 0.587 G + 0.114 B \text{ (black-white)} \\
 Cb &= 128 - 0.168736 R - 0.331264 G + 0.5 B \text{ (yellow-blue)} \\
 Cr &= 128 + 0.5 R - 0.418688 G - 0.081312 B \text{ (green-red)}
 \end{aligned}$$


The eye is less sensitive to fine detail in Cb and Cr than in Y, so these two are (optionally) downsampled 2 to 1 by averaging 2 by 2 blocks of pixels into 1 pixel. The inverse transform is:

$$\begin{aligned} R &= Y + 1.402(Cr - 128) \\ G &= Y - 0.34414(Cb - 128) - 0.71414(Cr - 128) \\ B &= Y + 1.772(Cb - 128) \end{aligned}$$

The [DCT](#) represents 8 by 8 blocks of pixels in the spatial frequency domain. The 64 DCT coefficients S_{uv} , u, v in $0..7$, of the 8 by 8 pixel block S_{xy} , x, y in $0..7$, are computed as follows:

$$S_{uv} = \alpha(u)\alpha(v) \sum_{x=0..7} \sum_{y=0..7} S_{xy} \cos[\pi/8 (x+1/2) u] \cos[\pi/8 (y+1/2) v]$$

where $\alpha(0) = 1/8^{1/2}$ and $\alpha(1..7) = 1/4$ are normalizing factors. u is the horizontal spatial frequency and v is the vertical spatial frequency. The image below shows the contribution of each of the 64 S_{uv} DCT coefficients to an 8 by 8 pixel block with u reading across and v reading down. The S_{00} coefficient is at the top left.



Left: 8 by 8 DCT (from [Wikipedia](#)). Right: zig-zag encoding order of the S_{uv} coefficients (from [Wikipedia](#)).

The inverse DCT has the same form:

$$S_{xy} = \alpha(x)\alpha(y) \sum_{u=0..7} \sum_{v=0..7} S_{uv} \cos[\pi/8 (u+1/2) x] \cos[\pi/8 (v+1/2) y]$$

Each of the 64 coefficients in Y and the 64 in Cb and Cr are quantized by dividing by one of 128 values from two quantization tables and rounding. Because the eye is less sensitive to high spatial frequencies (u and v large), especially in the two chroma components, these divisors can be larger.

The coefficient S_{00} is the average brightness of the 8 by 8 block. It is called the "DC" coefficient. It is the only one that depends significantly on neighboring blocks, so it is delta coded by subtracting the DC coefficient of the last coded block of the same color. The other 63 coefficients are called "AC".

For most images, the high spatial frequencies will be small except in regions with fine detail. Therefore the coefficients are reordered in zigzag order by increasing $u+v$, resulting in the largest coefficients appearing first.

The coefficients are grouped into runs of R zeros followed by one nonzero value, where R ranges from 0 to 15. The nonzero coefficient is a 12 bit signed number, but is usually near 0. It is coded as an S bit signed number, where S ranges from 1 to 12, followed by S extra bits. For example, the sequence 0,0,0,0,0,3 would be coded as $R=5, S=2, 11$. The RS code (52) would be Huffman coded, and the two bits "11" (binary 3) would follow uncompressed. Negative numbers are coded by subtracting 1 and sending the same number of bits. For example, -3 would be coded as "00", which are the last 2 bits of -4. After the last nonzero coefficient, a RS code of 00 marks the end of block.

There are 4 Huffman tables for the RS codes, one each for DC-Y, DC-color, AC-Y, and AC-color. The tables are transmitted by sending 16 lists of RS codes (1 byte each) having code lengths of 1 through 16. Each list is

preceded by one byte to indicate the length of the list. Other data such as the quantization tables are sent uncompressed. Huffman codes are packed in MSB to LSB order.

JPEG supports inserting restart codes into the Huffman coded data to mark the start of independently compressed image slices for error recovery. There is also an end of image symbol. The Huffman code is designed so that symbols can be found without decoding. Symbols are marked with a starting FF byte (11111111 binary). No symbol is assigned a Huffman code of all 1 bits. Also, if a byte of all 1 bits is coded, then it is followed by a 0 byte which the decoder skips.

The JPEG specification describes several modes in addition to baseline, described above. About 95% of images are baseline JPEG. The rest are mostly progressive mode. In this mode, a coarse approximation of the image is sent first so that the receiver can start displaying it before the rest of it is received. Progressive mode uses two techniques to do this. One is spectral selection, in which the low frequency DCT coefficients are sent first. The other is progressive approximation, in which the high order bits of the coefficients are sent first. Usually both techniques are combined.

JPEG allows up to 4 colors (for an alpha channel) and 12 bits per pixel. These modes are rare, but are supported by the [IJG reference implementation](#). Images may also be grayscale by dropping the Cb and Cr components.

The JPEG standard also describes arithmetic coding as an alternative to Huffman coding, and a lossless hierarchical mode in which successively higher resolution images are sent as differences from the previous one. Neither of these two were implemented by IJG or any subsequent software because the methods were patented at the time.

In 2002, [Forgent](#) claimed U.S. patent 4,698,672 on JPEG, specifically the invention of using a single code to represent a run length followed by a second value, which is used in the RS codes. By April 2004, Forgent announced that it had collected US\$90 million from 30 companies and filed patent infringement suits against 31 others. In May 2006 the USPTO ruled that the claims of the patent related to JPEG were invalid due to prior art. The patent expired 5 months later.

The images below show the effects of the color transform and DCT. The first image was created using [IJG's](#) public domain software `cjpeg` to convert `lena.bmp` (786,486 bytes) to JPEG (23,465 bytes). This is 17.5 times smaller than the best result obtained with lossless compression.




Left: lena.bmp, 786,486 bytes. Right: JPEG created with `cjpeg -quality 50 -optimize`, 23,465 bytes.

The -quality setting sets the quantization tables. These range from 16 for the Y-DC coefficient to 99 for high frequency coefficients. The -optimize parameter creates the best possible Huffman tables.

The images below show a JPEG image separated into its three color components by setting all of the other coefficients to 0. All images below are high quality (-quality 100) without chroma downsampling as above.




Y coefficients only.



Left: Cb only. Right: Cr only.

The images below separate some of the different frequency coefficients.



S_{00} (DC)




First two AC coefficients. Left: S_{01} . Right: S_{10} . Note that these are *color* images.



Left: S_{20} . Right: S_{11} .

The following image illustrates the eye's insensitivity to fine detail in Cb and Cr. All of the 63 AC coefficients in Cb and Cr are set to 0, and yet the effect is barely noticeable. Compare with S_{00} above when the Y AC coefficients are also removed.



All 63 AC chroma coefficients set to 0.

6.1.6. JPEG Recompression

Even though JPEG is already compressed, there are lossless algorithms that can improve the compression by 20% to 30% and decompress to the original JPEG image.

6.1.6.1. Stuffit

On Jan. 7, 2005, Darryl Lovato of Allume (now Smith Micro) [announced](#) that their [Stuffit](#) archiver would compress JPEG files by 30%. Until then, it was widely believed that JPEG files could not be compressed further because they were already compressed. Indeed, most compression algorithms will not compress JPEG files more than about 1 to 2%. On the [ACT benchmark](#) by Jeff Gilchrist, an early version compressed 3 test images by 25% to 30%. The [Maximum Compression JPEG benchmark](#) gives similar results for newer versions.

According to [U.S. Patent 7,502,514](#), filed Jan. 4, 2005, the compressor uses a transform that partially decompresses the JPEG image back to the quantized DCT coefficients, undoing all of the lossless steps. These values are then compressed with a proprietary model. Some details are described in the patent. There are separate models for S_{00} (the DC coefficient), S_{0v} , S_{u0} (DC in one dimension), and S_{uv} for $u, v > 0$. The DC coefficient is modeled with a predictive filter. The decompressor decompresses the coefficients and inverts the transform by repeating the normal lossless JPEG compression steps to restore the original image. JPEG encoding from the DCT coefficients onward is deterministic, so the result is bitwise identical. The patent was issued to Lovato and Yaakov Gringeler. Gringeler was the developer of the (now unsupported) [Compressia](#) archiver.

6.1.6.2. PAQ

PAQ versions beginning with [PAQ7](#) in Dec. 2005 also include a JPEG model for baseline images (but not the less common progressive mode). Both PAQ and Stuffed compress to about the same size, but the PAQ compressor is much slower. PAQ is open source, so its algorithm can be described in better detail. It uses a context model instead of a transform. The context model decodes the image back to the DCT coefficients like the Stuffed algorithm, but instead this information is used as context to predict the Huffman coded data for both compression and decompression without a transform. The PAQ algorithm is not patented.

The PAQ model has evolved over time but all versions work basically the same way. In [paq8px_v67](#) released in Nov. 2009, the Huffman codes are predicted using 28 indirect context models. These are mixed with 3 mixers. The output of those 3 are adaptively mixed and passed through 2 SSE stages and then arithmetic encoded. All contexts are computed on Huffman code boundaries and combined with earlier bits in the same Huffman code.

The most important context is the combination of the component (Y, Cb, Cr) and the u,v coordinates (horizontal and vertical spatial frequency) to one of 192 values. JPEG uses effectively only 4 values because it uses the same Huffman tables for Cb and Cr, and the same tables for all 63 AC components. The PAQ model distinguishes them. Coefficients with large u and v tend to be smaller. Note that what is being modeled is a Huffman code representing a run of zeros along a zig-zag path followed by a nonzero coefficient with extra uncompressed bits appended.

The second most important context is a prediction based on the partial DCT in one dimension of the 8 pixels along each of the two adjacent edges in the neighboring 8 by 8 blocks to the left and above. This model was added by Jan Ondrus beginning in July 2007 (starting with [paq8fthis](#)). (Ondrus, 2009). The model reflects the constraint that neighboring pixels in adjacent blocks should have similar values. More precisely, the 1 dimensional vertical DCT of the two adjacent columns of 8 pixels in two horizontally adjacent blocks should have similar values. Likewise, the horizontal DCT of adjacent rows of pixels in vertically adjacent blocks should be similar. A partial DCT is computed using one of the two summations in the DCT equation above, for example in the horizontal direction:

$$S_u = \alpha(u) \sum_{x=0..7} S_x \cos[\pi/8 (x+1/2) u],$$

where $\alpha(0) = 1/8^{1/2}$, $\alpha(u) = 1/4$, $u > 0$.

At the beginning of a block, the context is computed for the two neighboring blocks. For the block to the left, the DCT is inverted in the horizontal direction producing 8 columns each containing a DCT of the 8 pixels in each column. The rightmost column coefficients $C_{70}..C_{77}$ provide the context. For each coefficient S_{uv} being predicted, only the coefficient C_{7v} of the neighboring block is used as context. Furthermore, the context C_{7v} is updated by computing the partial inverse DCT of the current block in the horizontal direction and incrementally subtracting from C_{7v} the coefficient of the leftmost column, S_{0v} . Thus, $C_{70}..C_{77}$ represents at all times the difference between the vertical DCT of the two neighboring columns of pixels. For a smooth image, these coefficients will all tend toward small values as the higher frequency coefficients of the current block are coded. A similar calculation is done on the top row of the current block and the bottom row of the block above.

There are many other contexts which have a smaller effect on compression. These include preceding coefficients or their RS or Huffman codes within the same block, or corresponding coefficients in neighboring blocks or in earlier components (colors) in the same block. All of these contexts may be rounded or combined to provide additional contexts. The two most important contexts are also used as contexts for the mixers and SSE stages.

6.1.6.3. WinZIP

[WinZIP](#) added a JPEG compression algorithm to version 12.0 in Sept. 2008. The format is [specified precisely](#) (with [errata](#) discovered by reverse engineering) to allow third party developers to decompress WinZIP compressed JPEG files. The algorithm is very fast, but does not compress as well as PAQ or Stuffed. On the [Maximum Compression JPEG test](#), it is ranked third at about 18% compression, vs. 24% for PAQ and Stuffed. The algorithm handles the most common formats: baseline, progressive, and extended sequential with 8 or 12 bits per pixel and 1 to 4 component colors.

WinZIP uses a transform to partially decode a JPEG image back to the DCT coefficients. The Y, Cb, and Cr components are compressed in separate streams. Within each 8 by 8 DCT block, the coefficients are coded in reverse zigzag order from highest spatial frequency (S_{77}) to DC (S_{00}) starting with the first nonzero coefficient. The DC value is predicted using a more sophisticated model and replaced with a prediction error. The coefficients are then encoded using a variable length code similar to an exponential Golomb code. The numbers are coded in the order: range, magnitude, and sign, each modeled with different contexts, with special cases for magnitudes of 0 and 1. Contexts depend on u , v , and other coefficients in the same block and the blocks to the left and above the current block. Only one context is computed for each bit. Then the context is passed to a binary arithmetic coder described in [U. S. patent 4,791,403](#) by IBM (now expired). The context model is tightly integrated into the coder. It is essentially a direct context model with a fixed learning rate. The encoder is modified to support encoding extra bits without compression with a fixed probability of 1/2.

The WinZIP algorithm describes 3 functions used to compute contexts:

- **sum(S_{uv})** or $\text{sum}(u,v)$ predicts the magnitude of S_{uv} based on the magnitude of higher frequency coefficients in the same block. $\text{sum}(u,v)$ is the sum of the quantized absolute values of all previously coded coefficients $S_{u..7,v..7}$ in the lower right corner of the coefficient array. For example $\text{sum}(5,6) = |S_{57}/Q_{57}| + |S_{66}/Q_{66}| + |S_{67}/Q_{67}| + |S_{76}/Q_{75}| + |S_{77}/Q_{77}|$ where Q_{uv} is the quantization scale factor of S_{uv} . Coefficients are represented in the JPEG as integers S_{uv}/Q_{uv} .
- **avg(u,v)** predicts the magnitude of S_{uv} based on the neighboring blocks L to the left of S and U up. $\text{avg}(u,v)$ is the average of $|U_{u,v}|$, $|U_{u,v-1}|$, $|U_{u-1,v}|$, $|U_{u-1,v-1}|$, $|L_{u,v}|$, $|L_{u,v-1}|$, $|L_{u-1,v}|$, and $|L_{u-1,v-1}|$. (The specification implies $|U_{u,v}|$, $|U_{u,v-1}|$, $|U_{u-1,v}|$, $|U_{u-1,v-1}|$, and $4|L_{uv}|$, but this is probably a typo. Source code to confirm is not available). The context only applies where $u > 0$ and $v > 0$. On the top and left edges of the image, nonexistent U_{uv} or L_{uv} are omitted from the average.
- **bdr(u,v)** is a predictor for coefficients S_{u0} and S_{0v} for AC coefficients where $\text{avg}(u,v)$ does not apply. It uses the general principle that L_{uv} predicts S_{uv} when u is even and $-S_{uv}$ when u is odd. Likewise, U_{uv} predicts S_{uv} when v is even and $-S_{uv}$ when v is odd. The prediction is weaker at higher frequencies so it is used only for $u = 0$ or $v = 0$. The predictor is:
 $\text{bdr}(0,v) = L_{0v} - L_{1v} - S_{1v}$.
 $\text{bdr}(u,0) = U_{u0} - U_{u1} - S_{u1}$.

Magnitudes uses as contexts are typically scaled logarithmically and quantized to n discrete levels using the function **cat(x, n)** = $\min(n-1, \text{ceil}(\log_2(x+1)))$.

Coefficients are represented using a variable length code of the form (zero, pivot, range, magnitude, sign) with each part coded using different contexts. The parts have the following meaning:

- zero = 0 to encode 0, in which case the other parts are omitted.
- pivot = 0 to encode -1 or 1, in which case only a sign bit follows.
- range = 0, 10, 110, 1110,... = a unary encoding of a range of size 2, 4, 8, 16...
- magnitude = $n+1$ bits, where n is the unary number encoded by the range.
- sign = 0 if positive, 1 if negative.

For example:

Number	zero	pivot	range	magnitude	sign
-2	1	1	0	1	1
-1	1	0			1
0	0				
1	1	0			0
2	1	1	0	0	0
3	1	1	0	1	0
4	1	1	10	00	0
5	1	1	10	01	0
6	1	1	10	10	0
7	1	1	10	11	0
8	1	1	110	000	0
9	1	1	110	001	0

For each block, the position of the first nonzero coefficient is coded as an ordinary 6 bit number. The context is the previous 0 to 5 bits and $\text{cat}(\text{sum}(U_{00}) + \text{sum}(L_{00}), 13)$, i. e. the sum of the absolute values of the 126 adjacent AC coefficients of the two neighboring blocks. For the AC coefficients S_{uv} , the contexts are as follows:

- For the zero bit: $u, v, \text{cat}(\text{sum}(u,v), 6)$ and either $\text{cat}(\text{bdr}(u,v), 3)$ if $u = 0$ or $v = 0$, or $\text{cat}(\text{avg}(u,v), 3)$ otherwise.
- For the pivot bit: as above with 7 and 5 levels instead of 6 and 3 respectively.
- For the range: the bit position, whether $u = 0$, whether $v = 0$, $\text{cat}(\text{sum}(u,v), 9)$ and either $\text{cat}(\text{bdr}(u,v), 9)$ or $\text{cat}(\text{avg}(u,v), 9)$.
- For the magnitude: whether $u = 0$, whether $v = 0$, and $\text{cat}(\text{zigzag}(u,v)-4, 6)$, where $\text{zigzag}(u,v)$ means the position of S_{uv} in zigzag order from 0 to 63.
- For the sign bit: a prediction is computed as $+, -, \text{ or } 0$. If 0, then the sign bit is coded with fixed probability $1/2$. Otherwise the prediction, u and v are the context. A prediction is made only for $u, v < 2$. For $u = 0$ or $v = 0$, the prediction is the sign of $\text{bdr}(u,v)$. For $u = 1$, the prediction is the sign of U_{1v} . For $v = 1$, the prediction is the sign of L_{u1} . For S_{11} the prediction is the sign of both U_{11} and L_{11} if they are the same, or 0 if they differ.

The DC coefficient is predicted by the weighted average of predictions p_U from the block above and p_L from the block to the left.

$$\text{prediction} = (w_U p_U + w_L p_L) / (w_U + w_L).$$

The prediction is estimated from the neighboring block and the gradient given by the first AC coefficient of both blocks.

$$p_U = U_{00} - 2.2076(U_{01} + S_{01})$$

$$p_L = L_{00} - 2.2076(L_{10} + S_{10})$$

The prediction weights account for the smoothness of the image in the horizontal and vertical directions. This is indicated by the absence of high frequency coefficients in the predicted direction.

$$w_U = 2^{-|\sum_{i=1..7} (|U_{i0}| - |S_{i0}|)|}$$

$$w_L = 2^{-|\sum_{i=1..7} (|L_{0i}| - |S_{0i}|)|}$$

The DC coefficient prediction error is coded as above. The context for the zero, pivot, range, and magnitude is $\text{cat}(\text{sum}(0,0), 13)$. The context for the sign is the 3 bits from the signs of the two neighboring DC signs and the sign of the prediction.

Note that contexts do not include the previously coded bits of the current number. This context is implicit in the bit position for the zero, pivot, range, and sign, but not in the magnitude. This does not hurt compression too much because the magnitude bits tend to be random.

The context and bits are fed to the IBM binary arithmetic coder. The coder has a tightly integrated context model which is essentially a direct context model with a fixed learning rate. The coder is designed to avoid multiplication and division operations by representing the range in the logarithmic domain and using addition and subtraction instead. It uses lookup tables to convert to and from the linear domain when addition or subtraction is required. In 1987 when the coder was designed, table lookup was faster than multiplication, although this is no longer true on modern processors.

The coder maintains a range $(\text{low}, \text{low}+R)$. It receives as input a bit to be coded, y , a prediction for the most probable symbol, $\text{MPS}=0$ or $\text{MPS}=1$, and a prediction, p , for the MPS between $1/2$ and 1. It always codes the MPS in the lower part of the range. Thus, most of the updates have the form $(\text{low}, R) := (\text{low}, pR)$, which is a simple addition in the log domain. In the less likely case that $y \neq \text{MPS}$, then the update is $(\text{low}, R) := (\text{low}+(1-p)R, (1-p)R)$, which requires a table lookup to convert low to the linear domain and back. Some care must be taken with the tables so that the two symbol probabilities don't add up to more than 1 due to rounding.

The coder represents the range and probabilities as 16 bit fixed point integers with 10 bit fractional parts in the linear domain and 12 bit fractional parts in the log domain. It outputs and normalizes one byte at a time. This

also requires a table lookup to convert to the linear domain.

The context model quantizes probabilities into 48 discrete values between $1/2$ and 1. This results in a coding inefficiency of less than 0.2% for most of the probability range.

Each entry in the context table contains a prediction, p , for the MPS, the choice of MPS, a counter k , and a threshold, LRM. The counter and threshold are used to decide when it is time to update the prediction. When the MPS or LPS (least probable symbol) is coded, the prediction should be increased or decreased respectively, but increasing or decreasing p by one step every time would be too fast.

Initially k is 0 and LRM is set to a fraction of the current range, R , depending on a table lookup of the current value of p . Each time an MPS is coded, the range gets smaller. When the range reaches LRM, p is increased and LRM is reset again as a fraction of R depending on p . The amount of the increase depends on k , the LPS count. If k is 5 or more, then p is unchanged. If k is in the range 2 to 4, then p is increased by 1 in the log domain, which has the effect of reducing the LPS probability $(1-p)$ by about 10% to 15%. If k is 1 then $(1-p)$ is halved. If k is 0 then $(1-p)$ is halved again.

When an LPS is coded, p should be decreased, but again not by a whole step. If k is less than 5 then it is incremented and p is unchanged. Otherwise the fraction of the distance moved by R toward LRM is used as an implicit count of MPS symbols to determine the update. If R has moved less than $1/4$ of the way to LRM then p is decremented once in the log domain, or about 10% to 15% in the linear domain. If R has moved between $1/4$ and $1/2$ of the way toward LRM, then p is decremented twice. If R has moved between $1/2$ and $3/4$ then p is halved in the linear domain (decreased by about 6 to 10). If R has moved more than $3/4$, then p is divided by 4. Then k is reset to 0 and a new LRM is chosen based on a fraction of the range depending on p . The table is designed so that for any p , the ratio of increments to decrements is $1-p$ to p , so that p reflects the true probability.

It should be noted that the IBM model and coder was developed for encoding binary and grayscale images. Because the context model is tightly integrated with the coder, it lacks the flexibility required to implement context mixing algorithms or to fine tune the learning rate for stationary or adaptive models. Its design goals were for speed, not maximum compression, at a time when computation was expensive.

6.1.6.4. PackJPG

[PackJPG](#) is a JPEG compressor by Mattias Stirner (Stirner and Seelmann, 2007). As of Dec. 2009 it was ranked fourth on the Maximum Compression JPEG benchmark with 17% compression, similar to WinZIP.

PackJPG uses a simple order 1 context model and arithmetic coder to encode the DCT coefficients. However, it changes the scan order. First, the image is divided into 4 approximately equal sized sets of blocks representing regions of high and low detail as indicated by the EOB (end of block) value. EOB indicates the last nonzero coefficient in zigzag order. EOB tends to be higher in parts of the image with a lot of fine detail. Each set of blocks is compressed as a separate stream.

Within each EOB set, the S_{uv} coefficients are encoded as separate streams. Instead of encoding the coefficients within one block in zigzag order, the scan order is from block to block for the same value of u , v , and EOB category. The scan direction depends on u and v . If $u > v$ (in the upper right corner of uv space in the first figure in [section 6.1.5](#)), then S_{uv} represents a low spatial frequency in the vertical direction and the image is scanned in columns because these coefficients are more likely to be correlated with their neighbors in the vertical direction. If $u < v$ (lower left corner of uv space), then the image is scanned horizontally. If $u = v$ (on the diagonal in uv space, including the DC coefficient), then the image is scanned in a meandering path described by a [Hilbert curve](#). In a Hilbert curve, the image is divided into 4 quadrants and scanned in a U shaped path. Each quadrant is further subdivided and scanned recursively, for example:

```

1       4
|       |
|       |
2-----3
(Quadrant scan order)

2  1 16 15
3  4 13 14
(Scan order within quadrants)
```

6 5 12 11
7 8 9 10

The DC coefficient S_{00} is first predicted using a Paeth filter. Recall that the prediction is whichever of its three neighbors is closest to (left + above - aboveleft). Then the prediction error is encoded.

Coefficients are encoded as a sign, category, and extra bits, similar to the variable length codes used in baseline JPEG. The context is the category of the previous coefficient.

The EOB list must also be stored to allow decoding. The EOB list is compressed using a 2-D context. For each block, its category is encoded using the categories of its two previously coded neighboring blocks as context.

6.2. Video Compression

Video approximates continuously moving images by using a sequence of still images, called frames. The neural circuitry of the human visual system has a delayed response to light on the order of tens of milliseconds. Thus, a frame rate of at least 24 to 30 per second produces a sensation that is nearly indistinguishable from continuous motion. However, simply flashing images at this rate would produce a noticeable [flicker](#). The eye can detect flicker at rates of up to about 75 flashes per second. Sensitivity to flicker increases in bright light, toward the blue end of the spectrum, and in peripheral vision away from the [fovea](#) where visual acuity is sharpest. Thus, a computer monitor viewed up close requires a higher refresh rate than a television viewed from a distance. Movie theatres display 24 frames per second and remove flicker by flashing each frame on the screen two to four times.

Video frames do not have to have as much resolution as still images. The eye moves in [saccades](#), jumping from one part of the image to another at a rate of 30 to 70 times per minute. In still images, the eye is attracted to regions of high contrast such as edges or corners, and to areas of interest such as faces. In text, the eye jumps from word to word. This requires all of the image to be displayed in fine detail. In video, there is not enough time to look at more than one part of a frame before the next frame is displayed. Thus, the rest of the frame can be displayed at a low resolution.

Between saccades, the eye tracks moving objects smoothly. If a frame is displayed more than once or for more than a small fraction of the frame interval, then the effect is to blur the object as the eye moves across each frame.


Time sampling of images can produce artifacts such as the [wagon wheel effect](#), where a spoked wheel appears to spin slowly backward. This artifact is analogous to the [Moire effect](#) caused by spatial sampling of a repeating pattern in still images.

6.2.1. NTSC

[NTSC](#) is one of three standards for analog television, the others being PAL and SECAM, used in different parts of the world. NTSC standardized black and white television in 1941 and color TV in 1953 in North America. It was used until 2009 for over the air broadcasts in the U.S., when it was replaced by HDTV.

NTSC is displayed at 29.97 frames per second. Each frame consists of 525 horizontal scan lines starting at the top left corner of the screen. To reduce flicker, the display is interlaced: each frame is divided into two fields which alternately display the even and odd numbered lines. (PAL and SECAM use 625 scan lines at a rate of 50 fields or 25 frames per second). NTSC is an analog format, so there is no concept of a "pixel". However, the luma (brightness) signal is transmitted over a band that extends about 4.5 MHz above the carrier. This corresponds to a Nyquist sampling rate of 9 million pixels per second, equivalent to about 571 pixels per scan line.

When color TV was introduced in 1953, a chroma signal was added without increasing the bandwidth or breaking compatibility with black and white TV sets. The spectrum allocation is shown below.



NTSC frequency allocation (from [Wikipedia](#)).

The video signal is split into three color components similar to YCbCr as in JPEG. The black-white (luma) signal is unchanged. It is amplitude modulated in the same band of 4.95 MHz. The blue-yellow and red-green signals are transmitted in a smaller band with a width of 2 MHz that overlaps the luma signal. A narrower band is possible because the eye is less sensitive to high spatial frequencies in chroma (especially blue-yellow) than luma. The two signals are amplitude modulated 90 degrees out of phase, allowing them to be separated by the receiver.

Because the luma and chroma overlap, the color signal can produce black and white artifacts and vice versa. The carrier frequencies are carefully chosen so that the artifacts of successive frames cancel out, reducing their visibility.

6.2.2. MPEG

[MPEG](#) is the most widely used format for video compression. The most commonly used versions are as follows.

- [MPEG-1](#) is the original version of the standard, published in 1993. It specifies non-interlaced video at bit rates up to 1.5 Mb/s. All patents on the video portion of the specification have expired. MPEG-1 layer 3 audio (MP3) is still patent protected.
- [MPEG-2](#) extends MPEG-1 to interlaced video and higher bit rates to support digital television. It is the format used for most DVD video and for HDTV. In spite of minor differences between MPEG-1 and MPEG-2, it is protected by about 600 patents by dozens of companies. Licenses are managed by the [MPEG Licensing Authority](#) (MPEG-LA).
- [MPEG-4 part 10](#), also known as H.264 or AVC (Advanced Video Codec) compresses video to about half the size of MPEG-1 or 2 at a similar quality level. It is widely used in [Youtube](#) and [Google Video](#). It is also patented and licensed by MPEG-LA.

Although video files can stand alone, they are more often embedded in a container format such as [AVI](#) or streamed through a [Flash](#) player. MPEG-2 defines a transport stream for over the air transmission of HDTV by encapsulating the data in [188 byte packets](#) with error correction.

MPEG-1 and MPEG-2 use a compression algorithm similar to JPEG, but obtain additional compression by delta coding between frames with motion compensation. There are 3 types of frames, designated I (inter-frame), P (predictive), and B (bidirectional). An I-frame can be decoded by itself. A P-frame is described in terms of differences from the previous frame. A B-frame is described in terms of difference from both the previous and next frame. A typical pattern is one I-frame every 0.5 seconds, repeating the sequence IBBPBBPBBPBBPBB. To facilitate decoding, the frames are sent out of order with the B frames sent after any future frame it depends on. The decoder then reorders the frames before displaying them. The reason for having I frames every 0.5 seconds is to allow decoding to start from the middle of a video after rewinding or fast forwarding.

Motion compensation is implemented by dividing a frame into 16 by 16 macroblocks. A macroblock in a P or B frame is decoded using a motion vector which points to a same sized region of the previous (or next) decoded image with a specified offset horizontally and vertically. After motion compensation, P and B frames are encoded using a JPEG-like algorithm as with I frames.

It is up to the encoder to find good matches in adjacent frames for encoding macroblocks. The encoder calculates the differences from the decoded image, not the original, by encoding and then decoding the adjacent frame.

Frame compression is like JPEG except that it uses fixed Huffman tables (called variable length codes) and fixed quantization tables with only a scale factor transmitted. MPEG is often transmitted at a constant bit rate, which is achieved by adjusting the quantization scale factor as needed. For DVDs, the maximum bit rate is about 10 Mbits/second. The result is that scenes with lots of motion are transmitted at a lower resolution.


MPEG-4/AVC (H.264) differs from MPEG-1/2 in that it supports arithmetic coding in addition to Huffman coding. It also supports variable sized macroblocks, from 4 by 4 to 16 by 16, with motion vectors pointing to any of 16 adjacent frames (or 32 fields) in 1/4 pixel resolution. Fractional motion vectors are obtained by using a 6 tap filter to infer half pixel intensities, followed by simpler interpolation. It also uses a deblocking filter.

6.3. Audio Compression

Lossy audio compression uses a [psychoacoustic model](#) to determine which part of the signal can be discarded without changing the original sound. The human ear can only perceive sounds in the range 20 Hz to 20 KHz. Sensitivity peaks around 1 KHz to 5 KHz, the middle of the range of human speech. Frequency resolution is 3.6 Hz in the range 1 KHz to 2 KHz.

Like the eye, the ear perceives sound on a logarithmic scale. The range of hearing is from 0 [decibels](#), the threshold of hearing, to 120 decibels, which is loud enough to be painful and cause hearing damage. An increment of 10 decibels (dB) represents an increase in power by a factor of 10, although we perceive it as closer to twice as loud. An increment of 20 dB represents an increase in amplitude by a factor of 10 (because power is proportional to amplitude squared).

The logarithmic scaling is partially due to the masking effect, in which a sound decreases sensitivity to other sounds at different frequencies that occur at the same time (frequency masking) or other sounds at the same frequency that occur earlier or later (temporal masking). The graph below illustrates how frequency masking affects the threshold of hearing. Temporal masking has an exponentially decaying effect, starting from about 20 milliseconds before the sound to 100 milliseconds afterward.



The effect of frequency masking on the threshold of human hearing (from [Wikipedia](#)).

Humans can perceive the direction of a sound source to an accuracy of about 3 degrees. [Stereoscopic sound perception](#) depends on two effects. First, a sound is louder in the ear closer to the source. Second, there is a time delay in reaching the further ear because sound travels at about 300 meters per second through air. Earphones can reproduce both of these effects, but stereo speakers typically do not. The sound seems to come from one speaker or the other or from some point in between. This has led to sound systems with more than 2 channels.

High frequency sounds are harder to locate when the distance between the ears is more than 1/2 wavelength because the phase shift is ambiguous and because neurons can't fire fast enough to transmit phase information. This occurs at around 1.5 KHz. This suggests an approach of transmitting stereo information (left minus right) at a lower bandwidth.

The simplest form of lossy audio compression is to filter out the high frequencies where most of the information is located. [AM radio](#) discards frequencies above 10.5 KHz. [FM](#) uses a bandwidth of 15 KHz for the mono signal (left plus right) and 13 KHz for the stereo signal.

[DS0](#) digital telephony uses a sampling rate of 8 KHz, which requires filtering out all audio above the Nyquist rate of 4 KHz. In practice, audio above 3.3 to 3.5 KHz is filtered out. Each sample is 8 bits which is [companded](#), or quantized on a logarithmic scale. Essentially, each 8 bit value is a floating point representation of a 14 bit integer (78 dB dynamic range) using a sign bit, 3 exponent bits and 4 mantissa bits. This 64 Kbit/second signal is sufficient to reproduce speech in spite of the fact that some sounds such as /s/ are almost entirely outside the bandwidth (4 to 8 KHz).

[CD audio](#) is stored uncompressed. It consists of two channels sampled at 44.1 KHz (22.05 maximum frequency) at 16 bits per sample (90 dB range). The bit rate is 1411.2 Kbit/second.

Lossy audio formats such as [MP3](#), [AAC](#), [Dolby](#), and [Ogg Vorbis](#) are based on dividing the audio into blocks of samples, computing the modified (overlapped) discrete cosine transform (MDCT), quantizing, and transmitting the coefficients without further compression. Quantization uses the psychoacoustic model to determine the

appropriate precision at each frequency. All formats support a wide range of sampling rates, compressed bit rates, and number of channels. All but Ogg Vorbis are protected by patents.

All of these formats support [joint frequency encoding](#), a technique which compresses the stereo (left minus right) signal. Because the ear can detect intensity differences but not timing differences at high frequencies, this part of the stereo signal can be removed and replaced with information to control the overall intensity for each channel.

MP3 (MPEG-1 layer III) was the first widely used compressed format for encoding music on the Internet. Audio is divided into blocks of either 576 samples, or 192 samples to encode transients (rapid changes in the audio signal). Two channels (left and right) are supported. Good quality is achieved at a bit rate of 128 Kbits/seconds, or 9% of uncompressed CD audio. Bit rates can be variable.

AAC (Advanced Audio Codec, MPEG-2 part 7) is the default audio format used by Apple's iPod and iTunes. It is understood by most music players, phones, and video game consoles. AAC supports a greater range of bit rates and sampling rates than MP3. Block sizes are 1024 and 128 samples (or 960 and 120 depending on sampling rates). Good quality audio is about 96 Kbits/second.

Dolby Digital (AC-3, ATSC A/52) is the audio format used in DVDs and HDTV. It supports [5.1 Surround Sound](#). The "5.1" refers to 5 channels (left front, right front, left rear, right rear, center) and a low frequency subwoofer channel.

Ogg is a free, open source container for the Vorbis audio format. At a bit rate of 96 Kbits/second it has an audio quality slightly better than AAC and better than MP3.

The [MDCT](#) computes N frequency coefficients from $2N$ samples $x_0..x_{2N-1}$ in 2 adjacent blocks. The overlap is necessary to prevent artifacts where the blocks join together. The coefficients are computed:

$$X_k = \sum_{i=0..2N-1} w_i x_i \cos[(\pi/N)(i + (N+1)/2)(k + 1/2)], k = 0..N-1$$

The coefficients X_k represent frequencies ranging from $1/2\pi N$ up to the Nyquist rate $1/2$. The inverse transform (IMDCT) has the same form:

$$y_i = (w_i/N) \sum_{k=0..N-1} X_k \cos[(\pi/N)(i + (N+1)/2)(k + 1/2)], i = 0..2N-1$$

The inverse transform has N inputs and $2N$ outputs. To complete the transform, the two overlapping sets of samples, y_i and y_{i+N} , from adjacent blocks are added together. To minimize boundary artifacts, the window function weights $w_{0..2N-1}$ are selected such that the weights at the ends (near 0 and $2N-1$) go to 0 and are 1 in the middle, usually with a rounded shape. Because each weight is used twice in each block, they must satisfy $w_i^2 + w_{i+N}^2 = 1$. MP3 and AAC use the window:

$$w_i = \sin[(\pi/2N)(i + 1/2)].$$

Vorbis uses:

$$w_i = \sin\{(\pi/2) \sin^2[(\pi/2N)(i + 1/2)]\}$$

Dolby AC-3 uses a [Kaiser-Bessel derived window](#), which has a similar shape.

Conclusion

Data compression is the art of finding short descriptions for long strings. Every compression algorithm can be decomposed into zero or more transforms, a model, and a coder. Coding is a solved problem. Given a symbol with probability p , Shannon proved that the best you can do is code it using $\log_2 1/p$ bits. An arithmetic coder does this efficiently.

There is no general procedure for finding good models or prediction algorithms. It is both an art and a hard problem in artificial intelligence. There is (provably) no test to tell you if a string can be compressed or if a better

model exists.

Prediction is closely related to understanding. If you understand Chinese, then you can predict a sequence of Chinese symbols. This principle can be applied to context modeling. Useful contexts are semantically independent units, for example, words in text, instructions in executable code, fields in a database, or recognizable features in an image. Different symbols that have similar meanings should be treated as if they were the same context. For example, in text, it is useful to merge upper and lower case, spaces and newlines, or related words like "someone" and "somebody" or "cold" and "wet". A context model for images would distinguish blue from green pixels but ignore fine differences. The best compressors combine the predictions of many independent models.

Preprocessing transforms are optimizations that sacrifice compression for speed and memory. Often, the output can be compressed with a simple order 0 or low order model. A transform by itself does not compress. It may hide useful contexts and add arbitrary information that makes the output ultimately larger. A good transform should minimize these effects. Thus, a BWT compressor that works on words as units would compress text better than the usual case of sorting bytes. Likewise, an LZ77 compressor that replaced duplicate strings on whole word boundaries would be preferred. Dictionary preprocessors improve both BWT and LZ77 compression by forcing those transforms to split the input on word boundaries. The best compressors on the large text benchmark use dictionary preprocessing, but I believe that is because the benchmark is tightly constrained by memory. When computers with hundreds of gigabytes become available, I believe that the top ranked programs will no longer use large dictionaries.

Lossless compressors ignore meaningless data in selecting contexts. Meaningless or random data has no predictive value and is itself not compressible. A lossy compressor not only ignores the meaningless data, but also discards it completely. Deciding which data is meaningful is a hard AI problem that applies to both lossless and lossy compression. Both require a deep understanding of human cognitive psychology.

Coding theory says that the vast majority of strings do not have simple descriptions, so it is rather remarkable that compressible strings are so common in practice. Solomonoff, Kolmogorov, and Chaitin independently proposed that strings have a universal probability proportional to $2^{-|M|}$, where M is its shortest description, independent of the language in which M is written. Kolmogorov proved that there is no algorithm for finding such descriptions in any language. Hutter showed that the compression problem, if it were computable, would solve the general AI problem of optimizing arbitrary utility functions. In effect, he proved Occam's Razor, which is the foundation of all science: the simplest theory that explains the past is the best predictor of future events.

The prevalence of compressible strings, and thus science, depends on two facts. First, that all strings are the result of computable (or finitely describable) processes, and second, that shorter programs or descriptions are more likely than longer ones. To show the second, consider a probability distribution over the infinite set of all finite length descriptions. Any such distribution must favor shorter descriptions. Consider any description M having probability $p > 0$. There can be at most a finite number $(1/p)$ of more likely descriptions, and therefore an infinite number of less likely descriptions. Of the latter, there can be at most a finite number $(2^{|M|} - 1)$ that are shorter than M . Therefore there must be an infinite number of less likely descriptions that are longer than M , for all M .

The question remains whether all strings found in the real world are created by computable or finitely describable processes. This must be true for finite strings, but there are known to exist, at least in mathematics, infinite length strings such as [Chaitin's constant \$\Omega\$](#) (the probability that a random program will halt) that are not computable. In fact, the vast majority of infinite length strings do not have finite length descriptions. Could there exist phenomena in the real world that have infinite length descriptions that are not compressible? For example, would it be possible to take an infinite number of measurements or observations, or to measure something with infinite precision? Do there exist infinite sources of random data?

The laws of physics say no. At one time it was believed that the universe could be infinitely large and made up of matter that was infinitely divisible. The discoveries of the expanding universe and of atoms showed otherwise. The universe has a finite age, T , about 13.7 billion years. Because information cannot travel faster than the [speed of light, \$c\$](#) , our observable universe is limited to an apparent 13.7 billion light years, although the furthest objects we can see have since moved further away. Its mass is limited by the [gravitational constant, \$G\$](#) , to a value that prevents the universe from collapsing on itself.

A complete description of the universe could therefore consist of a description of the exact positions and velocities of a finite number (about 10^{80}) of particles. But quantum mechanics limits any combination of these two quantities to discrete multiples of [Planck's constant, \$h\$](#) . Therefore the universe, and everything in it, must have a finite description length. The entropy in [nats](#) ($1 \text{ nat} = 1/\ln(2) \text{ bits} = 1.4427 \text{ bits}$) is given by the [Bekenstein bound](#) as $1/4$ of the area of the event horizon in [Planck units](#) of area $hG/2\pi c^3$, a square of 1.616×10^{-35} meters on a side. For a sphere of radius $T_c = 13.7$ billion light years, the bound is 2.91×10^{122} bits.

We now make two observations. First, if the universe were divided into regions the size of bits, then each volume would be about the size of a proton or neutron. This is rather remarkable because the number is derived only from T and the physical constants c , h , and G , which are unrelated to the properties of any particles. Second, if the universe were squashed flat, it would form a sheet about one neutron thick. Occam's Razor, which the computability of physics makes true, suggests that these two observations are not coincidences.

Acknowledgements

I thank Ivo Danihelka, Glenn Davis, Szymon Grabowski, Aki Jäntti, Harri Hirvola, Christopher Mattern, Ilia Muraviev, Michael Newman, Jan Ondrus, Carol Porter, Alexander Rhatushnyak, Friedrich Regen, Steve Richfield, Sami Runsas, David A. Scott, Eugene Shelwien, Ali Imran Khan Shirani, Yan King Yin, and Bulat Ziganshin for helpful comments on this book. Thanks to Osman Turan for converging this book to XHTML compliant.

Please send corrections or comments to mattmahoneyfl@gmail.com.

References

- T. Bell, I. H. Witten. J. G. Cleary (1989), Modeling for Text Compression, ACM Computing Surveys (21)4, pp. 557-591.
- C. Bloom (1998), [Solving the Problems of Context Modeling](#).
- M. Burrows, D. J. Wheeler (1994), [A Block-sorting Lossless Data Compression Algorithm](#), Digital Systems Research Center.
- G. Chaitin (1966), On the length of programs for computing finite binary sequences. Journal of the ACM, 13:547-569.
- G. Cleary, W. J. Teahan (1995), Experiments on the zero frequency problem, Proc. Data Compression Conference, 480.
- G. Cormack, N. Horspool (1987), Data compression using dynamic Markov modeling, Computer Journal 30:6 (December).
- P. Deustch (1996), [RFC 1951 - DEFLATE Compressed Data Format Specification version 1.3](#).
- D. Dube, V. Beaudoin (2006), [Improving LZ77 data compression using bit recycling](#), Seoul: International Symposium on Information Theory and its Applications. ISITA2006.
- D. Dube, V. Beaudoin (2010), [Lossless data compression via substring enumeration](#), Proc DCC.
- J. Duda (2007), [Optimal encoding on discrete lattice with translational invariant constraints using statistical algorithms](#) (section 3).
- M. Gagliolo (2007), [Universal search](#). Scholarpedia, 2(11):2575.
- S. Grabowski, J. Swacha (2010), Language-independent word-based text compression with fast decompression, Proc. MEMSTECH, Polyana, Ukraine.
- D. Huffman (1952), [A Method for the Construction of Minimum-Redundancy Codes](#). Proc. I.R.E.: 1098-1101.

- M. Hutter (2004), Universal artificial intelligence: Sequential decisions based on algorithmic probability. Springer, Berlin.
- M. Hutter et al. (2007), [Algorithmic probability](#). Scholarpedia, 2(8):2572.
- H. Itoh, H. Tanaka (1999), An efficient method for in memory construction of suffix arrays, Proceedings of the IEEE String Processing and Information Retrieval Symposium, pp. 81-88.
- P. Ko, S. Aluru (2003), Space-efficient linear time construction of suffix arrays, Proceedings of the 14th Annual Symposium on Combinatorial Pattern Matching, pp. 200-210.
- A. Kolmogorov (1965), Three approaches to the quantitative definition of information. Problems Inform. Transmission, 1, 1-7.
- M. Kufleitner (2009) [On bijective variants of the Burrows-Wheeler transform](#), Proc. Prague Stringology Conference.
- T. Landauer (1986), [How much do people remember? Some estimates of the quantity of learned information in long term memory](#), Cognitive Science (10) 477-493.
- N. J. Larsson, K. Sadakane (1999), [Faster suffix sorting](#), Technical report LU-SC-TR:99:214, Lund University, Sweden.
- S. Legg, M. Hutter (2006), [A Formal Measure of Machine Intelligence](#), Proc. Annual machine learning conference of Belgium and The Netherlands (Benelearn-2006). Ghent.
- L. A. Levin (1973), Universal sequential search problems. Problems of Information Transmission, 9(3):265--266.
- L. A. Levin (1984), Randomness Conservation Inequalities: Information and Independence in Mathematical Theories. Information and Control, 61:15-37.
- M. Li and P.M.B. Vitanyi (2007) [Applications of algorithmic information theory](#). Scholarpedia, 2(5):2658.
- M. Mahoney (2000), [Fast Text Compression with Neural Networks](#), Proc. AAAI FLAIRS, Orlando.
- M. Mahoney (2002), [The PAQ1 Data Compression Program](#)
- M. Mahoney (2005), [Adaptive Weighing of Context Models for Lossless Data Compression](#), Florida Tech. Technical Report CS-2005-16.
- G. Manzini and P. Ferragina (2002) [Engineering a lightweight suffix array construction algorithm \(extended abstract\)](#).
- Y. Matias, S. C. Sahinalp (1999), [On the Optimality of Parsing in Dynamic Dictionary Based Data Compression](#), Proc. tenth annual ACM-SIAM symposium on Discrete algorithms, 943-944.
- C. Mattern (2012), [Mixing Strategies in Data Compression](#), Proc. 22'nd IEEE Data Compression Conference.
- D. T. Meyer, W. J. Bolosky (2011), [A Study of Practical Deduplication](#), FAST-2011.
- G. Nong, S. Zhang, W. H. Chan (2009) [Linear suffix array construction by almost pure induced sorting](#), Proc. 19'th IEEE Data Compression Conference.
- J. Ondrus (2009), [Beztrátová komprese JPEG grafiky](#) (Lossless compression of JPEG images), thesis, Univerzita Karlova v Praze (in Czech).
- S. J. Puglisi (2005), [Exposition and analysis of a suffix sorting algorithm](#), Technical Report Number CAS-05-02-WS (May, 2005), Dept of Computing and Software, McMaster University, Hamilton, Ontario, Canada.
- S. J. Puglisi, M. A. Maniscalco (2006a), Faster lightweight suffix array construction, 17th Australasian Workshop on Combinatorial Algorithms (AWOCA'06), Joe Ryan & Dafik (eds.) pp. 16-29.

- S. J. Puglisi, M. A. Maniscalco (2006b), An efficient, versatile approach to suffix sorting, ACM Journal of Experimental Algorithmics.
- S. J. Puglisi, W. F. Smyth, A. H. Turpin (2007), [A taxonomy of suffix array construction algorithms](#), ACM Computing Surveys 39(2).
- J. Rissanen (1976), Generalized Kraft Inequality and Arithmetic Coding, IBM Journal of Research and Development 20(3): 198–203.
- J. Schmidhuber, S. Heil (1996), Sequential Neural Text Compression, IEEE Trans. on Neural Networks 7(1): 142-146.
- C. Shannon, W. Weaver (1949), The Mathematical Theory of Communication, Urbana: University of Illinois Press.
- C. E. Shannon (1950), Prediction and Entropy of Printed English, Bell Sys. Tech. J. 3:50-64.
- D. Shkarin (2002), [PPM: one step to practicality](#), proc. DCC.
- P. Skibiński, S. Grabowski (2004), [Variable-length contexts for PPM](#), Proc. Data Compression Conference (DCC).
- P. Skibiński, S. Grabowski, S. Deorowicz (2005), [Revisiting dictionary-based compression](#), *Software - Practice and Experience*, Wiley.
- R. Solomonoff (1960), A Preliminary Report on a General Theory of Inductive Inference, Report V-131, Zator Co., Cambridge, Ma.
- R. Solomonoff (1964), A Formal Theory of Inductive Inference, Information and Control, 7(1) 1-22, 7(2) 224-254.
- M. Stirner and G. Seelmann (2007), [Improved Redundancy Reduction for JPEG Files](#). Lisboa: Picture Coding Symposium CD-ROM Proceedings.
- A. M. Turing (1950), [Computing Machinery and Intelligence](#), Mind, 59:433-460.
- E. Ukkonen (1995), [On-line construction of suffix trees](#), Algorithmica 14(3): 249-260.
- I. H. Witten, T. C. Bell (1991), The zero-frequency problem: estimating the probabilities of novel events in adaptive text compression, IEEE Trans. on Information Theory, 37(4): 1085-1094.
- J. Ziv, A. Lempel (1977), [A universal algorithm for sequential data compression](#), IEEE Trans. Information Theory 23(3) 337-343.
- J. Ziv, A. Lempel (1978), [Compression of Individual Sequences via Variable-Rate Coding](#), IEEE Transactions on Information Theory, 24 (5), 530-536.
- Zongjie Tu, Shiyong Zhang (2006) [A Novel Implementation of JPEG 2000 Lossless Coding Based on LZMA](#), Proc. Sixth IEEE International Conference on Computer and Information Technology, 140.

Note: Website links are current as of Feb. 2010.